

# Introduction to Hardware Programming

# Intro to Hardware Programming

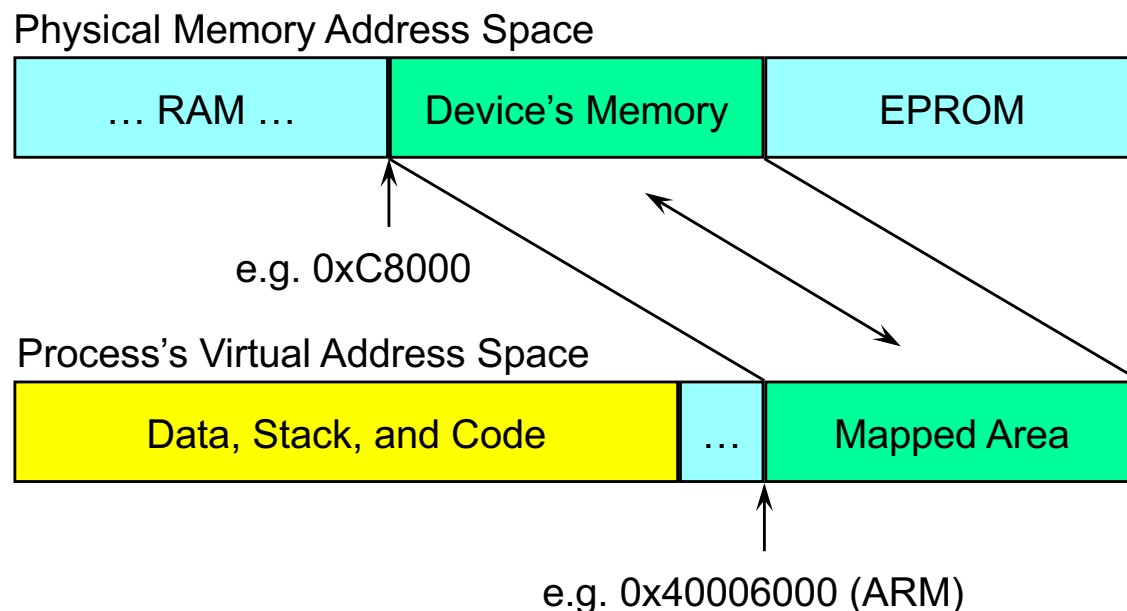
## Topics:

- **Hardware I/O**
- Programming PCI bus devices**
- Handling Interrupts**
- Conclusion**

# Memory Interface

To access memory on a hardware device:

- physical addresses must be mapped into your process's virtual address space:



```
// get a pointer to memory physically located at 0xc8000
vaddr = mmap_device_memory (0, 0x4000,
                             PROT_READ | PROT_WRITE | PROT_NOCACHE,
                             0, 0xc8000);
```

## DMA Safe Memory

DMA operations usually require physically contiguous RAM:

```
// for DMA:
// allocate len bytes of physically contiguous
// system memory

vaddr = mmap (0, len,
              PROT_READ | PROT_WRITE | PROT_NOCACHE,
              MAP_PHYS | MAP_ANON | MAP_SHARED, NOFD, 0);

// get the physical address for passing to the controller
mem_offset64 (vaddr, NOFD, len, &paddr, NULL);
```

# How you access control registers varies based on the platform:

- on ARM and AARCH64 platforms:
  - registers are mapped like memory
  - pointer dereferences are used for read/write operations
  - procnto knows based on system page that these are not RAM, sets CPU access flags differently
  - some x86\_64 devices may, also, be accessed this way
- on x86\_64 most devices continue to use “x86-style” I/O ports
  - a special address line must be active to change from memory to I/O addressing
  - requires a special set of assembly instructions for access

### Interfacing to I/O ports (x86\_64):

- QNX supplies cover functions for the inline assembly needed

```
#include <hw/inout.h> // header for in*() & out*() fns
```

```
// enable I/O privilege for this thread
```

```
ThreadCtl (_NTO_TCTL_IO, NULL);
```

```
val8  = in8  (ioport_addr); // read an 8 bit value
```

```
val16 = in16 (ioport_addr); // read a 16 bit value
```

```
val32 = in32 (ioport_addr); // read a 32 bit value
```

```
out8  (ioport_addr, val8 ); // write an 8 bit value
```

```
out16 (ioport_addr, val16); // write a 16 bit value
```

```
out32 (ioport_addr, val32); // write a 32 bit value
```



# Intro to Hardware Programming

## Topics:

**Hardware I/O**

**→ Programming PCI bus devices**

**Handling Interrupts**

**Conclusion**

To find and configure a PCI device:

- you must run the PCI server:

`pci-server`



## PCI - The calls

### The PCI calls include:

<code>pci_device_find()</code>	find hardware by Device ID and Vendor ID
<code>pci_device_attach()</code>	find hardware and get basic configuration information
<code>pci_device_reset()</code>	reset device
<code>pci_device_detach()</code>	detach device
<code>pci_device_read_*</code>	read configuration information
<code>pci_device_write_*</code>	write to device command or status register
<code>pci_device_map_as()</code>	translate between CPU and PCI addresses

The *pci\_device\_find()* call:

- fills in a `pci_bdf_t` base data type
  - encodes the bus, device and function of the PCI device
- various calls can determine
  - interrupt number
  - address translations for bus master PCI devices
  - PCI configuration space registers
  - the base addresses of a device

# Intro to Hardware Programming

## Topics:

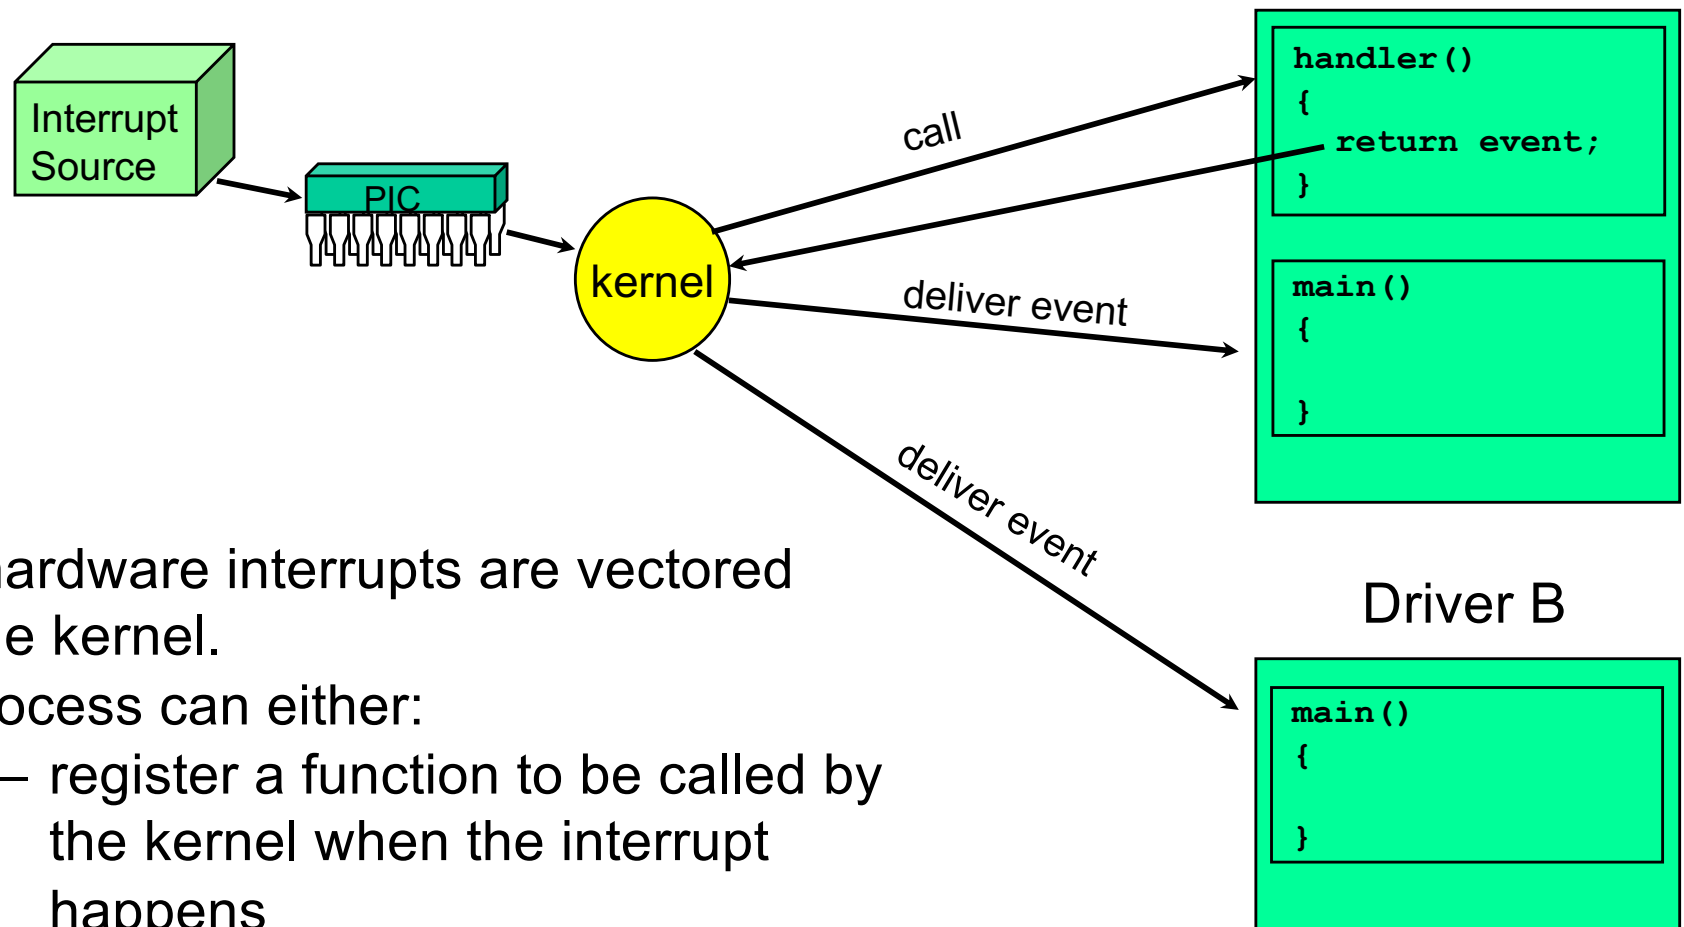
**Hardware I/O**

**Programming PCI bus devices**

**→ Handling Interrupts**

**Conclusion**

## Interrupt Handling: Two choices



All hardware interrupts are vectored to the kernel.

A process can either:

- register a function to be called by the kernel when the interrupt happens
- request notification that the interrupt has happened

# Interrupts - The Calls

## Interrupt calls:

```
id = InterruptAttach (int intr,  
                     struct sigevent *(*handler)(void *, int),  
                     void *area, int size, unsigned flags);  
id = InterruptAttachEvent (int intr, struct sigevent *event,  
                           unsigned flags);  
InterruptDetach (int id);  
InterruptWait (int flags, uint64_t *reserved);  
InterruptMask (int intr, int id);  
InterruptUnmask (int intr, int id);  
InterruptLock (struct intrspin *spinlock);  
InterruptUnlock (struct intrspin *spinlock);
```



Permissions are complicated... see next slide.



## Interrupt Calls – System Privileges

The privileges required for interrupt calls are complex:

- *InterruptAttach()*
  - **PROCMGR\_AID\_INTERRUPT**
- *InterruptAttachEvent()*
  - **PROCMGR\_AID\_INTERRUPTEVENT** or **PROCMGR\_AID\_INTERRUPT**
- *InterruptMask()*, *InterruptUnmask()*
  - I/O privilege is required to mask an interrupt the process has not attached
- *InterruptLock()*, *InterruptUnlock()*, *InterruptDisable()*, *InterruptEnable()*
  - I/O privilege which is gotten by calling:
- *ThreadCtl(\_NTO\_TCTL\_IO, 0):*
  - **PROCMGR\_AID\_IO**

## Handling an Interrupt

# Driver A example: interrupt handler function

```
struct sigevent event;

const struct sigevent *
handler (void *not_used, int id)
{
    if (check_status_register())
        return (&event);
    else
        return (NULL);
}

main ()
{
    SIGEV_INTR_INIT (&event);
    id = InterruptAttach (INTNUM, handler, NULL, 0, ...);
    for (;;) {
        InterruptWait (0, NULL);
        // do some or all of the work here
    }
}
```



# Handling an Interrupt

## Driver B example: interrupt event loop

```
struct sigevent event;

main ()
{
    SIGEV_INTR_INIT (&event);
    id = InterruptAttachEvent (INTNUM, &event, ...);
    for (;;) {
        InterruptWait (0, NULL);
        // do the interrupt work here, at thread priority
        InterruptUnmask (intnum, id);
    }
}
```





# InterruptAttachEvent

Telling kernel what code to run when an interrupt happens:

```
id = InterruptAttachEvent (intr, event, flags);
```

logical interrupt  
vector number

↑  
tell kernel what wakeup  
event to give us

additional  
information  
flags

↓  
handler function  
to run on interrupt

↓  
mem passed  
into handler

```
id = InterruptAttach (intr, handler, area, size, flags);
```

## InterruptAttach\*() Flags

*InterruptAttachEvent()* and *InterruptAttach()*'s flags parameter can contain:

— **NTO\_INTR\_FLAGS\_END**: If multiple interrupt handlers, specify we should execute last

— **NTO\_INTR\_FLAGS\_PROCESS**: for events types that are directed at a process rather than a thread, e.g. pulses

— **NTO\_INTR\_FLAGS\_TRK\_MSK**: request that the kernel adjust the interrupt mask when the attaching process terminates (ALWAYS set this flag)

— **NTO\_INTR\_FLAGS\_NO\_UNMASK**: must explicitly unmask the interrupt to enable

# An interrupt handler operates in the following environment:

- it is sharing the data area of the process that attached it
- the environment is very restricted:
  - cannot call kernel functions except *InterruptMask()*, *InterruptUnmask()* and *TraceEvent()* (see notes)
  - cannot call any function that might call a kernel function
    - the documentation for each function specifies whether or not that function is safe to call from an interrupt handler
    - there is also a section in the Library Reference manual called “Full Safety Information” that lists all safe functions
  - the interrupt handler is using the kernel’s stack, so keep stack usage small (if you have a lot of data, use variables defined outside of the handler rather than variables defined local to the function, and don’t call too many function levels deep)
  - shouldn’t do floating point or equivalent operations



# Should you attach a handler or an event?

- The kernel is the single point of failure for a QNX system; attaching a handler increases the size of the SPOF, an event does not
- debugging is far simpler with an event
  - ISR code can not be stepped/traced with the debugger
- full OS functionality when doing h/w handling in a thread
- events impose far less system overhead at interrupt time than handlers
  - no need for the MMU work to gain access to process address space if using an event
- scheduling a thread for every interrupt could be more overhead, if you could do some work at interrupt time and only need to schedule a thread some of the time
- handlers have lower latency than getting a thread scheduled
  - does your hardware have some sort of buffer or FIFO? If not, then you might not be able to wait until a thread is scheduled

You have the following notification methods:

– SIGEV\_INTR

- unblocks *InterruptWait()*
- simplest to use (least setup)
- fastest (lowest overhead, latency)
- not queued or counted
- must dedicate a thread

– SIGEV\_SEM

- unblocks *sem\_wait()* on a named semaphore
  - within driver, use an anonymous named semaphore
- counted
- only direct cross-process choice

*continued...*

### Notification methods (continued):

#### – SIGEV\_SIGNAL

- unblocks *sigwaitinfo()*
  - Do not use a signal handler, the overhead and latency are awful
- can be queued
- can carry data

#### – SIGEV\_PULSE

- unblocks *MsgReceive\*()*
- queued
- carries data
- most flexible
  - single threaded driver can handle hardware and clients
  - pool of threads
- highest overhead and latency

## EXERCISE

### Simple interrupt handler:

- in your `interrupt` project is a skeleton file called `intsimple.c`
- fill it in with the code for handling interrupts, the instructor will tell you which interrupt to attach to
- attach an interrupt handler that will return a `SIGEV_INTR` event
- in the loop, use *InterruptWait()* to wait for the interrupt notification

# Intro to Hardware Programming

## Topics:

**Hardware I/O**

**Programming PCI bus devices**

**Handling Interrupts**

**→ Conclusion**



## Conclusion

### You learned:

- That memory or port mappings have to be set up to access hardware devices
- that the kernel is the first handler for all interrupts
- that processes can register handlers or can register for notification of interrupts
- that interrupt handlers run in a very restricted environment