

Introduction to Resource Managers

You will learn:

- what a resource manager is
- how to use the QNX Neutrino resource manager framework:
 - initialization
 - handling read and write

Introduction to Resource Managers

Topics:

→ Overview

A Simple Resource Manager

- Initialization
- Handling *read()* and *write()*

Conclusion


What is a resource manager?

- a program that looks like it is extending the operating system by:
 - creating and managing a name in the pathname space
 - providing a POSIX interface for clients (e.g. *open()*, *read()*, *write()*, ...)
- can be associated with hardware (such as a serial port, or disk drive)
- or can be a purely software entity (such as queuing or logging)

Let's take a look at the pathname space

Overview

Name mapping:



RESMGR

PATHNAMESPACE

RESMGR

The prefix tree:


- is the root of the pathname space
 - every name in the pathname space is a descendant of some entry in the prefix tree
- is maintained by the Process Manager
 - stored as a table
- Resource Managers add and delete entries
- associates a **pid**, **chid**, and **handle** with a name
- is searched for the longest slash-delimited whole-word matching prefix



Overview

For example, to resolve the pathname:
`/dev/ser1`

```
fd = open("/dev/ser1", ...);
```




The longest match is `/dev/ser1`, which points to `devc-ser8250`

Overview

Or,

/home/bill/spud.dat

```
fd = open("/home/bill/spud.dat", ...);
```



The longest match is /, which points to **procnto** and **fs-qnx6.so**. **procnto** would fail the open, and **fs-qnx6.so** would then handle requests.

Overview

A Client requests a service:

```
fd = open ("/dev/ser1", O_RDWR) ;
```

or


```
fp = fopen ("/home/bill/abc", "w") ;
```

or

```
mq = mq_open ("/myqueue", O_RDONLY) ;
```

which results in the client's library code
sending a message to the process
manager...

Interactions:



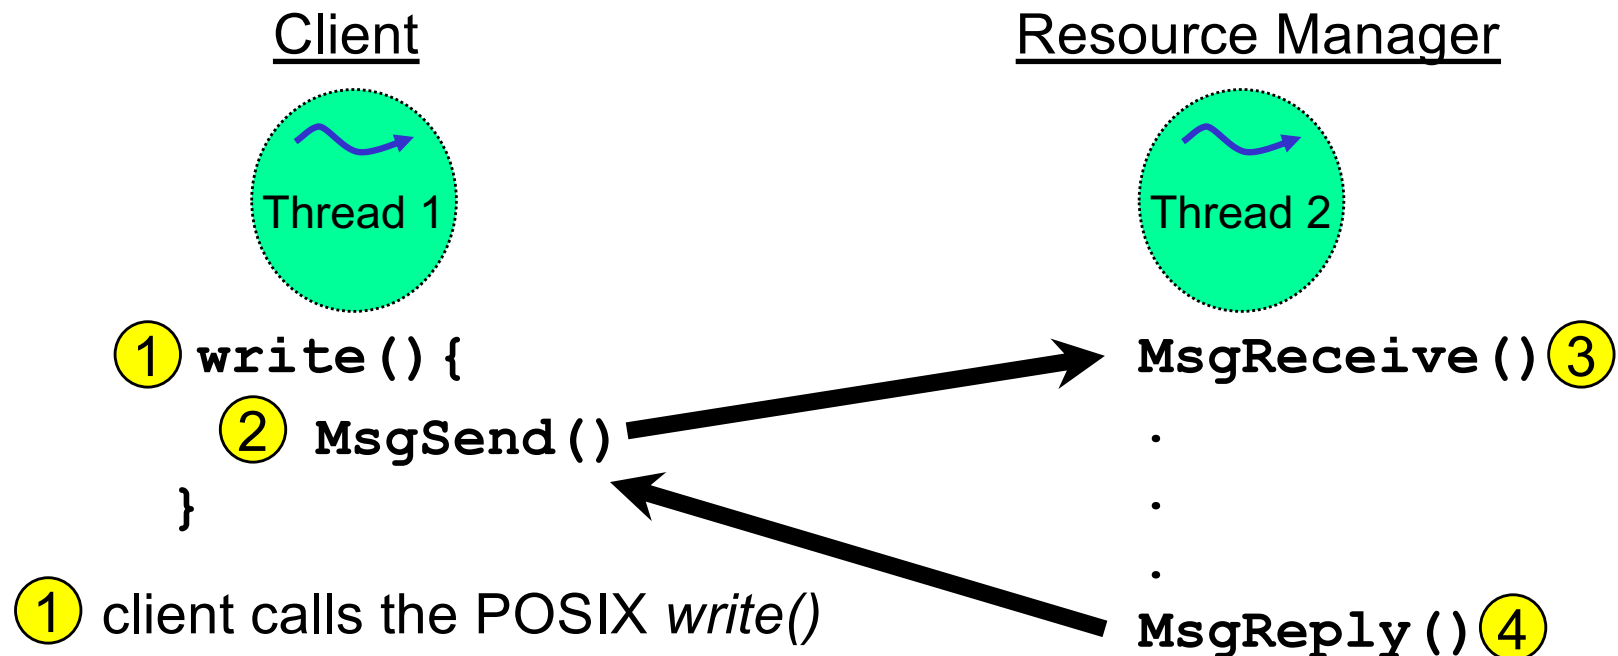
- 1 `open()` sends a "query" message
- 2 Process Manager replies with who is responsible (`pid`, `chid`, `handle`)
- 3 `open()` establishes a connection to the specified resource manager (`pid`, `chid`), and sends an open message (containing the `handle`)
- 4 Resource manager responds with status (pass/fail)

All further communication goes directly to the resource manager.



Resource Managers

Resource managers are built on message passing:



- ① client calls the POSIX *write()*
- ② this builds and sends a message
- ③ the resource manager receives and processes the **write** message appropriately
- ④ the resource manager replies with a result and this is returned to the client from *write()*



A resource manager performs the following steps:

- creates a channel
- takes over a portion of the pathname space
- waits for messages & events
- processes messages, returns results

There are three major types of messages:

Connect messages:

- pathname-based (eg: `open ("spud.dat", ...)`)
- may create an association between the client process and the resource manager, which is used later for I/O messages

I/O messages:

- file-descriptor- (`fd`-) based (eg: `read (fd, ...)`)
- rely on association created previously by connect messages

Other:

- pulses, private messages, etc

Resource Manager Messages

Connect Messages:

client call

open()

unlink()

rename()

message

`_IO_CONNECT`

`_IO_UNLINK`

`_IO_RENAME`

Defined in `<sys/iomsg.h>`

Resource Manager Messages

I/O Messages (frequently used):

client call

read()

write()

close()

devctl(), ioctl()

message

`_IO_READ`

`_IO_WRITE`

`_IO_CLOSE`

`_IO_DEVCTL`

Defined in `<sys/iomsg.h>`

continued...

Resource Manager Messages

I/O Messages (continued):

<code>_IO_NOTIFY,</code>	<code>_IO_STAT,</code>
<code>_IO_UNBLOCK,</code>	<code>_IO_PATHCONF,</code>
<code>_IO_LSEEK,</code>	<code>_IO_CHMOD,</code>
<code>_IO_CHOWN,</code>	<code>_IO_UTIME,</code>
<code>_IO_LINK,</code>	<code>_IO_FDINFO,</code>
<code>_IO_LOCK,</code>	<code>_IO_TRUNCATE,</code>
<code>_IO_SHUTDOWN,</code>	<code>_IO_DUP</code>

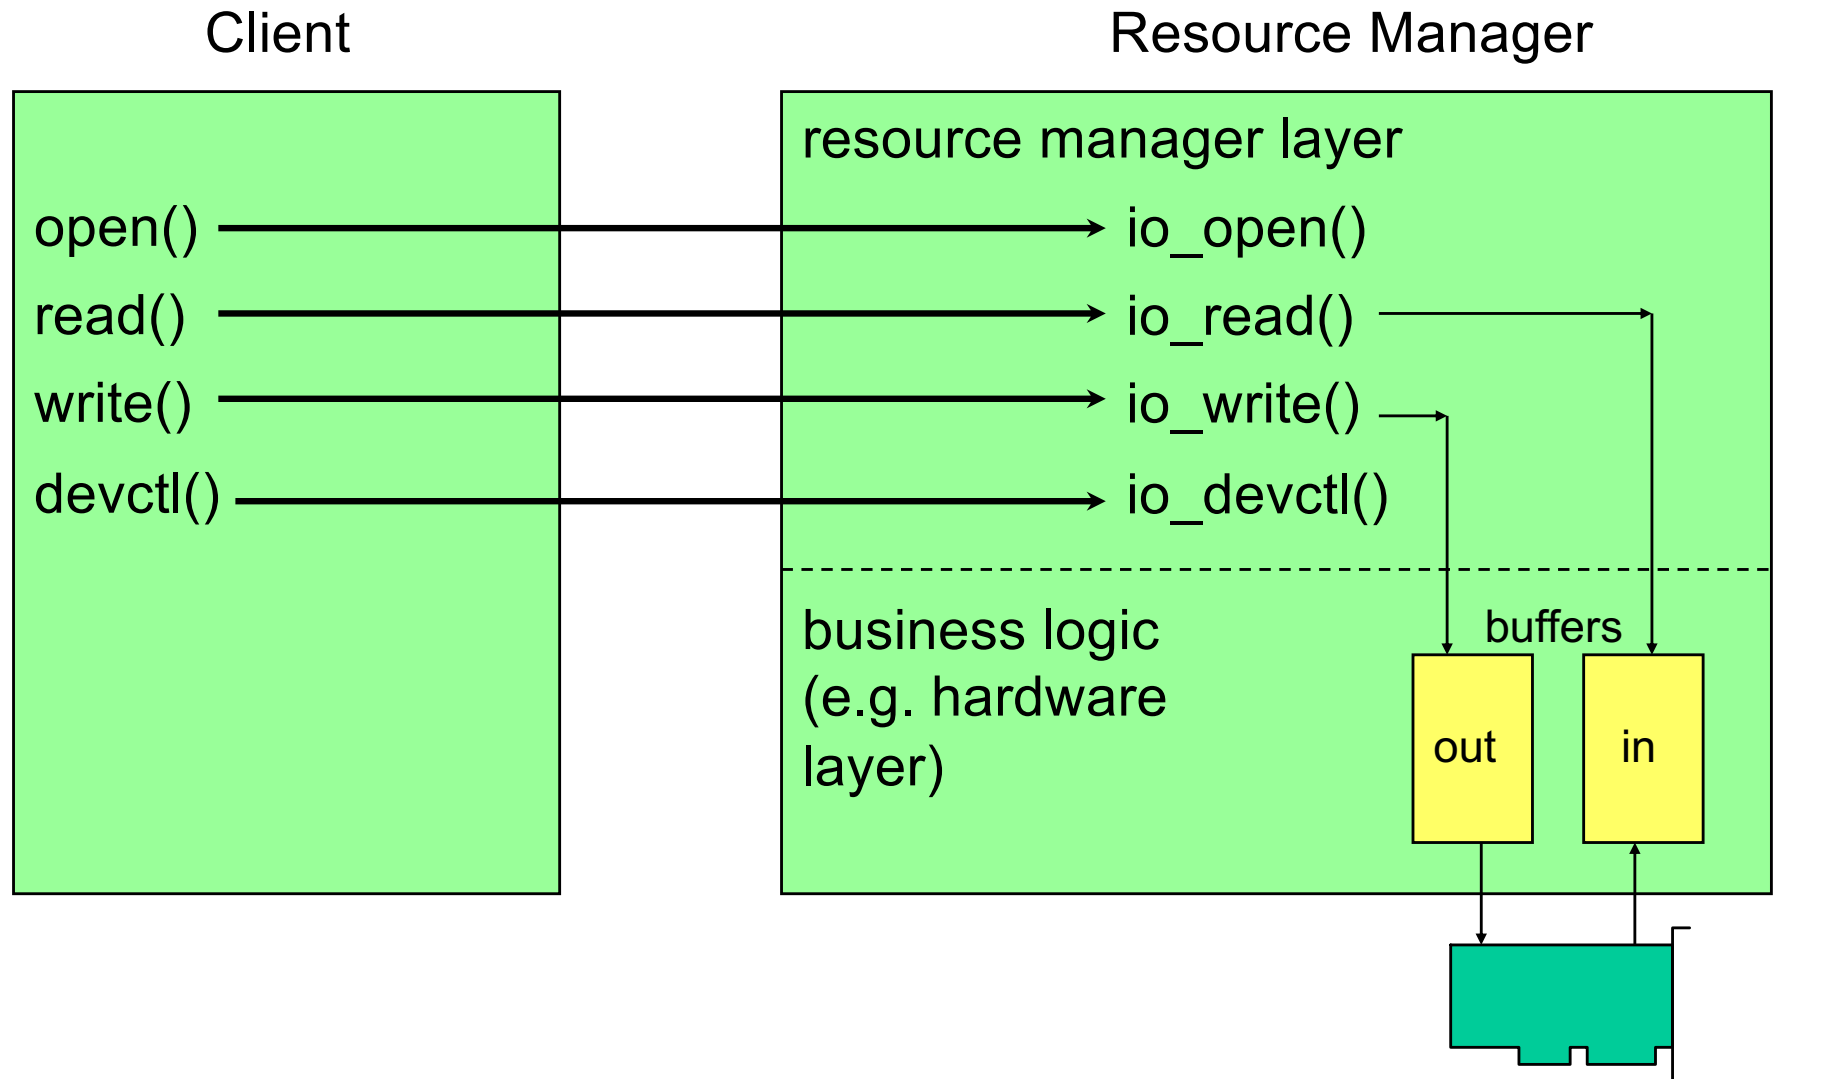
Writing resource managers is simplified greatly with a resource-manager shared library that:

- simplifies main receive loop (table-driven approach)
- has default actions for any message types that do not have handlers specified in tables



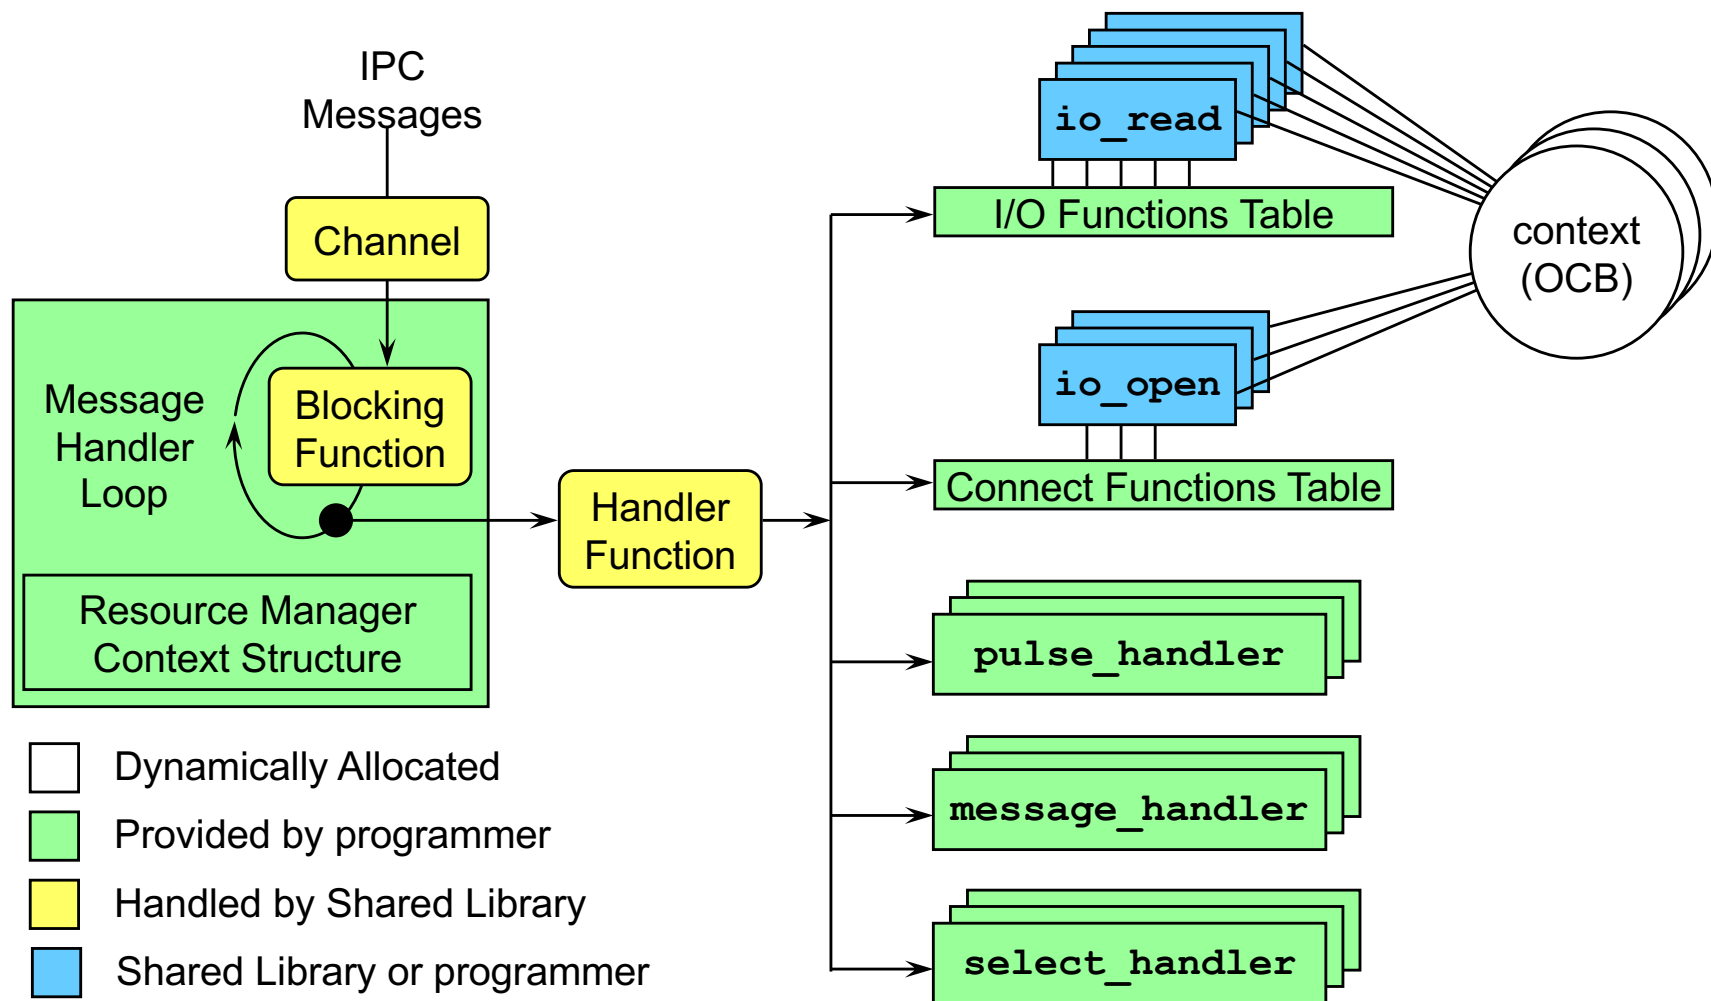
The Big Picture

Client calls go to handlers:



Resource Manager Structure

The resource manager layer:



Message-handling functions:

- may access client data
 - data already in the receive buffer
 - get more data if necessary
- must either:
 - reply to the client with:
 - an error
 - a success status without data
 - a success status and data
- or:
 - delay responding to the client until later



Introduction to Resource Managers

Topics:

Overview

A Simple Resource Manager

- – Initialization
- Handling *read()* and *write()*

Conclusion

The example Resource Manager

To talk about setting up a resource manager, we'll use an **example** resource manager:

Client side:

- Here's how it behaves from a client's point of view:
 - **read** always returns 0 bytes
 - **write** of any size always works
 - other things behave as expected



The example resource manager:

- create & initialize structures:
 - a dispatch structure
 - list of *connect* message handlers
 - list of *I/O* message handlers
 - device attributes
 - resource-manager attributes
 - dispatch context
- attach a pathname, passing much of the above
- from the main loop:
 - block, waiting for messages
 - call a handler function; the handler function handles requests and performs callouts to your specified routines.

Setting things up – Dispatch structure

First, create a dispatch structure:

```
dispatch_t *dpp;
```

dpp

```
dpp = dispatch_create_channel(chid, flags);
```

- this is the glue the resource manager framework uses to hold everything together
- the contents are hidden (it is an opaque type)

- This is usually called with the parameters:

```
dispatch_create_channel(-1, DISPATCH_FLAG_NOLOCK);
```

- create a channel
- disable some, usually unnecessary, mutex locks during message handling



Setting things up - Connect and I/O functions

Next, we set up two tables of functions:

- connect functions


- these are called as a result of POSIX calls that take a filename

e.g.: *open (filename, ...), unlink (filename), ...*

- I/O functions

- these are called as a result of POSIX calls that take a file descriptor


e.g.: *read (fd, ...), write (fd, ...), ...*



Setting things up - Connect and I/O functions - Example

Example of declaring and initializing the connect- and the I/O-functions structures:

```
resmgr_connect_funcs_t connect_funcs;  
resmgr_io_funcs_t      io_funcs;  
  
iofunc_func_init (_RESMGR_CONNECT_NFUNCS, &connect_funcs,  
                  _RESMGR_IO_NFUNCS, &io_funcs);  
  
connect_funcs.open = io_open;  
io_funcs.read      = io_read;  
io_funcs.write     = io_write;
```



iofunc_func_init() places default values into the passed connect- and I/O-functions structures, based on the number of values that you have specified via the first and third integer arguments. It is recommended that you use the **`_RESMGR_CONNECT_NFUNCS`** and **`_RESMGR_IO_NFUNCS`** constants for those two arguments.



Setting things up - `iofunc_attr_t`

Next, fill the device-attributes structure, for passing to `resmgr_attach()`:

```
iofunc_attr_t ioattr;
```

device structure
`iofunc_attr_t`


```
iofunc_attr_init (&ioattr, S_IFCHR | 0666, NULL,  
                 NULL) ;
```

- this is your device-specific data structure
- used by the *iofunc_*()* helper functions
- it is possible to extend this structure so that it can contain your own data too



Putting /dev/example into the pathname space

Next, put /dev/example into pathname space:



```
resmgr_attach(..., "/dev/example", ...);
```

Setting things up - resmgr_attach() details

The parameters are:

```
id = resmgr_attach (dpp, &rattr, path, file_type,  
                    flags, &connect_funcs, &io_funcs, handle);
```

dpp = pointer returned by *dispatch_create_channel()*

rattr = **NULL** or structure of further parameters

path = `"/dev/example"`

file_type = **_FTYPE_ANY** (the usual case)

flags = 0 or control flags...

connect_funcs and **io_funcs** point to the tables of functions we just created

handle = pointer to device attributes

id = id of this pathname, used for *resmgr_detach()* call

👉 Requires **PROCMGR_AID_PATHSPACE** to succeed.



The resmgr_attach(..., flags, ...):

_RESMGR_FLAG_BEFORE and

_RESMGR_FLAG_AFTER

this resource manager will handle the pathname
BEFORE or AFTER all others that have attached the
same pathname

_RESMGR_FLAG_DIR

allow pathnames that extend past the registered
pathname to be handled by this resource manager
used by filesystem resource managers (e.g.:
/cdrom/...)




resmgr_attach()

When you call *resmgr_attach()*, you are:

- creating your device
- associating data and handlers with it

```
id = resmgr_attach (dpp, NULL, "/dev/example", _FTYPE_ANY,  
                  NULL, &connect_funcs, &io_funcs, &ioattr);
```



👉 The library does not make copies of these structures

resmgr_attach()


resmgr_attach() puts the name in the pathname space:

- your resource manager becomes visible to clients
 - clients will expect you to be ready to handle messages
- before doing so, you should have completed most of your initialization:
 - hardware detection and initialization
 - buffer allocation and configuration
- if something fails, don't attach name

Setting things up - dispatch_context_alloc()

Lastly, allocate a dispatch context structure:

```
dispatch_context_t *ctp;  
ctp = dispatch_context_alloc (dpp);
```



- this is the operating parameters of the message receive loop
- it is passed to the blocking function and the handler function
- it contains things like the `rcvid`, pointer to the receive buffer, and message info structure
- it will be passed as the `ctp` parameter to your connect and I/O functions

Setting things up - What we have so far

Putting together what we have so far:

```
dpp = dispatch_create_channel (-1, DISPATCH_FLAG_NOLOCK);

iofunc_func_init (_RESMGR_CONNECT_NFUNCS, &connect_funcs,
                  _RESMGR_IO_NFUNCS, &io_funcs);
connect_funcs.open = io_open;
io_funcs.read      = io_read;
io_funcs.write     = io_write;

iofunc_attr_init (&ioattr, S_IFCHR | 0666, NULL,
                  NULL);

id = resmgr_attach (dpp, NULL, "/dev/example", _FTYPE_ANY,
                   0, &connect_funcs, &io_funcs, &ioattr);

ctp = dispatch_context_alloc (dpp);
```

Setting things up - The dispatch loop

So now that everything's set, the loop:

```
while (1) {  
    dispatch_block (ctp);  
    dispatch_handler (ctp);  
}
```

- *dispatch_block()* blocks waiting for messages,
- *dispatch_handler()* handles them, including calling any callout functions which you've provided (connect function, I/O functions)

EXERCISE

Exercise:

- in your `rm_writing` project, look at `example.c`
- it is missing much of the initialization, finish the initialization
- run it as:
`example -v`
- you will need a command line on your target to test it, try:
`echo Hello >/dev/example`
`ls /dev`
`ls -l /dev/example`
`cat /dev/example`
- you should see it printing that it is in the various io functions



Introduction to Resource Managers

Topics:

Overview

A Simple Resource Manager

- Initialization
- – Handling *read()* and *write()*

Conclusion

The example Resource Manager

Let's see what happens when a client uses the new `/dev/example` device (by, for example, doing `cat /dev/example`):

Internally, `cat` basically does:

```
fd = open("/dev/example", O_RDONLY) ;  
while (read (fd, buf, BUFSIZ) > 0)  
    /* write buf to stdout */  
close (fd) ;
```

The example Resource Manager

Which results in:

- Communications with the Process Manager:
 - an inquiry message to the process manager:
 - “who is responsible for `/dev/example`?”
 - returns a reply, “(pid, chid)” (our resource manager, **example**), “is responsible”
- Communications with **example**:
 - an open message
 - “open this device for read”
 - returns a reply, “yes, open succeeded, proceed”
 - the *open()* library call returns a file descriptor, **fd**
 - a read message
 - “get me some data”
 - returns a reply, “here are 0 bytes” (i.e. EOF)
 - a close message



The example Resource Manager's io

Let's look at example's I/O functions:

```
int  io_read (resmgr_context_t *ctp,  
              io_read_t *msg,  
              RESMGR_OCB_T *ocb) ;  
  
int  io_write (resmgr_context_t *ctp,  
               io_write_t *msg,  
               RESMGR_OCB_T *ocb) ;
```

- **msg** is always a pointer to the current message being handled

They both share the **ctp** and **ocb**...

I/O Functions Arguments -- ctp

ctp

- pointer to a resource-manager context structure
- information about the received message
- contains at least:

```
typedef struct _resmgr_context {  
    int                rcvid;  
    struct _msg_info    info;  
    resmgr_iomsgs_t    *msg;  
    size_t              msg_max_size;  
    size_t              offset;  
    size_t              size;  
} resmgr_context_t;
```




dispatch
context
ctp




I/O Functions Arguments -- ocb

ocb



- Open Control Block
- one **ocb** per *open()*
- maintains context between the *open()* call and subsequent I/O calls, i.e. *iofunc_open_default()* allocates and initializes it, I/O functions use it.
- will be either an **iofunc_ocb_t**, or
- an **iofunc_ocb_t** encapsulated within your own structure with additional state information, etc.
- points to the attribute structure for the device opened



READ

The message will be generated by the *read()* call which looks like:

```
bytes_read = read( fd, buf, nbytes )
{
    struct _io_read hdr;
    hdr.type = _IO_READ;
    hdr.xtype = _IO_XTYPE_NONE;
    hdr.nbytes = nbytes;

    ...

    return MsgSend( fd, &hdr, sizeof(hdr),
                    buf, nbytes );
}
```

READ

The message you'll receive is:

```
typedef union
{
    struct _io_read  i;
} io_read_t;

struct _io_read
{ // contains at least the following
    _Uint16t      type;    // message type = _IO_READ
    _Uint32t      xtype;   // extended type
    _Uint32t      nbytes;  // number of bytes to be read
};
```



For the reply:

- if successful:

```
MsgReply(ctp->rcvid, bytes_read, buffer,  
         bytes_read);  
return _RESMGR_NOREPLY;
```

- the reply message (buffer, bytes_read) would be your data (i.e. there is no header to worry about)
- the *MsgReply()* status would be the return value from the client's *read()*, the number of bytes successfully read

- if failed, do:

```
return errno_value;
```

The xtype (extended type) member

The xtype (extended-type) member:

- will most often be:

`_IO_XTYPE_NONE`

- most *read()* and *write()* handlers check for `_IO_XTYPE_NONE`. If xtype is not this, then they return `ENOSYS`



READ

example's *read()* handler:

```
int
io_read(resmgr_context_t *ctp, io_read_t *msg, RESMGR_OCB_T *ocb)
{
    int status;

    if ((status = iofunc_read_verify(ctp, msg, ocb, NULL)) != EOK)
        return status;

    if ((msg->i.xtype & _IO_XTYPE_MASK) != _IO_XTYPE_NONE)
        return ENOSYS;

    MsgReply(ctp->rcvid, 0, NULL, 0); /* 0 bytes successfully read */

    if (msg->i.nbytes > 0) /* mark access time for update */
        ocb->attr->flags |= IOFUNC_ATTR_ETIME;

    return _RESMGR_NOREPLY;
}
```



EXERCISE

Exercise:

- in your `rm_intro` project, go back to `example.c`
- modify the `read()` handler to return some data
 - make up some data
- run it as:
`example -v`
- test it from a command line on your target:
`cat /dev/example`



WRITE

The message will be generated by the *write()* call, which looks like:

```
bytes_written = write( fd, buf, nbytes )
{
    struct _io_write hdr;
    iov_t iov[2];
    hdr.type = _IO_WRITE;
    hdr.nbytes = nbytes;
    ...
    SETIOV(&iov[0], &hdr, sizeof(hdr));
    SETIOV(&iov[1], buf, nbytes );
    return MsgSendv( fd, iov, 2, NULL, 0 );
}
```

WRITE

The message you'll receive is:

```
typedef union
{
    struct _io_write i;
} io_write_t;
/* the data to be written usually follows the io_write_t */

struct _io_write
{ // contains at least the following
    unsigned short    type;    // message type = _IO_WRITE
    long              nbytes;  // number of bytes to write
    uint32_t          xtype;   // extended type
};
```



For the reply:

- if successful:

```
MsgReply(ctp->rcvid, bytes_written, NULL, 0);  
return _RESMGR_NOREPLY;
```

- there is no data to reply with
- the *MsgReply()* status would be the return value from the client's *write()*, the number of bytes successfully written

- if failed, do:

```
return errno_value;
```

WRITE

example's *write()* handler:

```
int io_write (resmgr_context_t *ctp, io_write_t *msg,
              RESMGR_OCB_T *ocb)
{
    int status;

    if ((status = iofunc_write_verify(ctp, msg, ocb, NULL)) != EOK)
        return status;

    if ((msg->i.xtype & _IO_XTYPE_MASK) != _IO_XTYPE_NONE)
        return ENOSYS;

    // msg -> i.nbytes is the number of bytes to be written,
    // we are telling it that we wrote everything (msg -> i.nbytes)
    MsgReply(ctp->rcvid, msg->i.nbytes, NULL, 0);

    if (msg->i.nbytes > 0) /* mark times for update */
        ocb->attr->flags |= IOFUNC_ATTR_MTIME | IOFUNC_ATTR_CTIME;

    return _RESMGR_NOREPLY;
}
```




WRITE - Getting the data


example doesn't do anything with the data to be written, but what if you want to?

- the data usually follows the `io_write_t` in the message buffer.
- but it may not all have been received, what happens in the following case?

```
MsgSend(coid, smsg, sbytes, ...);
```

`smsg =` 
└──────────────────┘
`sbytes = 3000`

```
MsgReceive(chid, rmsg, rbytes, ...);
```

`rmsg =` 
└──────────┘
`rbytes = 1000`

As you know, the kernel copies the lesser of the two sizes, so in this case only 1000 bytes will have been received. How do you handle this?

WRITE – Getting the data

You need to:

- figure out where to put the data
 - pre-allocated ring buffer
 - hardware cache buffers
 - allocate a buffer
- put the data where it goes:
 - for a single destination buffer:
 - *resmgr_msgget()*
 - for multiple destination buffers:
 - *resmgr_msggetv()*

Getting the data – resmgr_msgget() details

Get client data with:

```
bytes_copied =  
    resmgr_msgget(ctp, buf, bytes, offset);
```

- copy **bytes** of data from client **ctp->rcvid**
- starting at **offset** bytes after the start of **msg**
 - remember, **msg** is a standard parameter to our I/O functions
- to **buf**
- copies locally if possible
- calls *MsgRead()* when needed

EXERCISE

Exercise:

- we're going to further modify `example.c` in your `rm_intro` project
- modify the `io_write` function to:
 - print out the number of bytes written
 - print out all the data written
- run it as:

```
example -v
```

- test it from a command line on your target:

```
echo Hello >/dev/example  
cp /etc/services /dev/example  
cp /etc/termcap /dev/example
```



Introduction to Resource Managers

Topics:

Overview

A Simple Resource Manager

- Initialization
- Handling *read()* and *write()*

→ Conclusion

Conclusion

You learned:

- that a resource manager is a device driver framework
- how to initialize and register a resource manager
- how to handle *read()* and *write()* client requests