

System Profiling

You will learn:

- how detailed data can be gathered from the kernel, about many things that are going on, all the way down to the interrupt and kernel call level
- how to add your own data
- how to control this data gathering from:
 - the IDE
 - the command line
 - your own code
- how to analyze this data

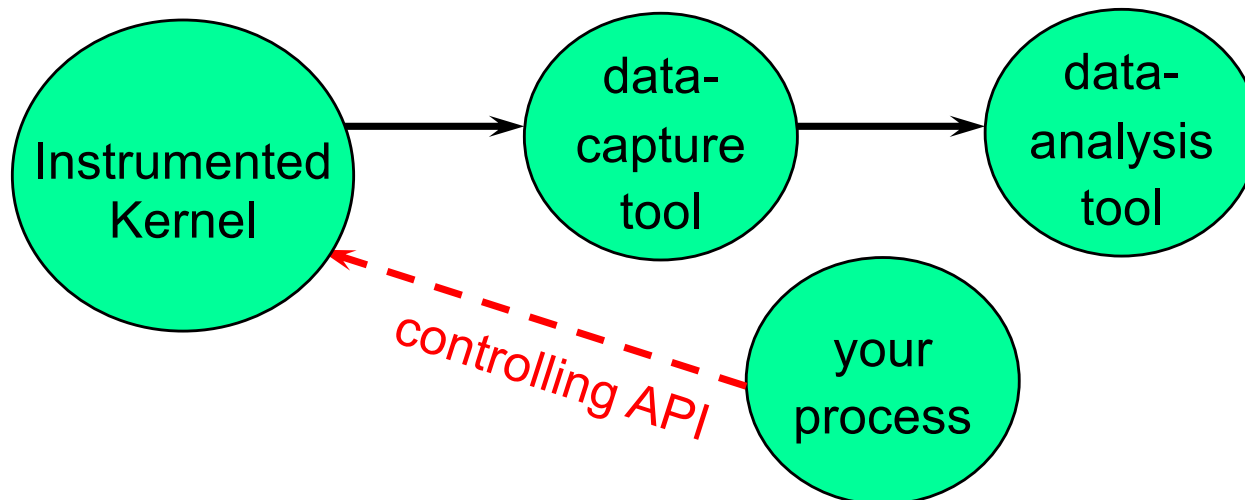
System Profiling

Topics:

- ➔ **Overview**
- Creating a Log**
- Log Summary**
- A Quick Tour**
- Filtering Events and Event Owners**
- Navigating through a log with the Timeline Pane**
- Statistics**
- CPU Activity Pane and CPU Usage Pane**
- Tying The Trace to Your Code**
- Multi-core Related Features**
- Adaptive Partitioning: Partition Summary**
- Conclusion**

System profiling consists of:

- an *Instrumented Kernel* that logs many different types of *events*, as they happen
- tools for capturing and analyzing that log
- an optional API for controlling logging



Examples of event types:

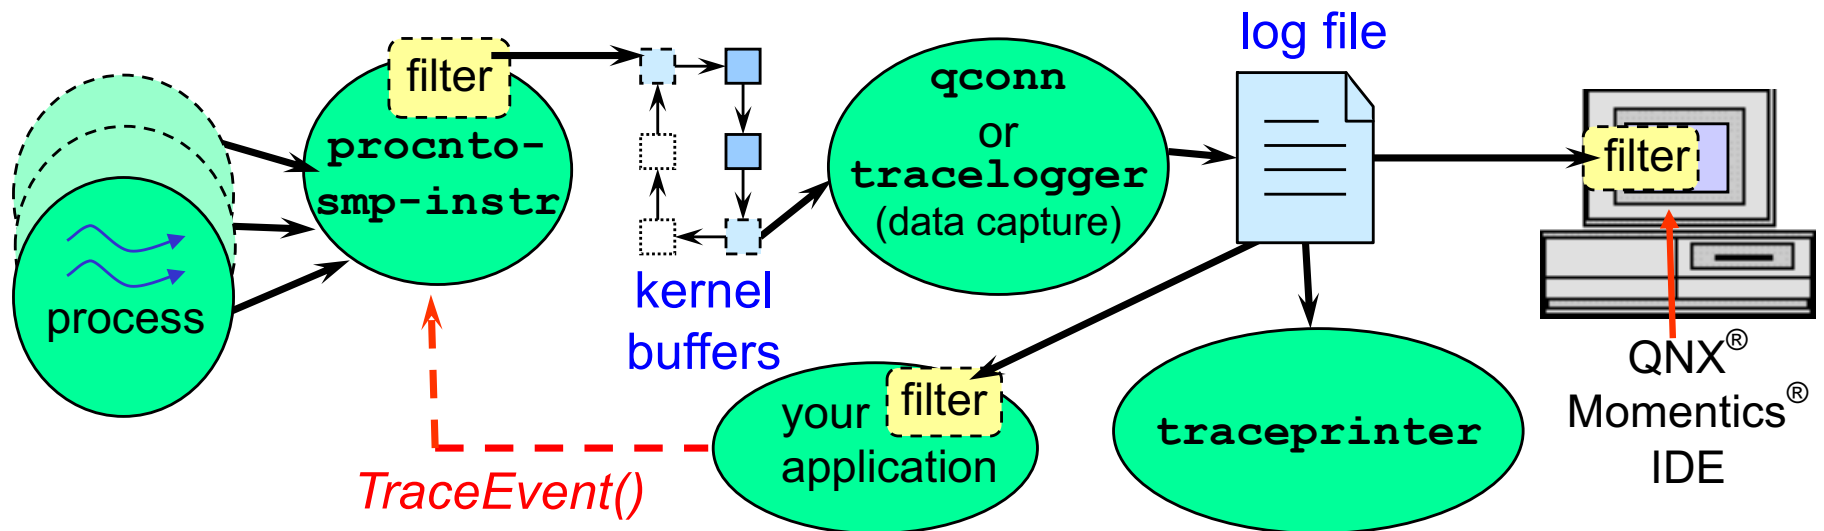
- kernel calls
- process manager activity (e.g. process creation)
- interrupts
- rescheduling (thread state changes)
- context switches
- user-defined trace events

What system profiling does for you:

- gives you a way to analyze how the different processes and/or threads in your system interact
 - this goes beyond traditional debugging, which gives you a process level view, not a system level one
- gathers data for post-mortem analysis
- you can choose:
 - what types of events are logged
 - where the data is stored

Overview

Information logging involves several aspects:



- ways to control logging:
 - **qconn**
 - **tracelogger**
 - custom application (using **TraceEvent()**)
- choices for analyzing log:
 - Momentics® IDE System Profiler perspective
 - **traceprinter**
 - custom application (using traceparser API)

The events:

- an event will include the following:
 - what category the event belongs to (event class):
 - kernel call, interrupt servicing, process and thread management, user-generated
 - when it happened
 - event-specific data
- events can be fast or wide

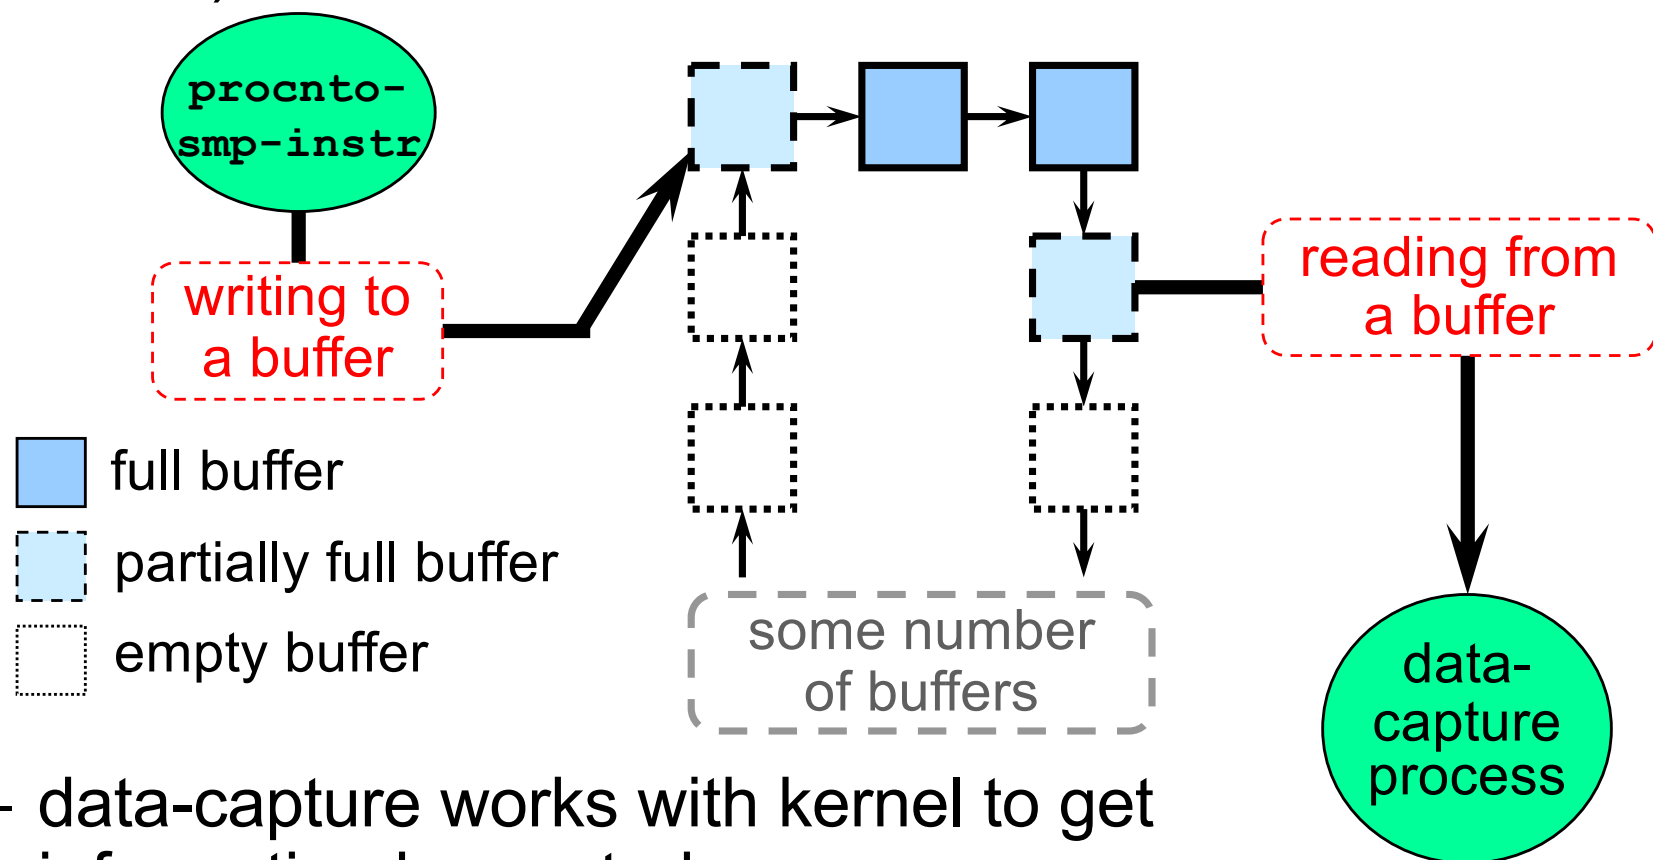
	Amount of information	Work for kernel	Size of log files
Fast events	less	less	smaller
Wide events	more	more	larger



Overview - The instrumented kernel

Kernel buffer management:

- events are stored in a circular list of buffers (a ring buffer)




- data-capture works with kernel to get information harvested

continued... ↓

Overview - The instrumented kernel


Kernel buffer management (continued):

- buffer size is fixed




- the number of buffers that can be successfully created may vary over time

☞ the kernel allocates buffers from physically contiguous memory, if there isn't enough for the requested number of buffers then it will simply use less



The data-capture process:

- retrieves events logged in the kernel buffers
 - **qconn** is used to give data to the IDE
 - **tracelogger** is run for automated logging, or logging under program control
 - you can create your own data-capture process
- can do these three things:
 1. tell the kernel what sorts of events to log
 2. interface with the kernel
 3. move events from the kernel's event buffers into some other storage location

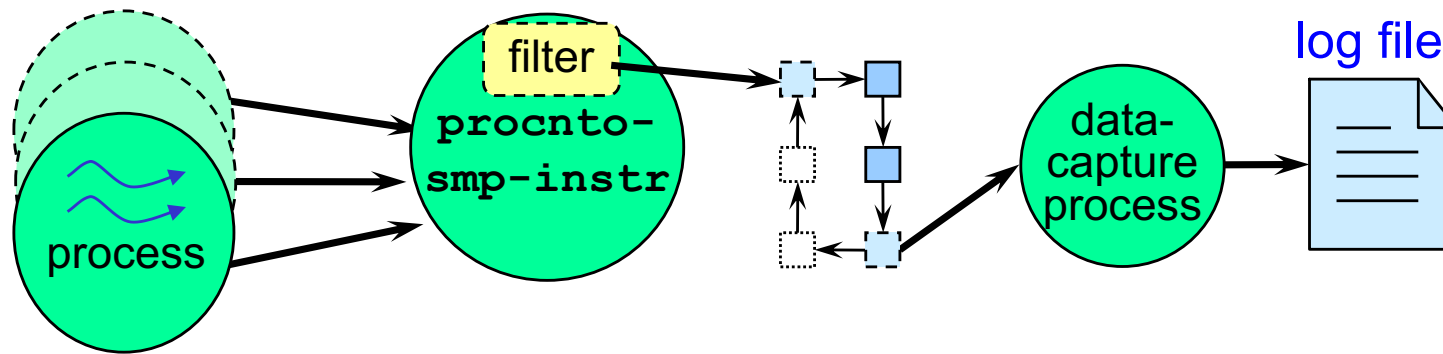


Overview - Data capture

Choosing which events to log:

- as you work on your analysis, you may find that only some data is helpful
 - in many cases, only some types of events are relevant
 - you might need only a subset of the event information
- you can make **procnto-smp-instr** filter out data
 - allows you to capture logs of longer time duration
 - improves system performance during capture
 - easier to digest the info (both for you and the IDE).

👉 no record of this data will be available



System Profiling

Topics:

Overview

→ Creating a Log:

- From the IDE
- Exercise
- Using tracelogger
- Under program control

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

...

Ways to create a log:

- from the IDE's Target Navigator view
- from the command-line using **tracelogger**
- using **tracelogger**, but under program control
- if you want to use the System Profiler to analyze the log, it has to go in an IDE Project
 - doesn't matter how the log was created, e.g. System Profiler can analyze log from **tracelogger**
 - a project is just a directory, so log can go into any one
 - some suggestions:
 - your QNX Target project
 - a project relevant to your problem

Topics:

Overview

Creating a Log:



- From the IDE
- Exercise
- Using tracelogger
- Under program control

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

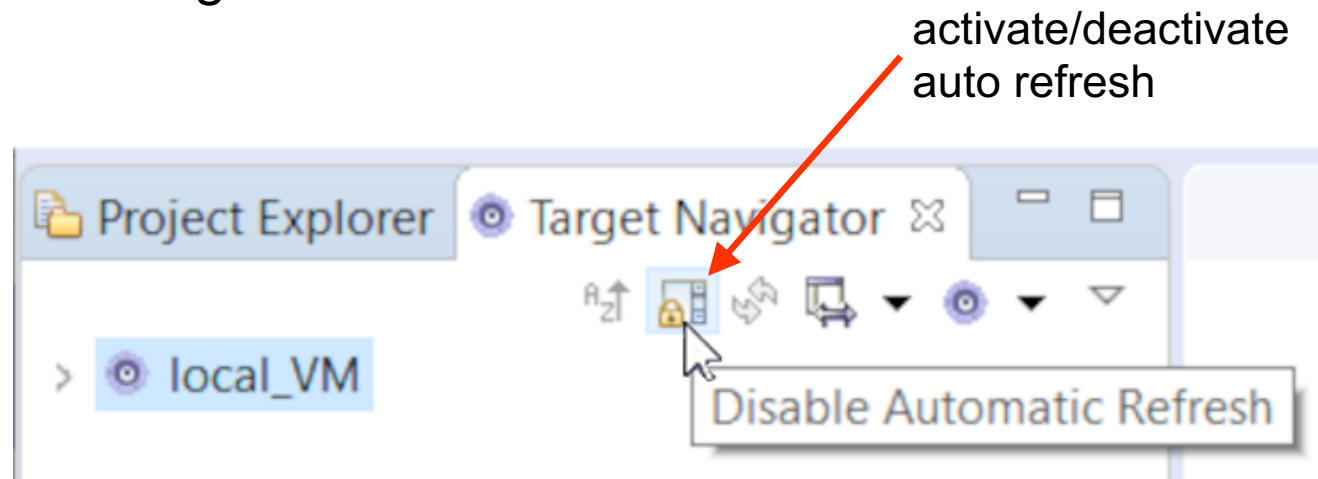
Tying The Trace to Your Code

...

Creating a Log from the IDE

Consider eliminating things that could interfere:

- the IDE periodically gathers system information
 - this generates irrelevant events

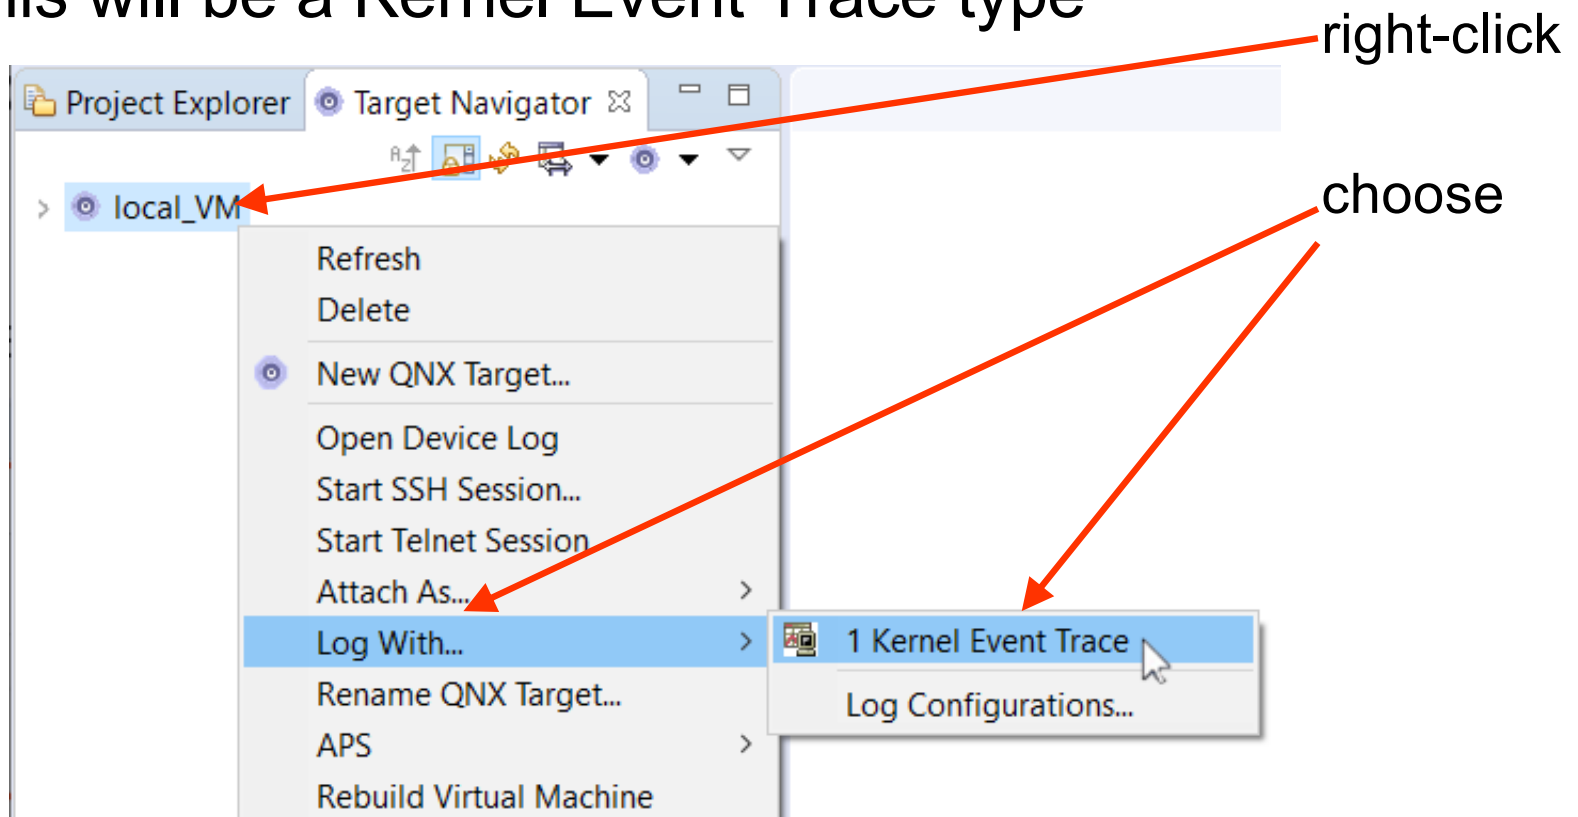


- consider deactivating Automatic Refresh (at least during capture):

Creating a Log from the IDE

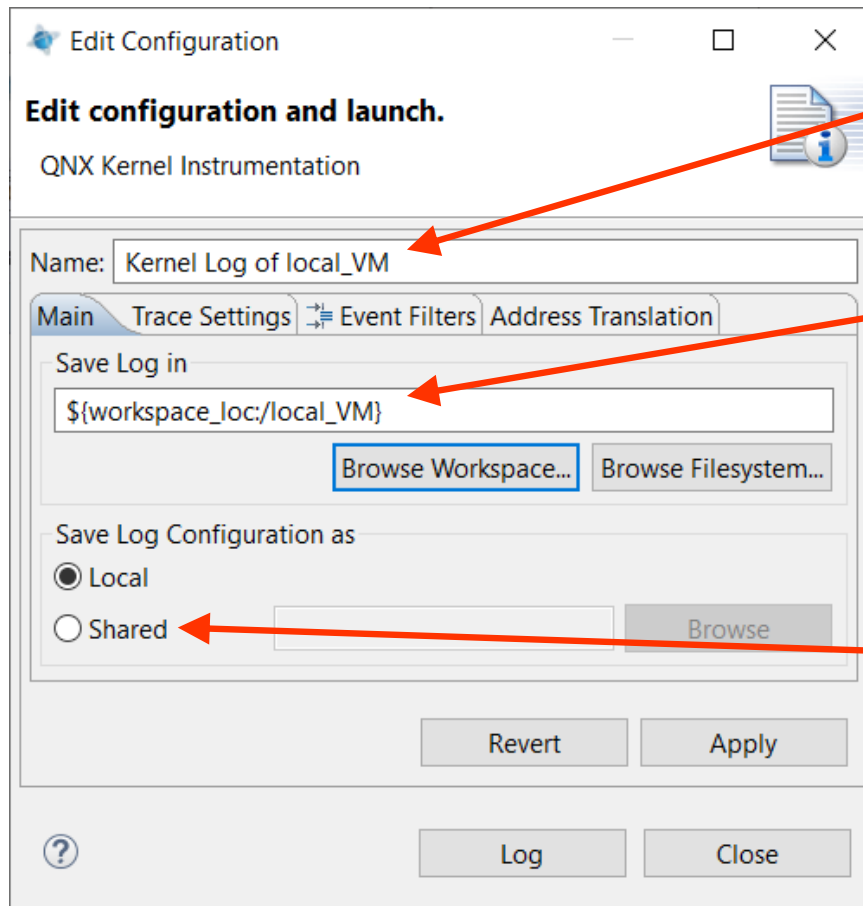
Next, create a Launch Configuration:

- the IDE uses launch configurations to remember logging settings as well as running programs
- this will be a Kernel Event Trace type



Creating a Log from the IDE

Log Configuration settings (Main tab):



give the configuration a name

where to save log file

you can make it easier to give the log configuration to someone else by making it a shared configuration

Creating a Log from the IDE

Log Configuration settings (Trace Settings tab):

The screenshot shows the 'Trace Settings' dialog box with the following configuration:

- Tracing duration:** ☒ Period of time (in seconds): 3.0; ☐ Number of iterations: 32
- Trace Collection:** ☒ Save on target then upload; File path on target: /dev/shmem/tracebuffer.kev; ☐ Memory mapped file then upload; Filename on target: tracebuffer.kev; Max file size: 41943040B
- Buffers:** Number of kernel buffers: 32; Number of qconn buffers: 128

choose whether log capture completes after:

- specified time is reached, or
- specified amount of data is captured

choose where the log is stored:

- file on target, or
- Memory mapped file -- useful if no storage on target

the default buffer settings are usually fine

Creating a Log from the IDE

Log Configuration settings (Event Filters tab):

– works as a hierarchy of choices


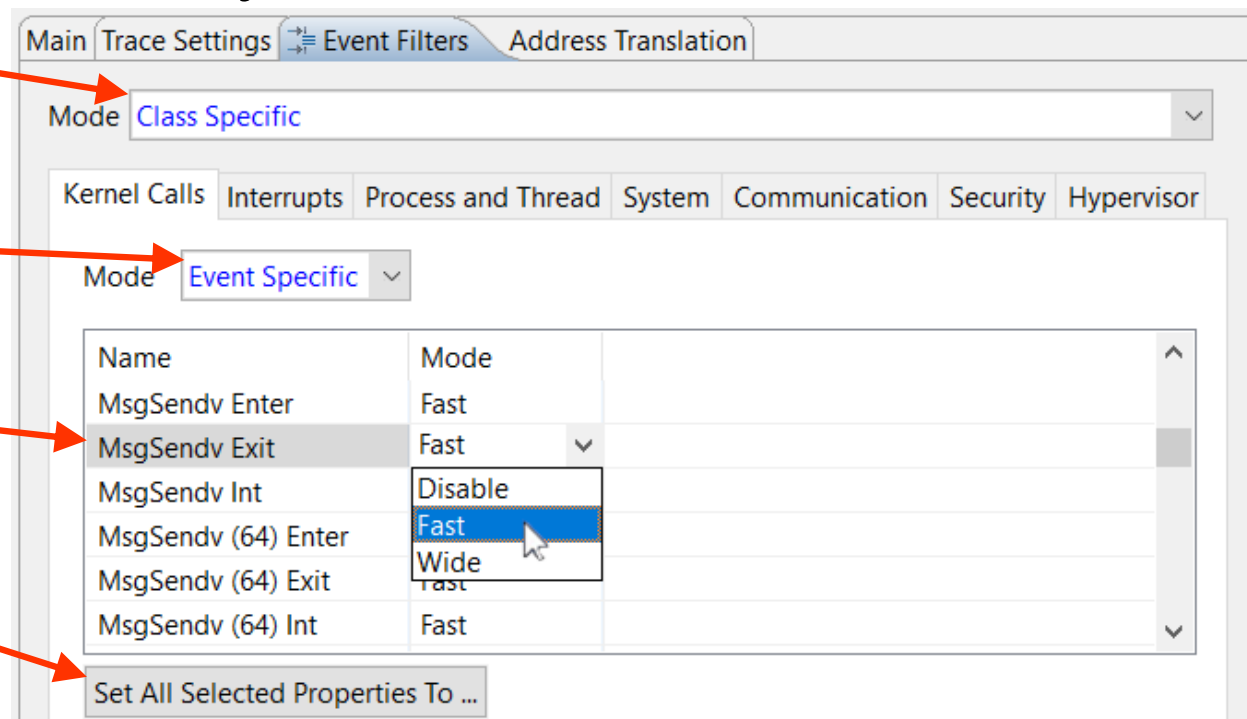
Class Specific:
enables setting
modes within the tabs

Event Specific:
enables setting modes
for events within this class

you can configure
specific events here,
but only if Event Specific
is set above

Or, after multi-select, you
can set all selected at
once

this filtering affects
this stage
of the logging

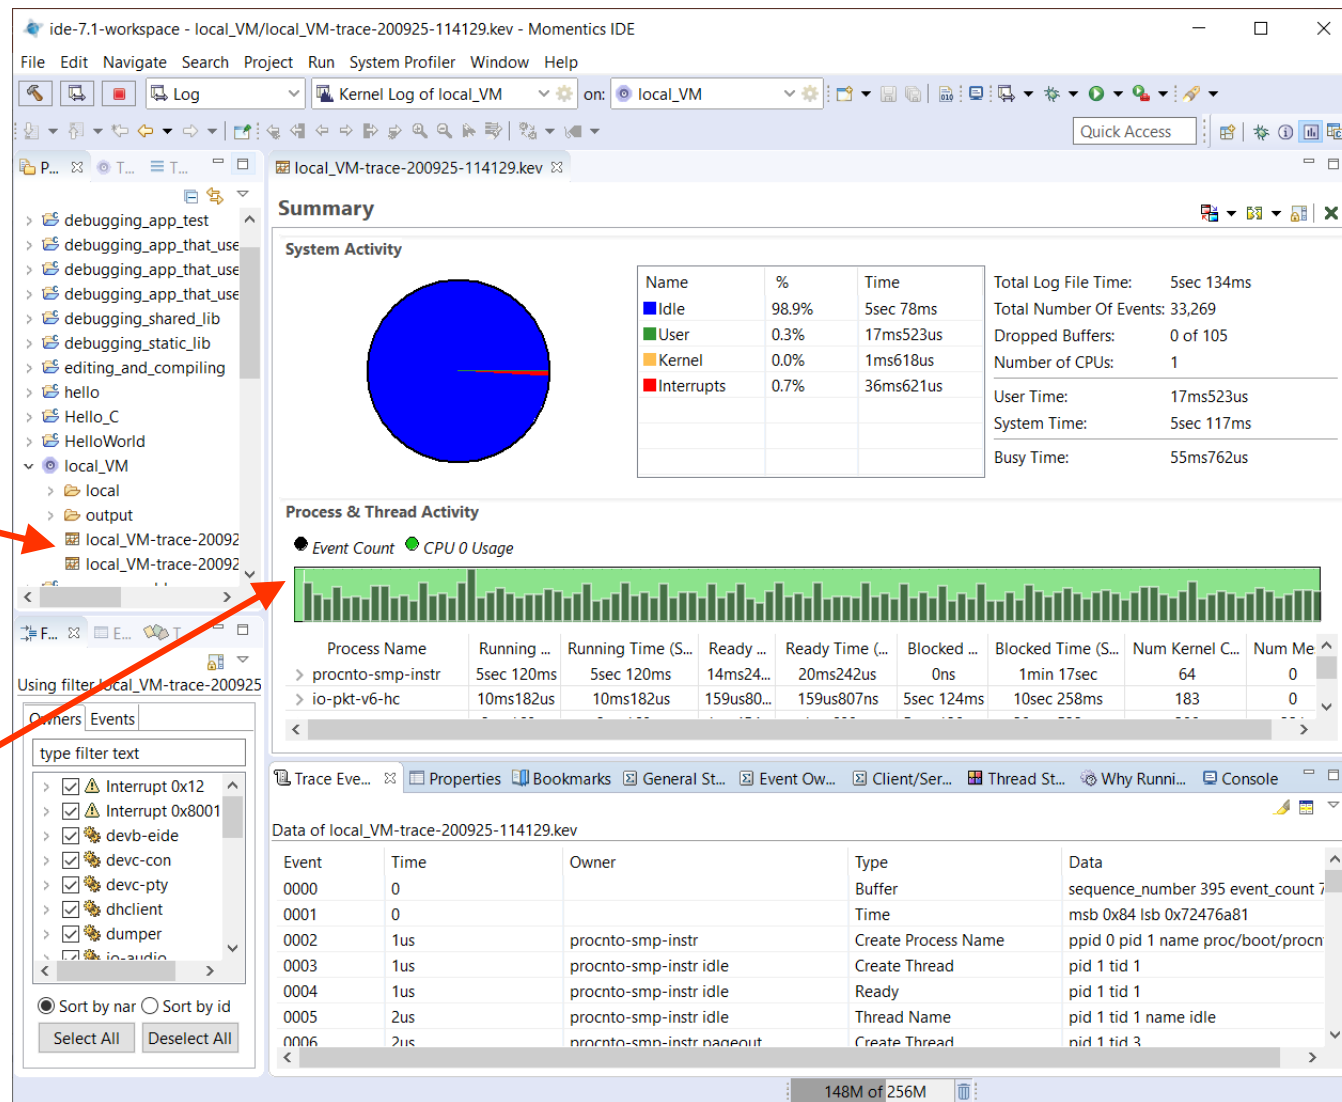


Creating a Log from the IDE

The resulting log file displayed in the IDE:

resulting log file is in the Target System project

double-click on it to open...
...Summary is displayed to give you an overview of the log



Topics:

Overview

Creating a Log:

- From the IDE
- – Exercise
- Using tracelogger
- Under program control

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

...

EXERCISE - Creating a Log from the IDE

Experiment with log creation:

1. capture a 1 second log

- pay particular attention to the configuration of the log capture

2. open it in the system profiler

- don't worry too much about what you are looking at, we'll get into that later

3. capture another 1 second log

- this time, use filtering to tell the kernel not to log interrupt or kernel call event classes
- compare the log sizes (without and with filtering)
 - (right-click on log file, choose properties)

Topics:

Overview

Creating a Log:

- From the IDE
- Exercise
- – Using tracelogger
- Under program control

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

...


Creating a Log using tracelogger

tracelogger is useful if:

- you don't have the IDE handy
 - e.g. to capture what happened during system initialization
- you want to create the log under program control

Working with **tracelogger** you can:

- control how, where and when data is logged
- set **procnto-smp-instr**'s event filter



tracelogger can be run in 5 different modes:

- n iterations** : (the default), log for a specified number of kernel buffers (default 32)
- s seconds** : log for the specified number of seconds
- c** : continuous, log until tracelogger is terminated (SIGINT preferred)
- d1** : daemon mode, run in background waiting for program to control things
- r** : ring mode, generate data, but don't dump until later

We'll look at daemon and ring mode in more detail.

Creating a Log using tracelogger


Running **tracelogger**:

– example command-line:

```
tracelogger -f /dev/shmem/mylog.kev -s3
```



file to log to



how many sec.
to log for

Topics:

Overview

Creating a Log:

- From the IDE
- Exercise
- Using tracelogger
- – Under program control

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

...

Program Control of Log Creation

tracelogger can do logging when told to by a program (daemon mode):


```
tracelogger -d1 -E -f logfile.kev
```

- can configure basic filtering, fast/wide, etc on the command line
 - the **-E** option allows this
- detailed configuration and control can be done using *TraceEvent()* calls
- event logging begins only after the appropriate *TraceEvent()* call is made

Program Control of Log Creation

TraceEvent():

- manages the flow of data to **tracelogger**



- lets you control:

- when to start/stop logging
- what events to generate/filter
- how many data buffers **procnto-smp-instr** is asked to use
- fast vs. wide events
- where data is sent (**/dev/shmem/tracebuffer** by default)

- can also be used to insert your own events



Controlling the logging

- start logging by the kernel

```
TraceEvent (_NTO_TRACE_START) ;
```

- stop logging:

```
TraceEvent (_NTO_TRACE_STOP) ;
```

Ring mode vs. linear mode


- **-r** option sets ring mode; default is linear
- ring mode
 - kernel continuously goes around its ring, storing events
 - events only read by **tracelogger** on demand, when:
 - a program does **TraceEvent(_NTO_TRACE_STOP)** ;
 - **tracelogger** receives SIGINT
 - good for situations where you don't know in advance when you might need to analyze events, e.g. SIGSEGV
 - when SIGSEGV happens, have **tracelogger** flush the buffers, then look back in the log
 - you can spawn **tracelogger**, then do **kill(tracelogger_PID, SIGINT)** to stop logging

continued...

Ring mode vs. linear mode (continued)

–linear mode

- kernel continuously goes around its ring storing events, but when each buffer becomes full, it notifies **tracelogger** to read the data
- good for situations where you know something interesting will happen at some point in the future, and you want to start **tracelogger** storing events

 the kernel always stores events in its ring of buffers, regardless of the linear vs. ring option

For more information on creating a log under program control:

- have a look at the `TraceEvent()` function documentation in the Library Reference documentation
- as well as the System Analysis Toolkit User's Guide

Topics:

Overview

Creating a Log

→ Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

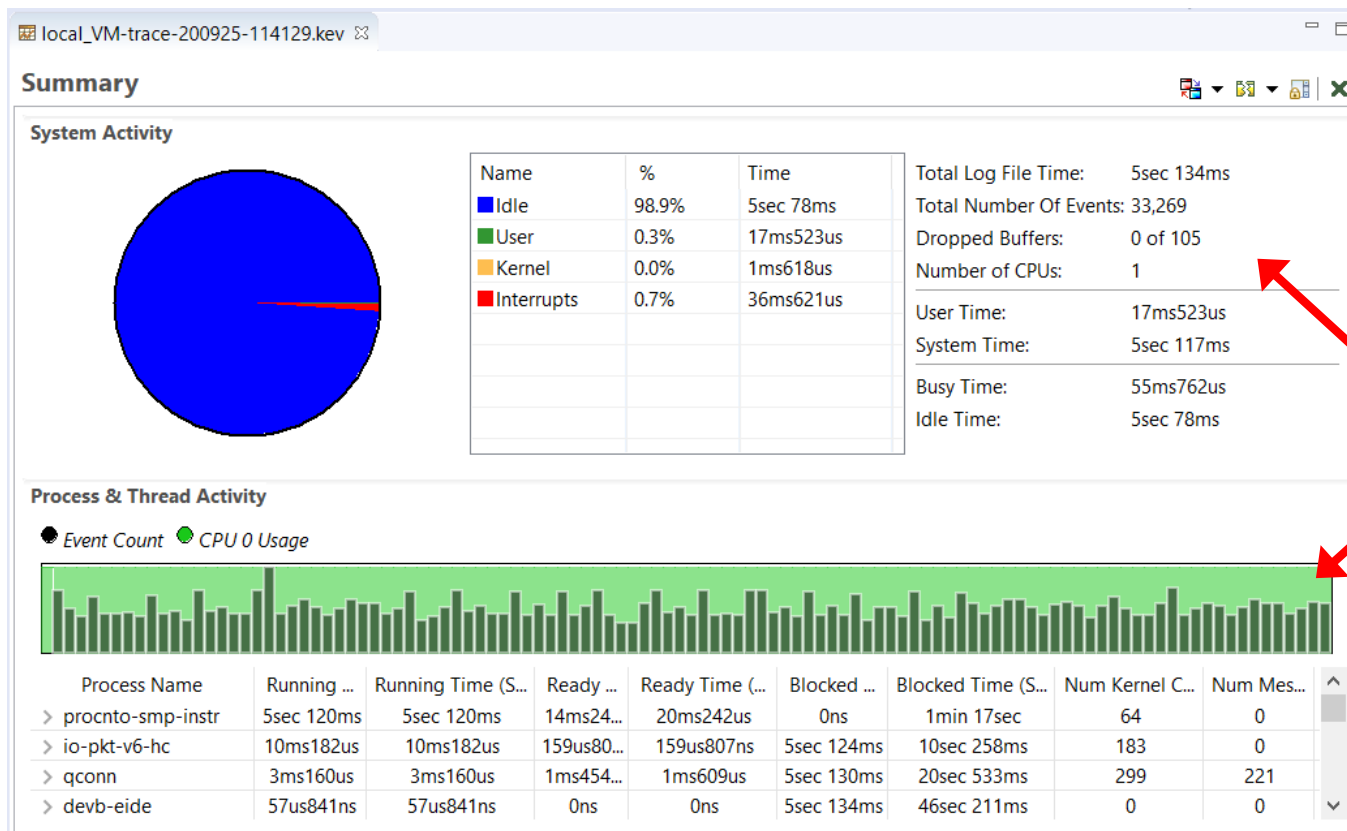
Multi-core Related Features

Adaptive Partitioning: Partition Summary

Conclusion

Event Log Summary

After opening a log file, first thing you see is summary information:

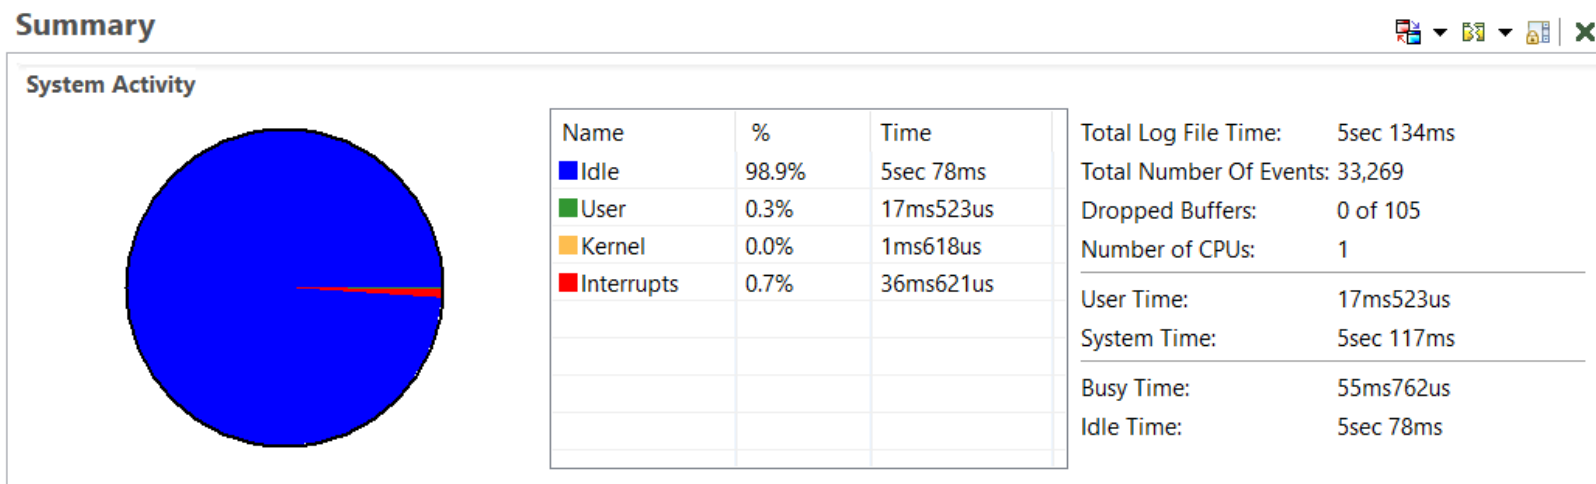


let's look at what these indicate...

Event Log Summary

Summary: System Activity

- indicates how busy the system was during the log



- Idle is the time that the idle thread (or threads) was (were) executing
- User is the time spent in threads in processes
- System is time spent in the kernel
- Interrupts is time spent handling interrupts
 - this should generally be fairly low; a high load could indicate:
 - faulty hardware
 - bad driver
 - application profiling for too many applications




Event Log Summary

Summary: Process & Thread Activity

- shows change in CPU usage and event rate

Process & Thread Activity

● Event Count ● CPU 0 Usage



Process Name	Running ...	Running Time (S...	Ready ...	Ready Time (...)	Blocked ...	Blocked Time (S...	Num Kernel C...	Num Mes...
> procnto-smp-instr	5sec 120ms	5sec 120ms	14ms24...	20ms242us	0ns	1min 17sec	64	0
> io-pkt-v6-hc	10ms182us	10ms182us	159us80...	159us807ns	5sec 124ms	10sec 258ms	183	0
> qconn	3ms160us	3ms160us	1ms454...	1ms609us	5sec 130ms	20sec 533ms	299	221
> devb-eide	57us841ns	57us841ns	0ns	0ns	5sec 134ms	46sec 211ms	0	0

- table provides metrics on how busy each individual process/thread was

Topics:

Overview

Creating a Log

Log Summary

→ A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

Multi-core Related Features

Adaptive Partitioning: Partition Summary

Conclusion

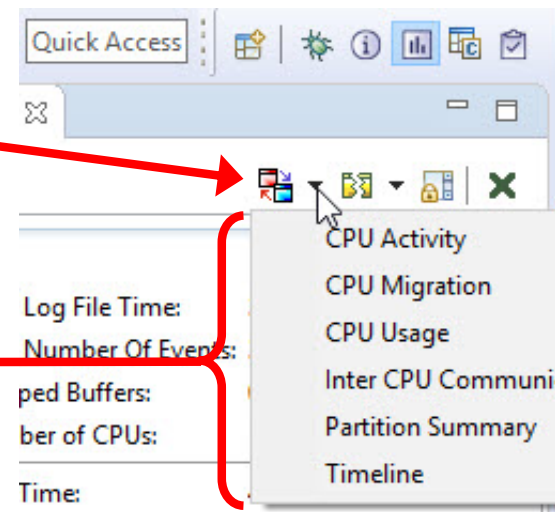
Switching Panes

Switching Panes:

- you can switch to different ‘Panels’ in the System Profiler that give you different methods for visualizing trace data

click this icon to switch the display

let's go through the different displays...



The different display types:

- **Summary:** gives an overall picture of the log
- **CPU Activity:** tracks total CPU usage over time
- **CPU Migration:** provides information about how the kernel migrated threads between CPUs in a multi-core or multi-CPU system
- **CPU Usage:** shows the average % CPU usage, per process (or thread), for the time period in the display
- **Inter-CPU Communication:** indicates how many times the kernel migrated threads due to message passing
- **Partition Summary:** shows how much CPU usage each scheduling partition consumed
- **Timeline:** graphically shows timing of events

We'll also be looking at various views



System Profiling

Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

→ Filtering Events and Event Owners

– Exercise

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

Multi-core Related Features

Adaptive Partitioning: Partition Summary

Conclusion

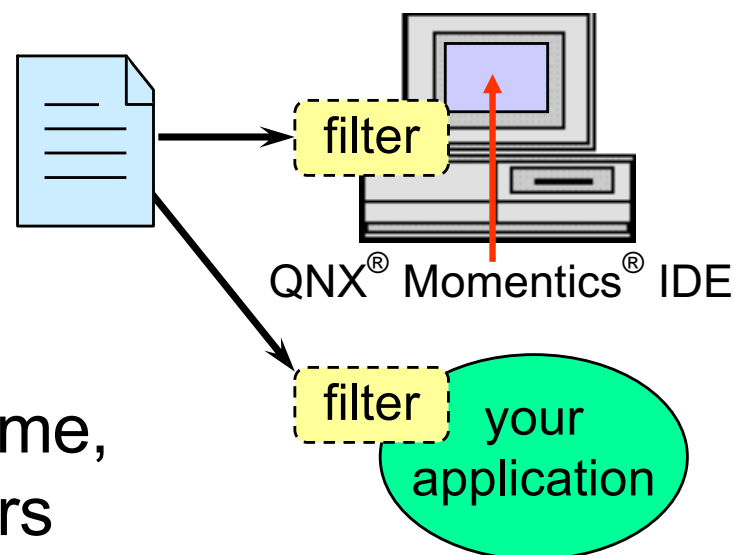
Filtering Events

You can filter events during data-analysis:

- this is useful when you want to look at a subset of your captured data

- the IDE allows you to filter out events based on type, time, and various other parameters
- the *traceparser*()* library allows similar control for applications you write yourself


👉 events filtered out at this stage are not lost, unless you save the filtered log file



Filtering Events

Kernel filter vs. System Profiler filter:

the filtering we're talking about now affects this stage of the logging



Event Owner Filter vs. Event Filter:

– Event Owner Filter

- filters in or out events based on the owner of the event
- event owners are:
 - processes
 - threads
 - interrupts
- e.g. show only events for `procnto` and `HelloWorld`, or remove events for `procnto`, `qconn` and `io-pkt`

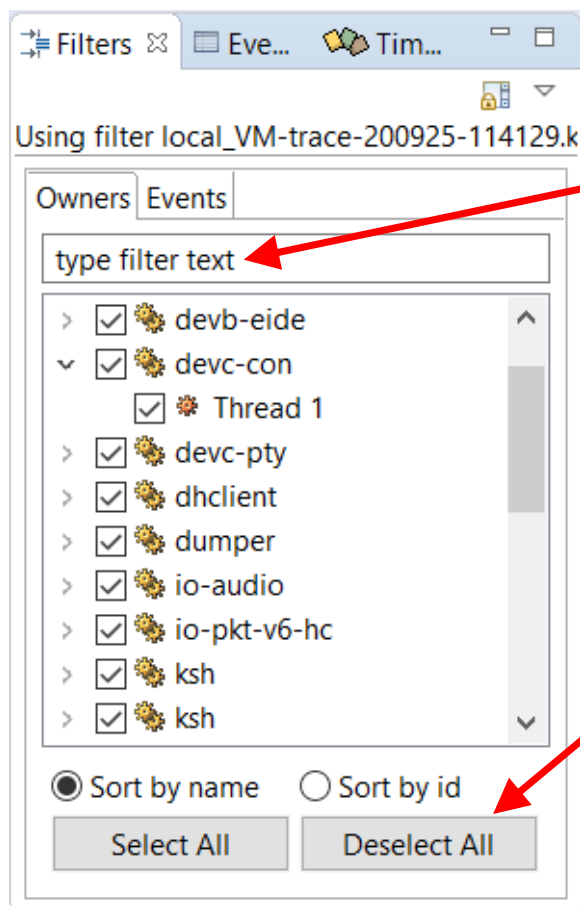
– Event Filter

- filters in or out certain types of events
- e.g. show only message-passing events

Filtering Events

Owner Filters:

- to remove event owners you don't care about uncheck in 'Owners' tab



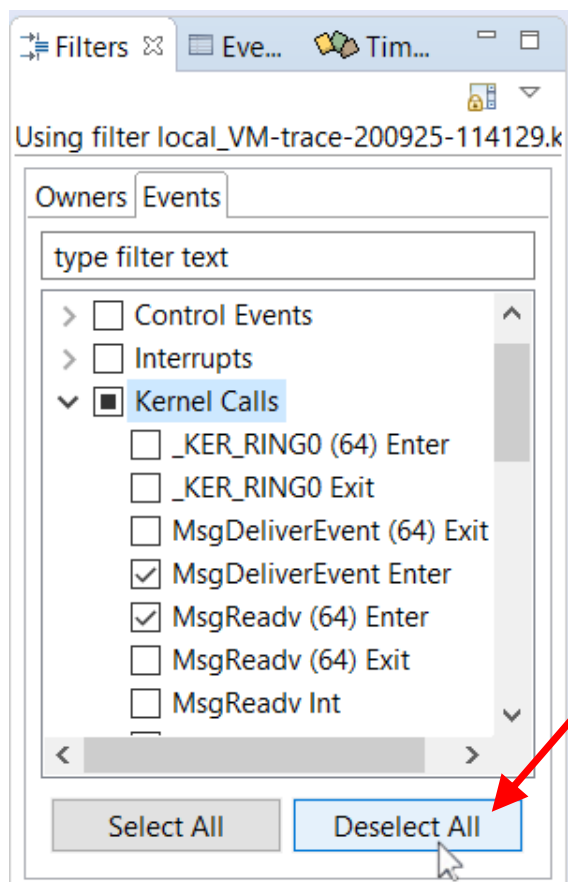
in a large list of processes, you can type some text to find the processes you want

if you just want to see a few owners, remove everything then add back in

Filtering Events

Event Filters:

- select/select event classes or specific event types you want or don't want

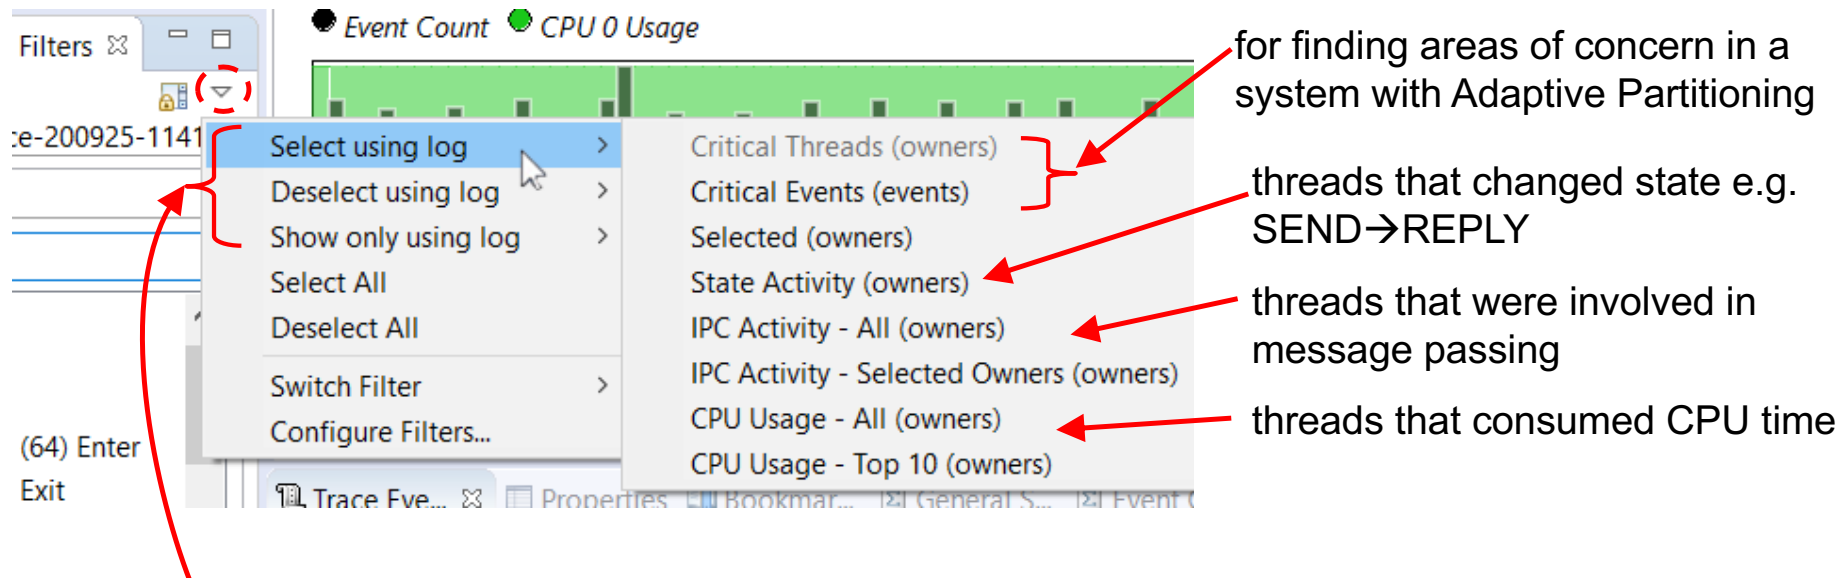


it's useful to be able to deselect all events and start from nothing, adding in events of interest

Filtering Events

Built-In Filters:

- a collection of useful pre-existing filters



- how to apply the filter:
 - Select using – add data that matches filter to data set
 - Deselect using – remove data that matches the filter from the data set
 - Show only – only include data that matches the filter in the data set

Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners



– Exercise

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

Multi-core Related Features

Adaptive Partitioning: Partition Summary

Conclusion

EXERCISE – Filtering

Try out the various filters.

Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

→ Navigating through a log with the Timeline Pane

– Exercise

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

Multi-core Related Features

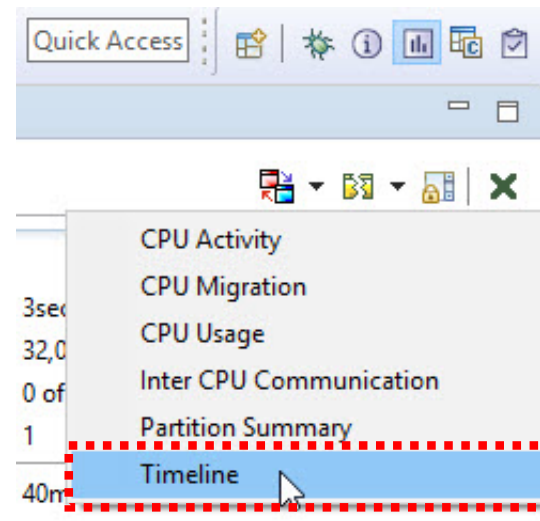
Adaptive Partitioning: Partition Summary

Conclusion

The Timeline

The Timeline pane is probably the best visualization tool:


– Let's look at it...



The Timeline

The timeline graphically shows the timing of:

- QNX native message passing
- thread states
- events that occurred, .e.g.
 - interrupts
 - entry into all kernel calls made
 - many others
- events are represented by vertical ‘ticks’



Timeline - Zooming

Zooming in/out:

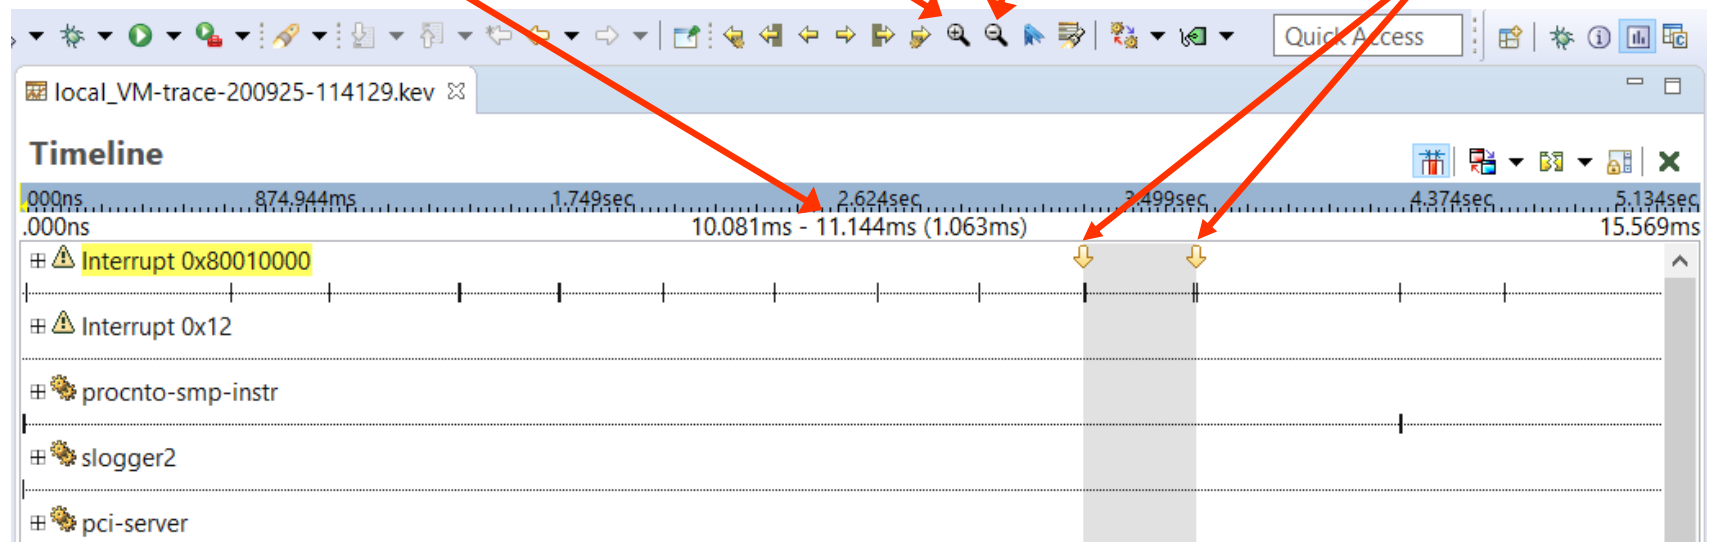
time range of
your selection
is shown here

clicking on
the + magnifying
glass zooms in

the -
zooms out

to zoom in,
select the range
you want to magnify

it automatically
snaps
to the nearest event

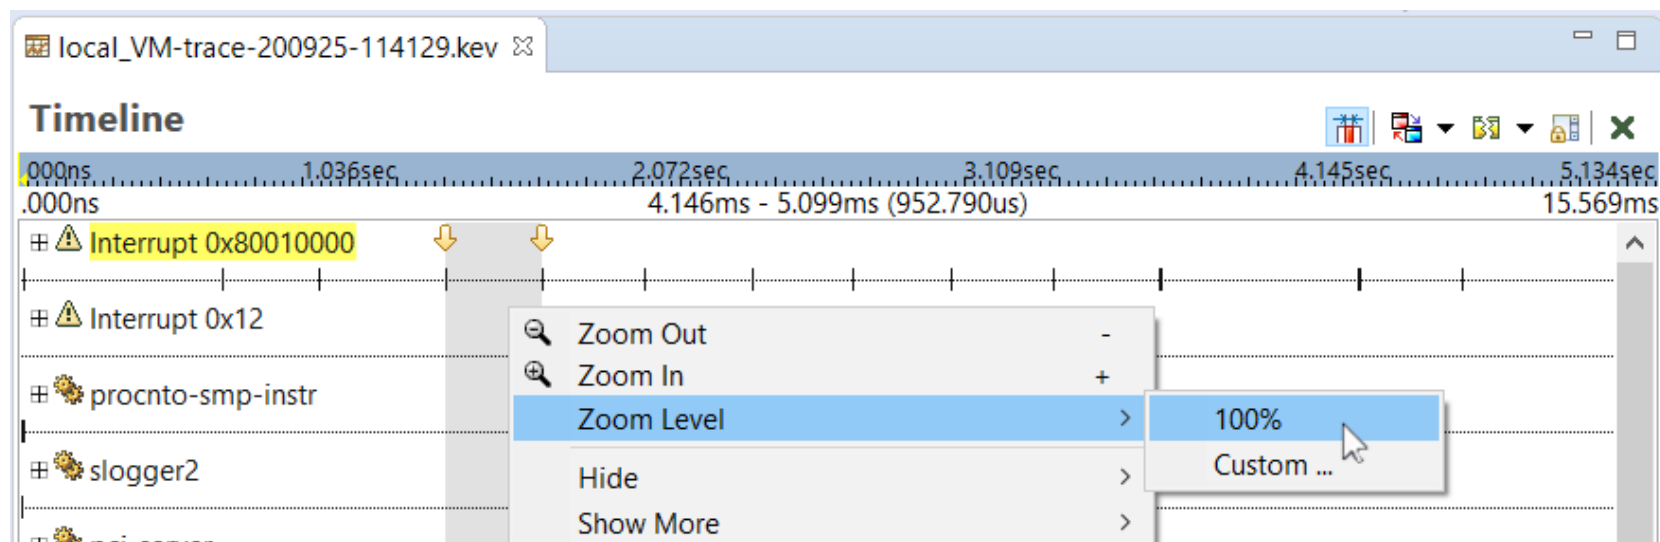


continued...

Timeline – Fixed Zoom

Zooming in/out (continued):

- for other zooming choices, right-click, choose Zoom Level:

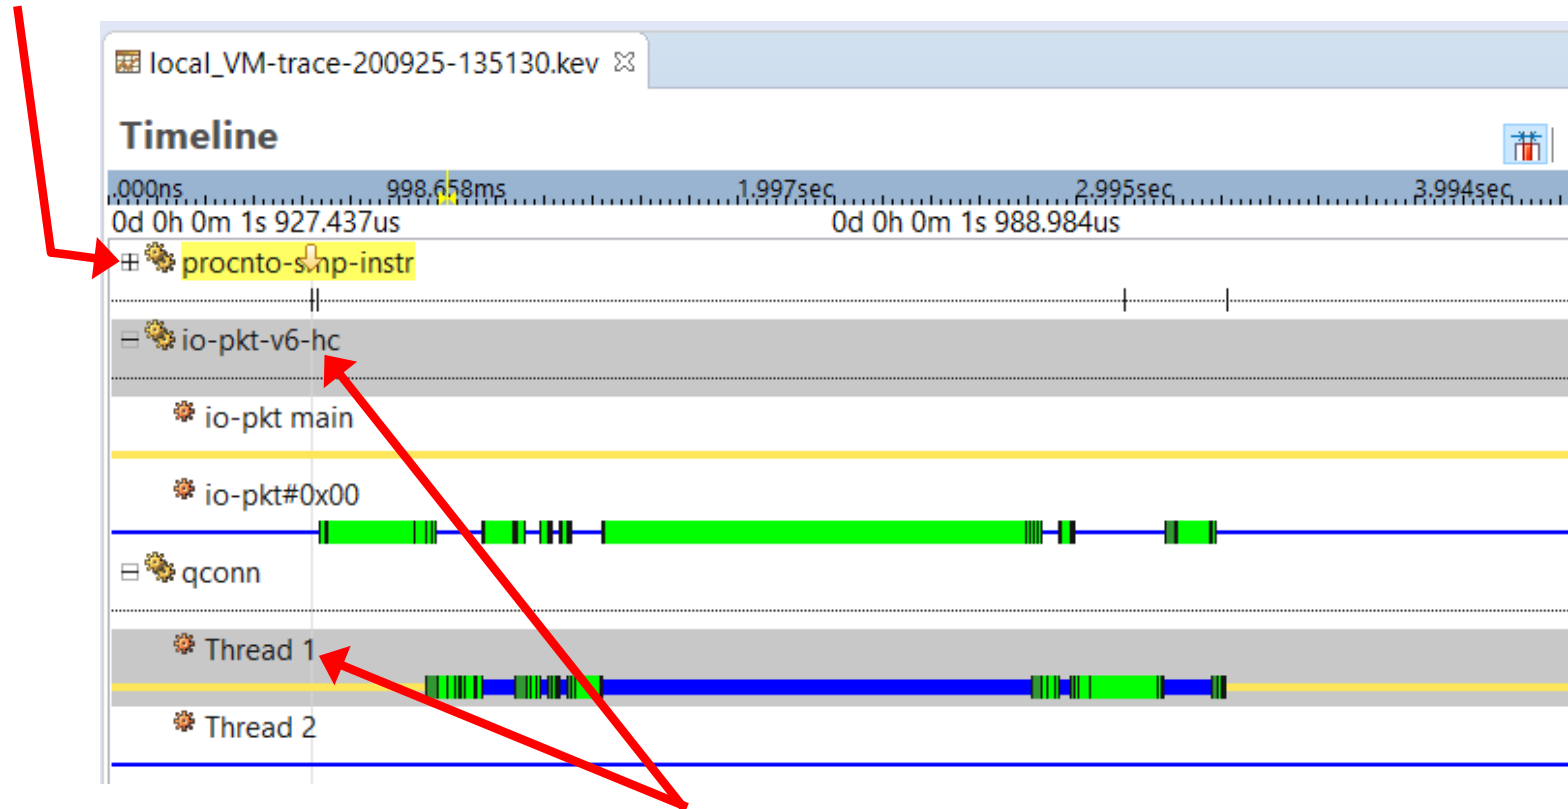


- you don't need to select a range to use these

Timeline – Viewing Threads

Viewing thread states:

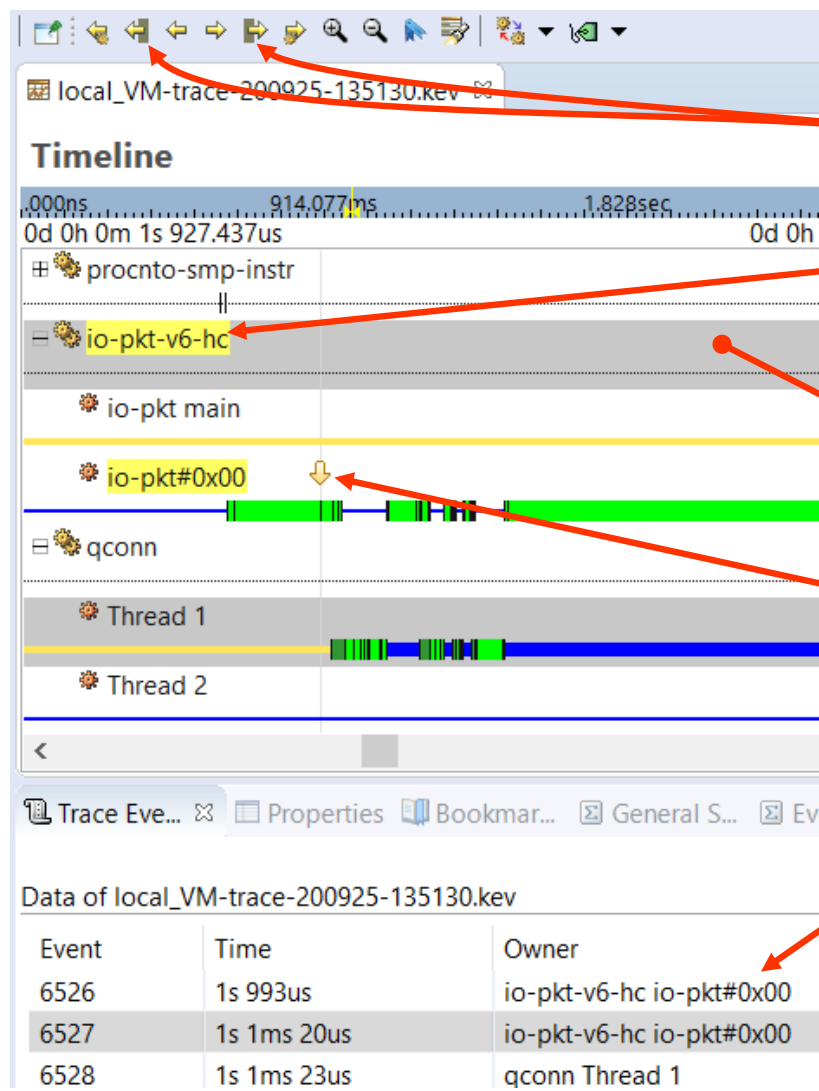
- click the '+' to show threads:



- click a process or thread to select it (grey highlight) for performing a specific operation on it (**Shift** and/or **Ctrl** for multiple selections), deselect by **Ctrl**-clicking on one

Timeline – Getting Around

Getting around:



click on these buttons to move between events in the selected owner(s)

click in here and use ctrl-left and ctrl-right arrow keys to move between events

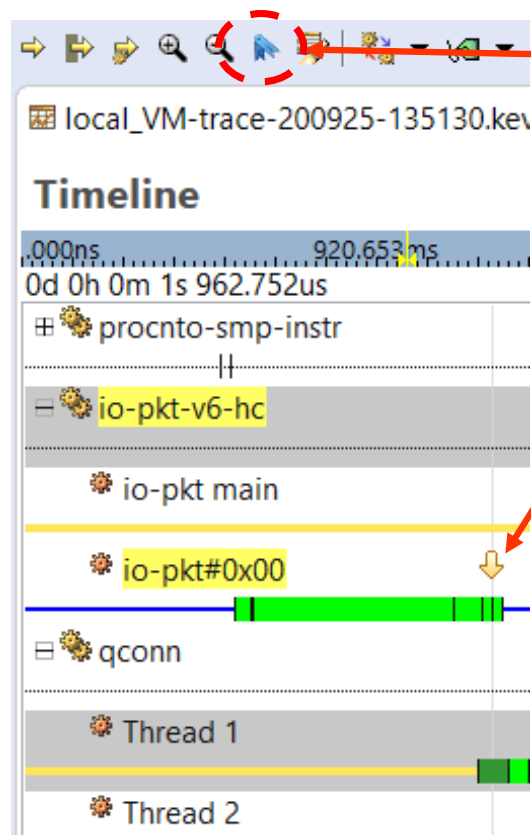
arrow points to currently selected event

walk through the events in the Trace Event Log view using the up & down arrow keys

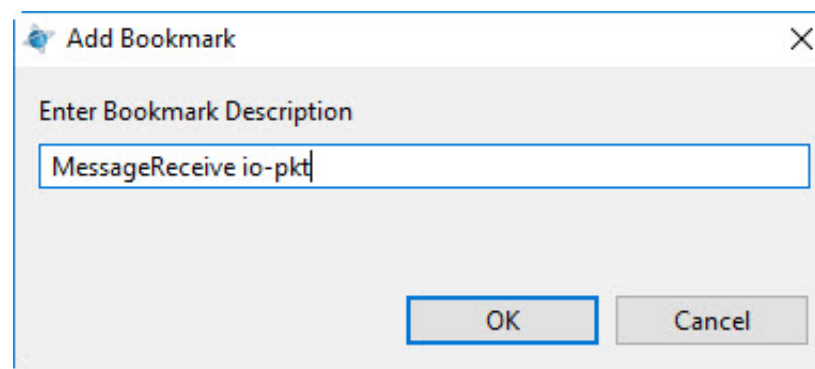
Bookmarks

It is easy to lose your place in the log:

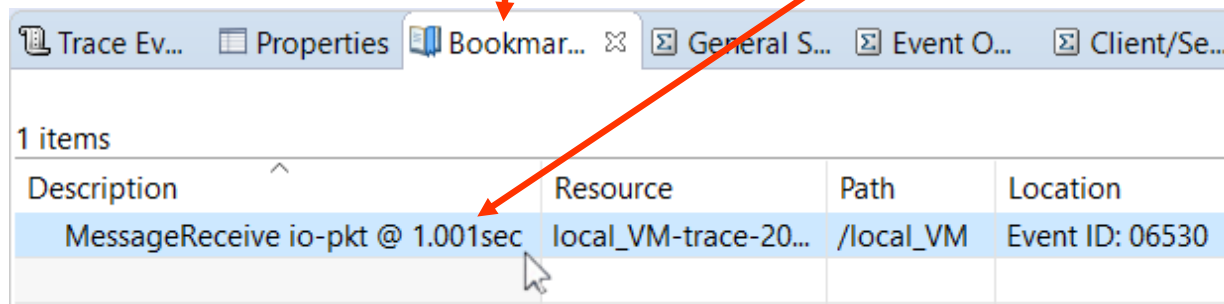
- bookmarks can help you find your way back



Bookmark a selected event:

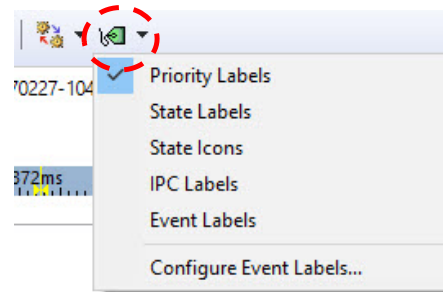


use the Bookmarks view to get back to it

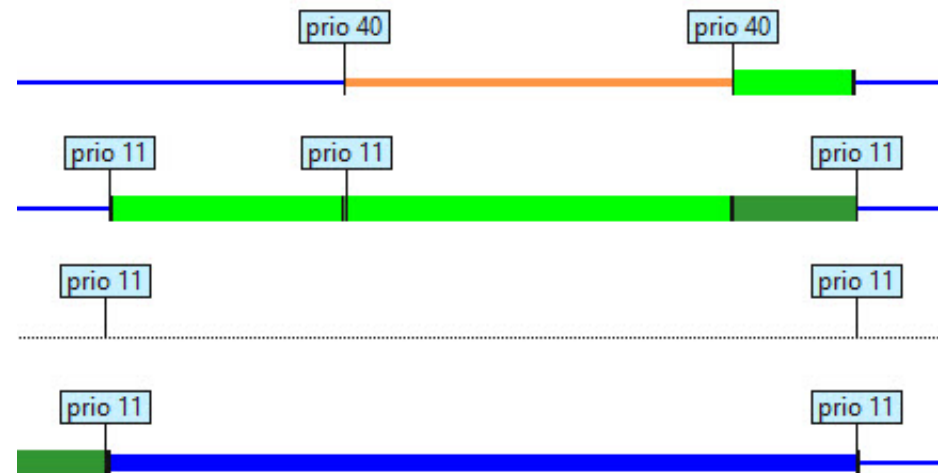


Timeline - Labels

The Timeline provides a few ways of labeling things to make life easier:



- priority labels mark running threads with priority values

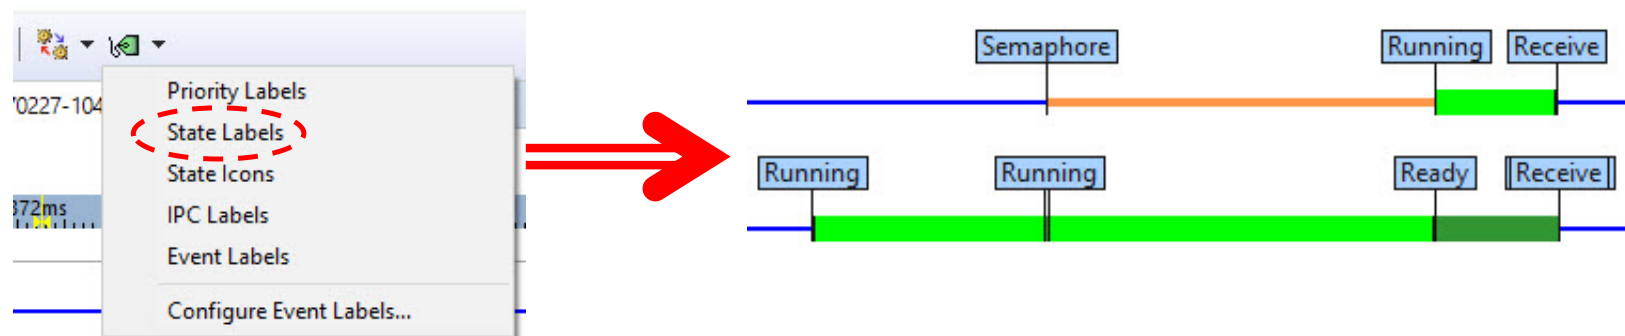


- thread running events must be wide to include this data


Timeline – State Labels

Colours indicate thread states:

- enabling ‘state labels’ means you don’t have to remember the colours

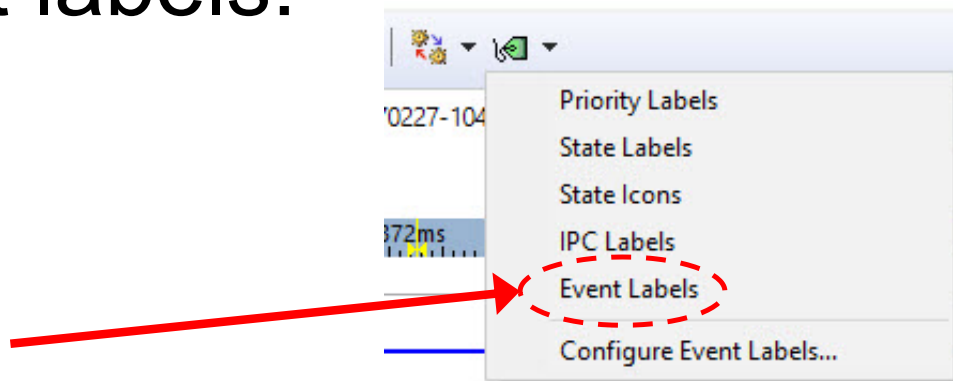


- state icons use less space, but you have to recognise the icons



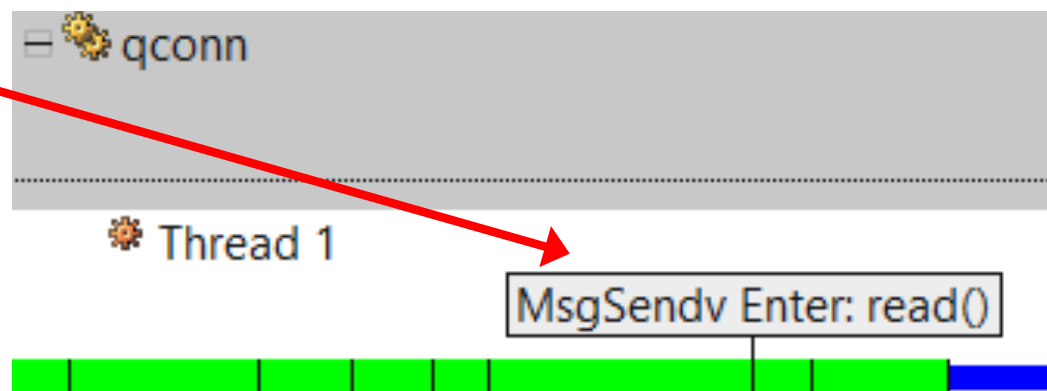
Timeline – Event Labels

Event labels:



– enabling causes some common message pass events to be labeled, e.g.

- *read()*
- *write()*
- *open()*
- *close()*

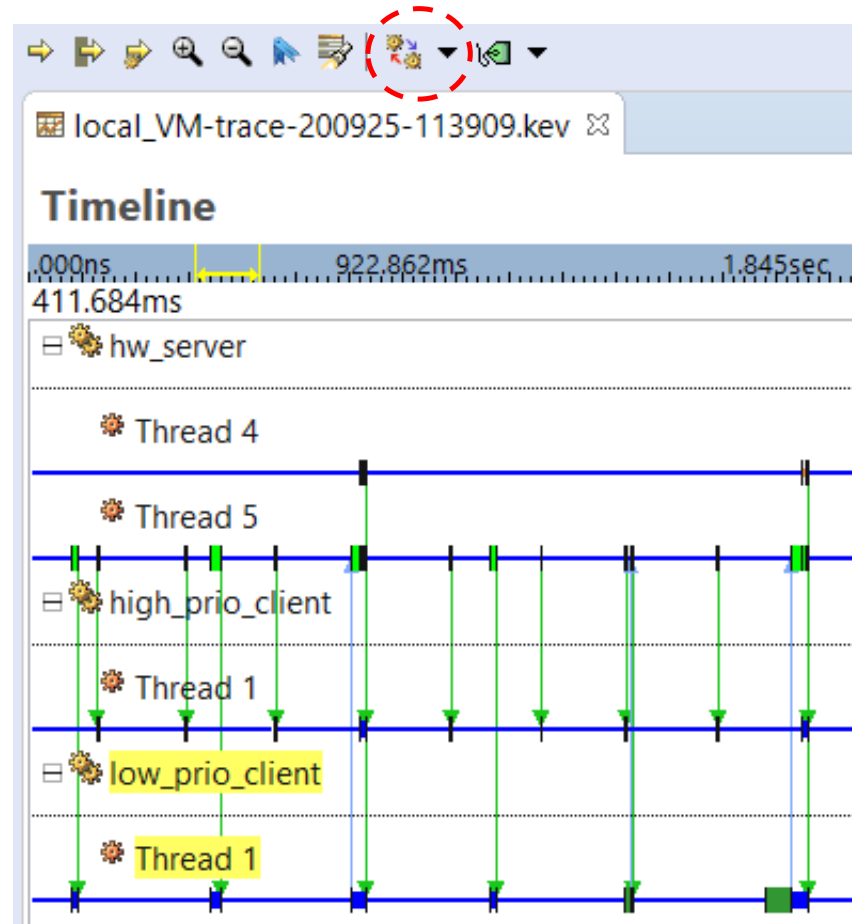


Timeline – IPC Lines

To see what QNX native messaging is occurring:

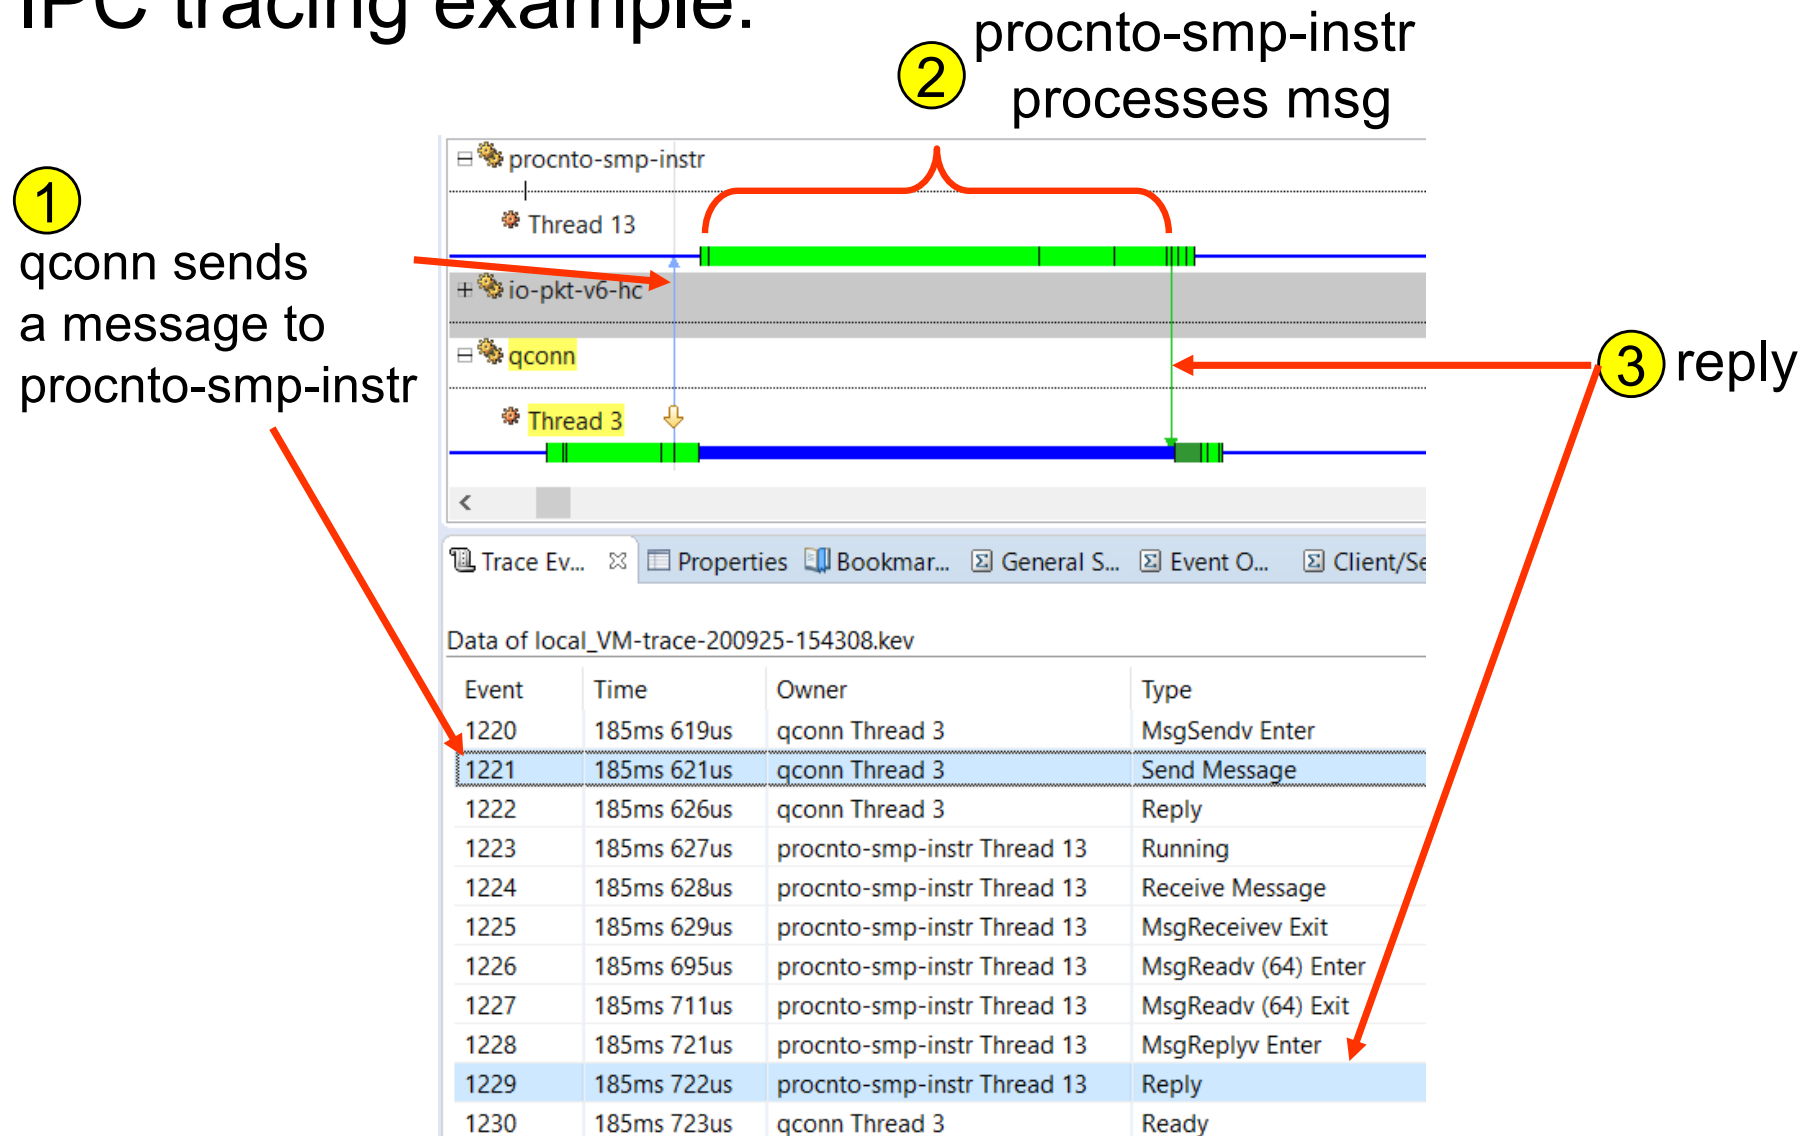
- with 'IPC lines' turned on, vertical lines are drawn between owners to represent IPC

- orange shows pulses
- blue shows sends
- green shows replies
- red shows error replies



Timeline – IPC Tracing

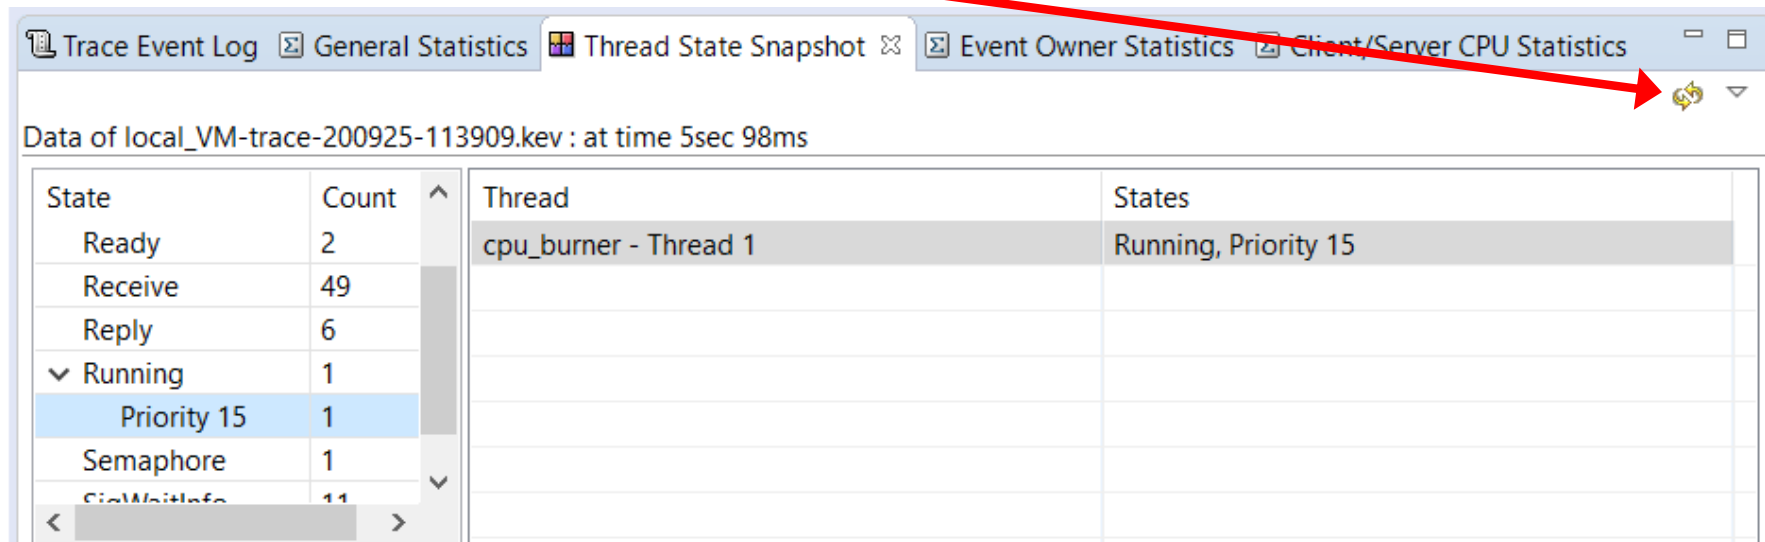
IPC tracing example:



Timeline – Thread State Snapshot

What are the other threads doing?

- the Thread State Snapshot answers this question
- click somewhere in the timeline, then click the refresh button to get the list:



Trace Event Log | General Statistics | Thread State Snapshot | Event Owner Statistics | Client/Server CPU Statistics

Data of local_VM-trace-200925-113909.kev : at time 5sec 98ms

State	Count	Thread	States
Ready	2	cpu_burner - Thread 1	Running, Priority 15
Receive	49		
Reply	6		
Running	1		
Priority 15	1		
Semaphore	1		
SigWaitlife	11		

Timeline - Trace Event Log

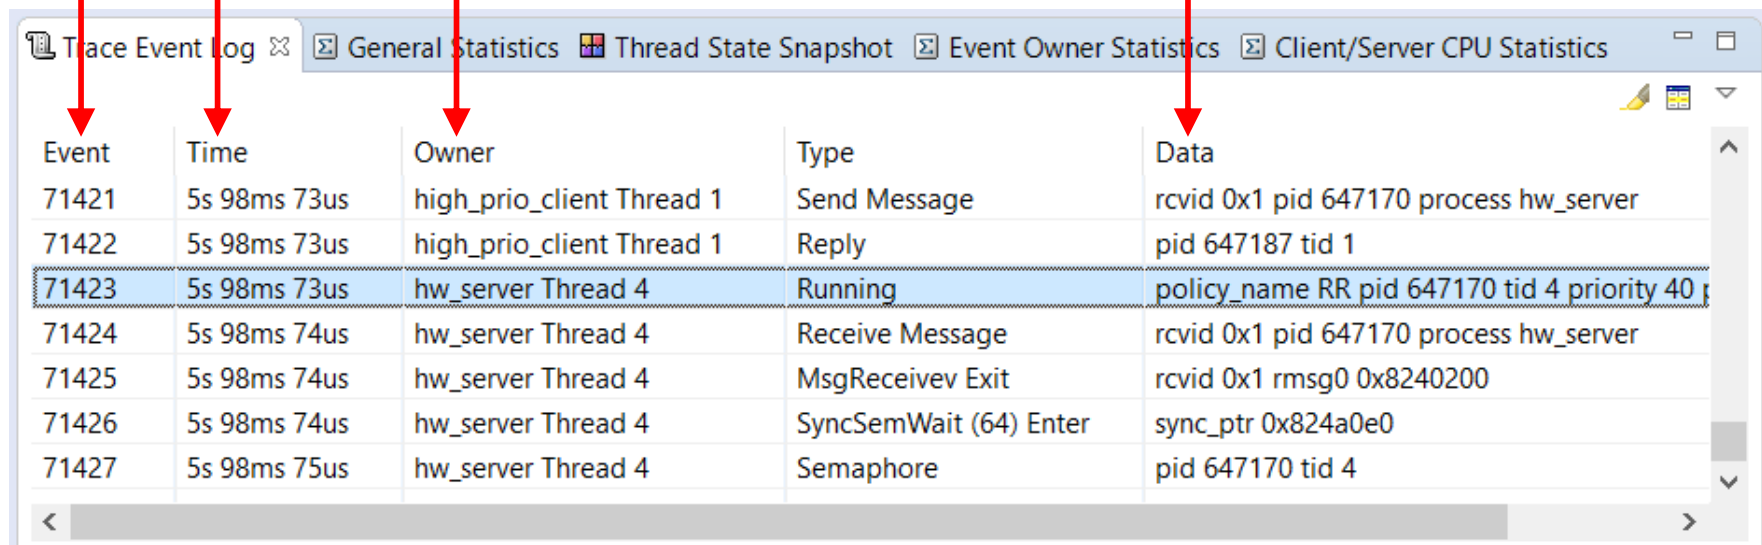
If you select an event in the Timeline, details can be found in the Trace Event Log:

sequence number

high resolution time stamp
based on CPU cycle counter

process, thread, or interrupt
that the event is related to

data associated with event



Event	Time	Owner	Type	Data
71421	5s 98ms 73us	high_prio_client Thread 1	Send Message	rcvid 0x1 pid 647170 process hw_server
71422	5s 98ms 73us	high_prio_client Thread 1	Reply	pid 647187 tid 1
71423	5s 98ms 73us	hw_server Thread 4	Running	policy_name RR pid 647170 tid 4 priority 40
71424	5s 98ms 74us	hw_server Thread 4	Receive Message	rcvid 0x1 pid 647170 process hw_server
71425	5s 98ms 74us	hw_server Thread 4	MsgReceivev Exit	rcvid 0x1 rmsg0 0x8240200
71426	5s 98ms 74us	hw_server Thread 4	SyncSemWait (64) Enter	sync_ptr 0x824a0e0
71427	5s 98ms 75us	hw_server Thread 4	Semaphore	pid 647170 tid 4

Two ways to locate specific events in a log:

– Find

- quicker to use, simpler
- good for most times you need to locate an event
- CTRL-F, or, Edit Menu → Find
- moves selection to first event that is found
- works same way as typical 'Find' in office apps

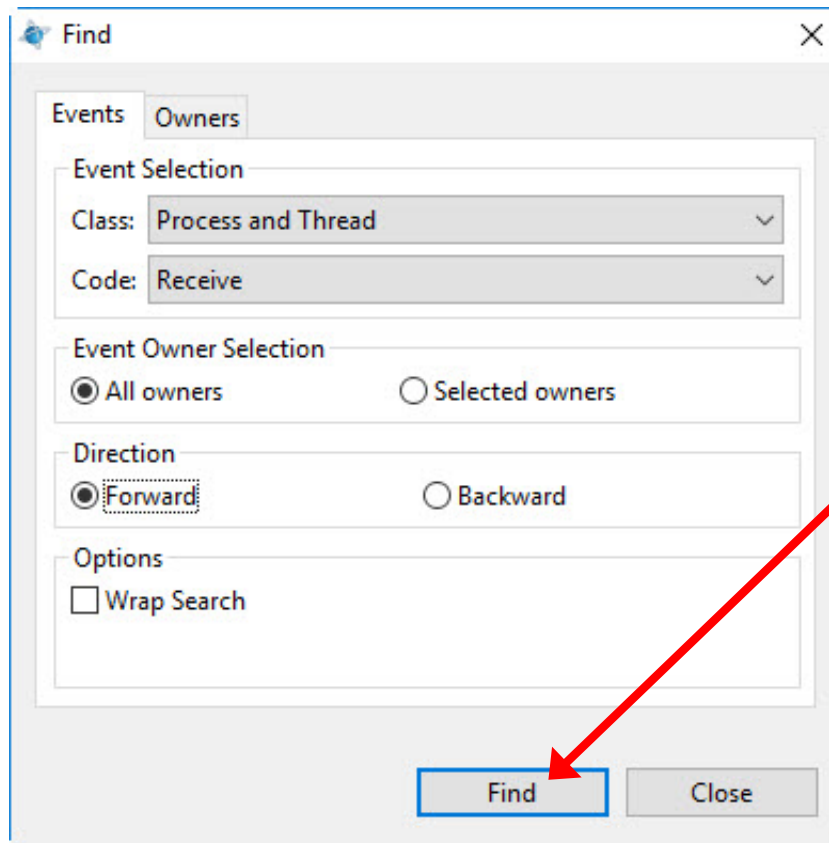
– Search

- CTRL-H, or, Search Menu
- more features
- requires that you set up 'search conditions' that it will look for
- Search view provides you with list of all matches

Finding

Using 'Find':

- setup information about event you want to find

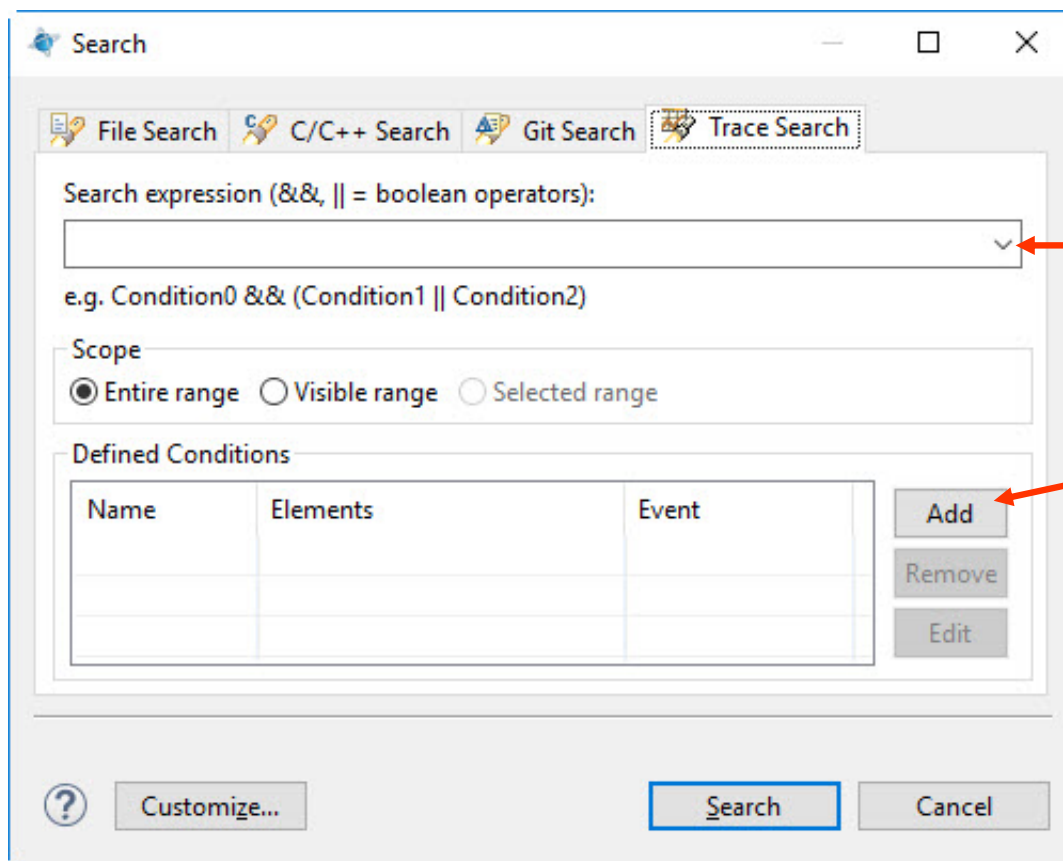


clicking find takes you to next event that is found

Searching

Using 'Search':

- use the Search menu or press Ctrl-H

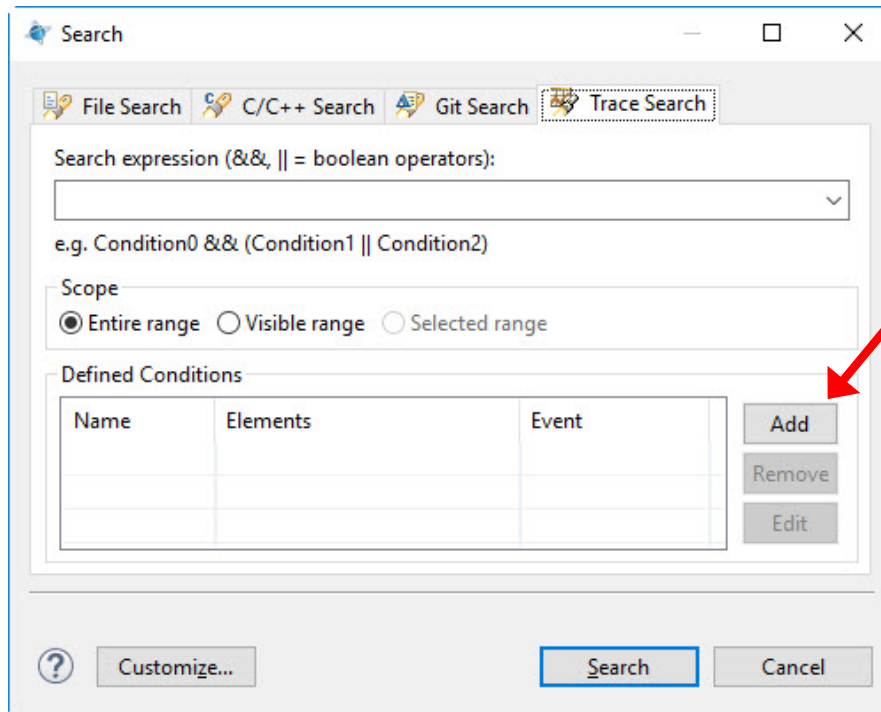


- 1 first add your condition(s) e.g. search for all MsgSends
- 2 then fill in your search criteria by making use of your conditions

Searching

Searching example:

- say we have a process called “**high_prio_client**”, which isn’t getting as much CPU time as we expect
- let’s find circumstances where **high_prio_client** wanted to run (i.e. was ready), and then we can look at what else was running instead



① click ‘Add’ to create a search condition

② let’s look at condition definition on next page...


continued...

Searching

Searching example (continued):

The screenshot shows the 'Trace Condition Wizard' window. The 'Trace Condition Definition' section is active, with the instruction 'Select event and elements for condition'. The 'Name' field contains 'PriorityClient_ready'. The 'Event Selection' section shows 'Class' set to 'Process and Thread' and 'Code' set to 'Ready'. The 'Event Details' section has input fields for 'pid', 'tid', and 'priority'. The 'Element Selection' section has a checkbox for 'Select all elements' which is unchecked, and a list of elements: 'dumper' (unchecked), 'high_prio_client' (checked), and 'hw_server' (unchecked). Red arrows point from the numbered instructions to these specific fields.

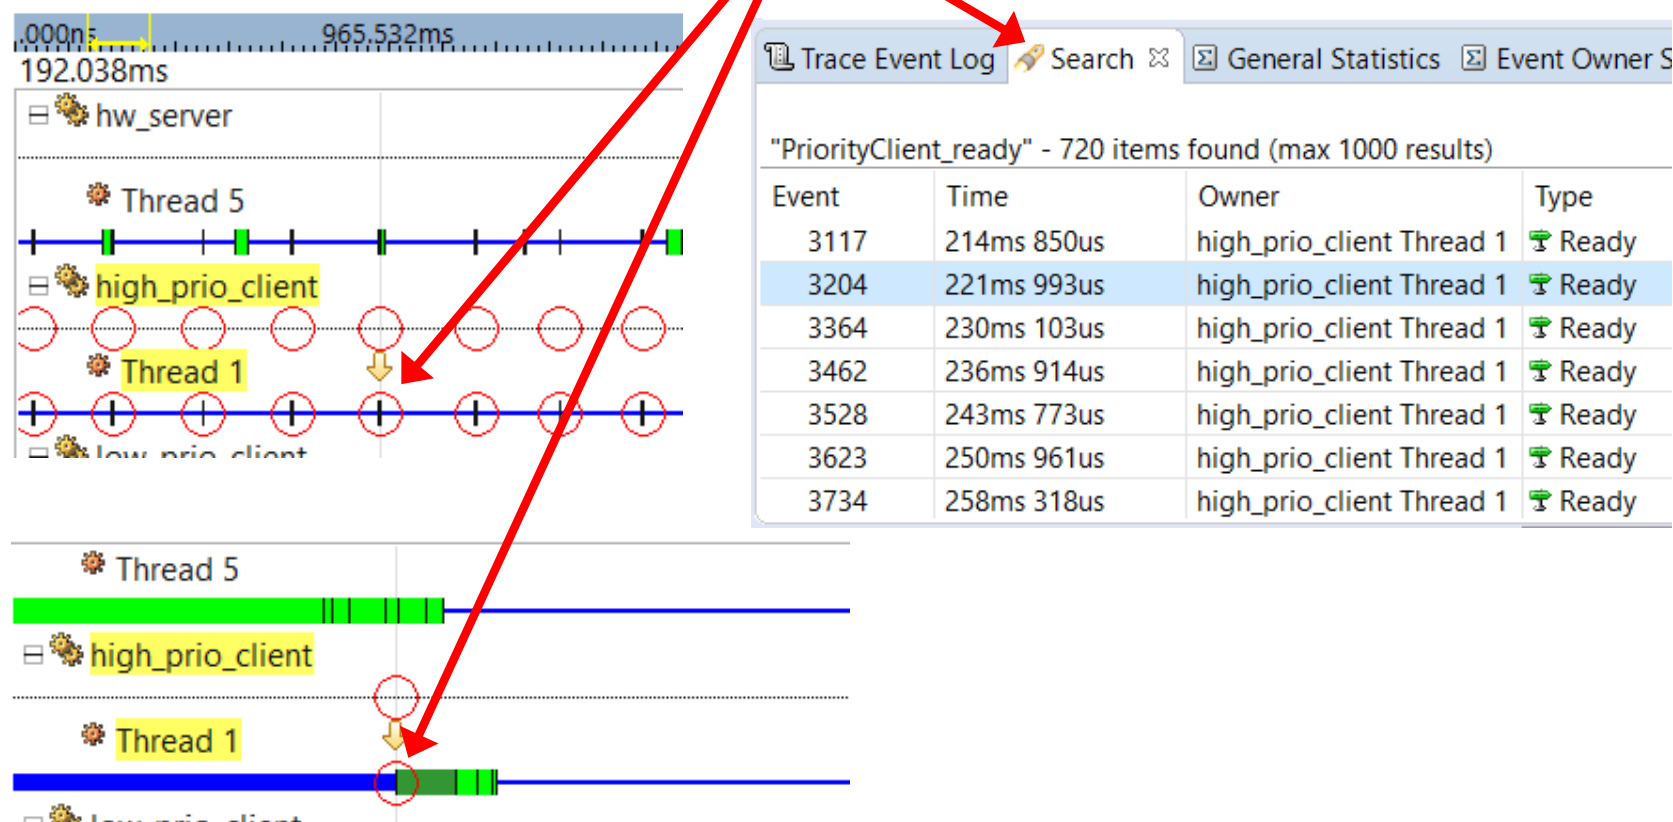
- 1 give the search condition a name
- 2 select event type, if you choose a specific type of event, rather than a whole class, 'Event Details' allows you to refine the search
- 3 unchecking this to be useful to restrict the search to a specific process or thread



Searching

Results from a search:

- events are shown as:



Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

→ – Exercise

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

Multi-core Related Features

Adaptive Partitioning: Partition Summary

Conclusion

EXERCISE – Navigating a Log

Try out the log navigation techniques:

1. capture a new log, or use an existing one
2. switch panes e.g. CPU Activity, CPU Usage
3. zoom in/out, scroll
4. open some processes to see threads and thread states
5. match events in Timeline with Event Log, try and become familiar with format of events in log
6. bookmark an event
7. use Find and Search to locate an event where qconn goes into the running state.
8. turn on labels and look at the resulting information



Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

→ Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

Multi-core Related Features

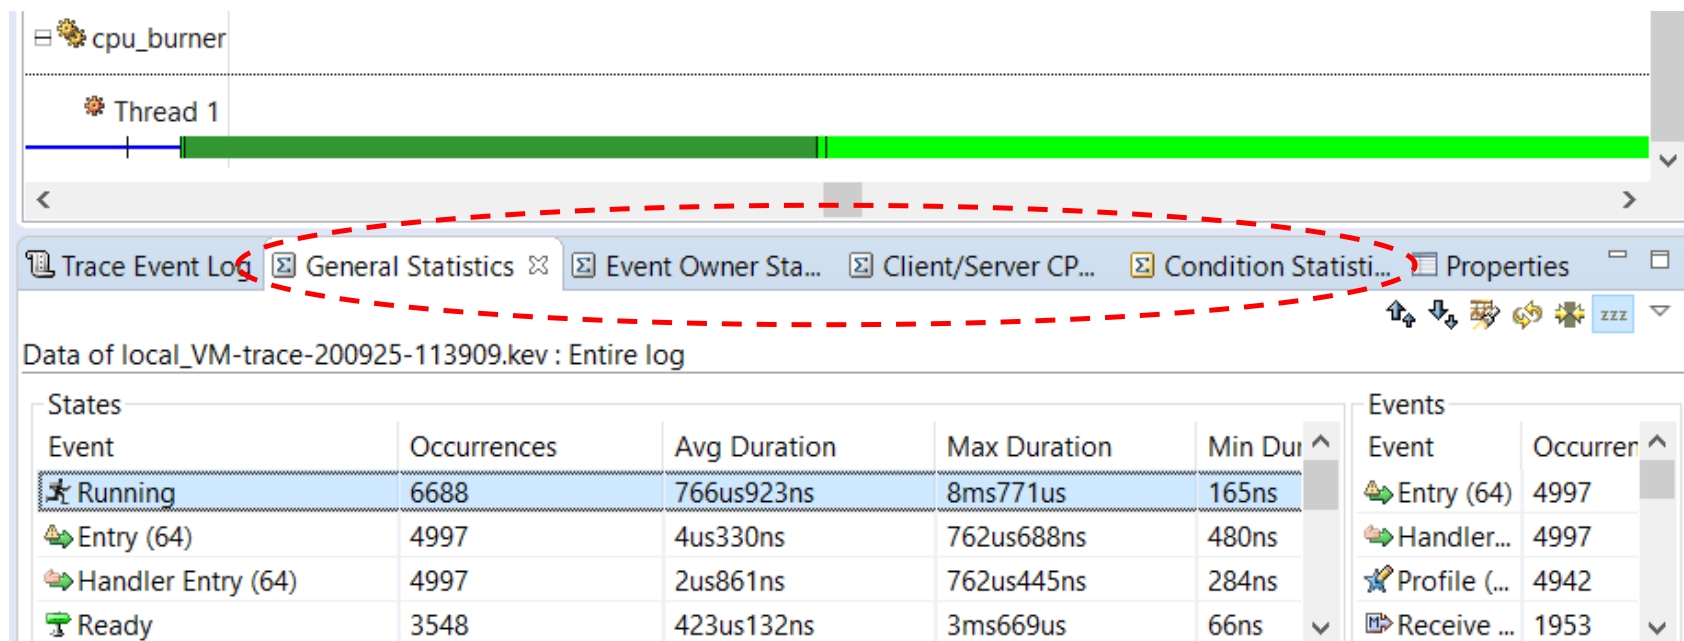
Adaptive Partitioning: Partition Summary

Conclusion

Statistics Views

Statistics views:


- System Profiler can provide various statistical info.
- useful for:
 - getting an overall feel for what went on during a log
 - seeing some kinds of problem situations
 - e.g. a thread spend a really long time waiting for a mutex to be




Statistics Views

The Statistics Views all share some common behaviors:


- they don't display any data until you click the refresh statistics button:



- they can all operate on either the entire log, or a selected area, controlled by a toggle:



- all can generate a CSV report of their data:



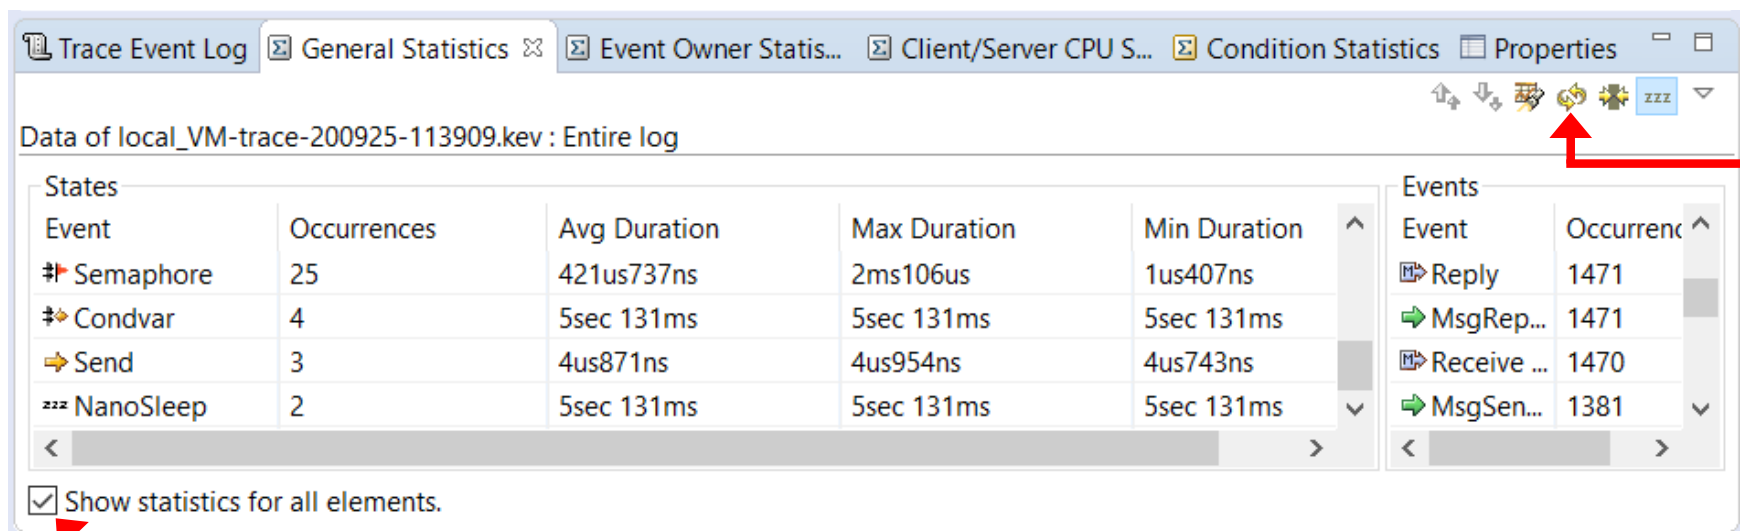
The Statistics Views are:

- General Statistics
 - give a bunch of general statistics for all owners, or selected owners
- Event Owner Statistics
 - collect statistics for a configurable list of events on a per owner basis
- Condition Statistics
 - do a time range count of configured conditions
- Client/Server CPU Statistics
 - accumulate server time back to a client when appropriate

Statistics

General Statistics: 2 ways to use it:

- show stats for all processes/threads
- show stats for only selected process/thread (or multiple)



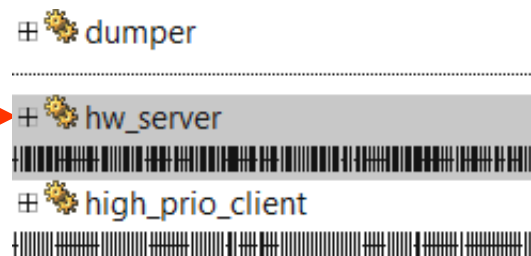
- 1 choose to show stats for all processes, or only a selection
if you uncheck this, make sure to have one or more processes or threads
selected in the timeline
- 2 click Refresh

General Statistics

Let's use the General Statistics view:

- we'll try to find the longest time a thread in *hw_server* is semaphore blocked for, and where this occurred

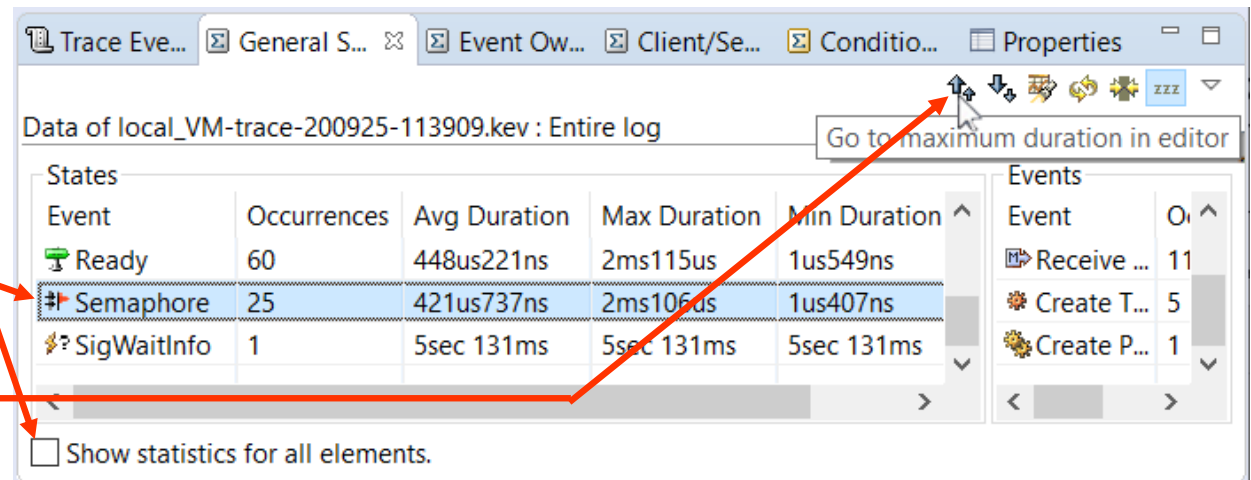
we'll start by selecting *hw_server* in the timeline



unclick all elements and refresh if needed

select Semaphore

and click on “go to maximum duration”



States	Event	Occurrences	Avg Duration	Max Duration	Min Duration
	Ready	60	448us221ns	2ms115us	1us549ns
	Semaphore	25	421us737ns	2ms106us	1us407ns
	SigWaitInfo	1	5sec 131ms	5sec 131ms	5sec 131ms

Go to maximum duration in editor

Show statistics for all elements.

Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

→ CPU Activity Pane and CPU Usage Pane

– Exercise

Tying The Trace to Your Code

Multi-core Related Features


Adaptive Partitioning: Partition Summary

Conclusion

CPU Activity

CPU Activity:

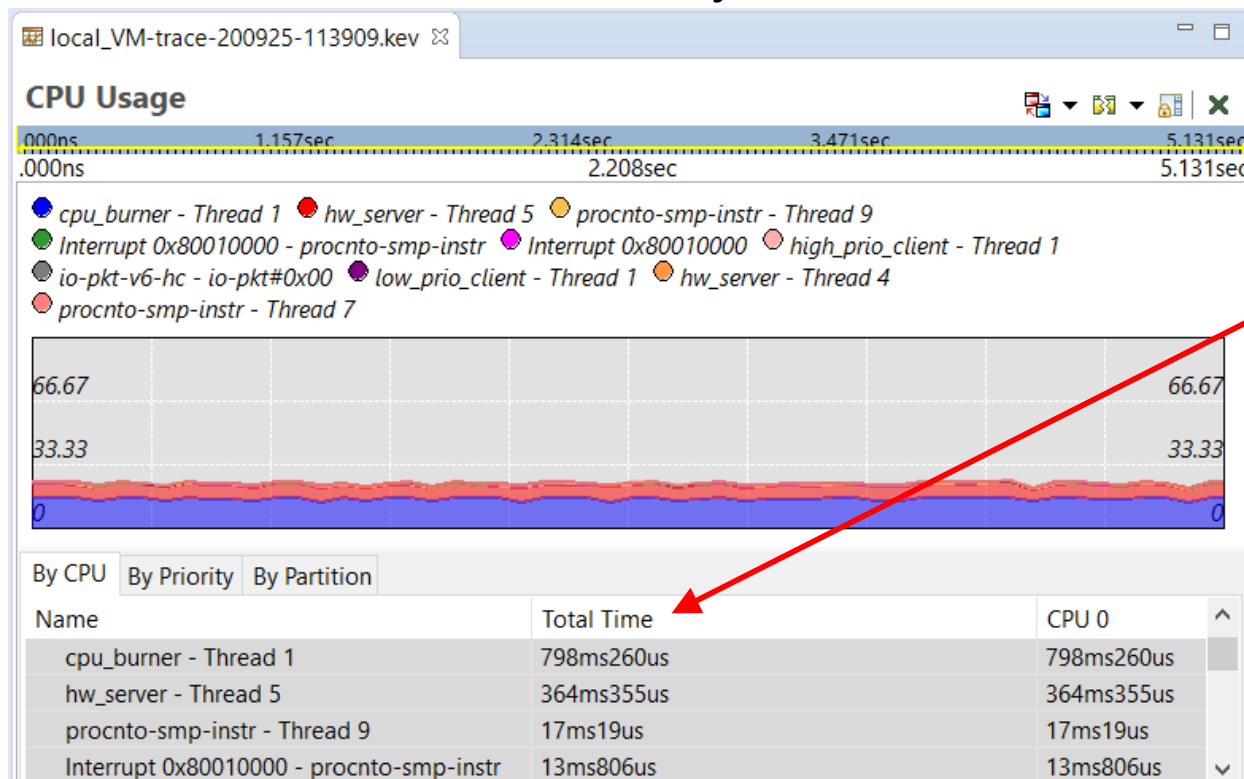
- tracks total CPU usage over time



CPU Usage

CPU Usage:

- shows the CPU consumption for the time period selected in the Timeline, showing the top ten CPU users by default
- can be shown in %, time, or both (configured in Window menu→Preferences→System Profiler→CPU Usage)

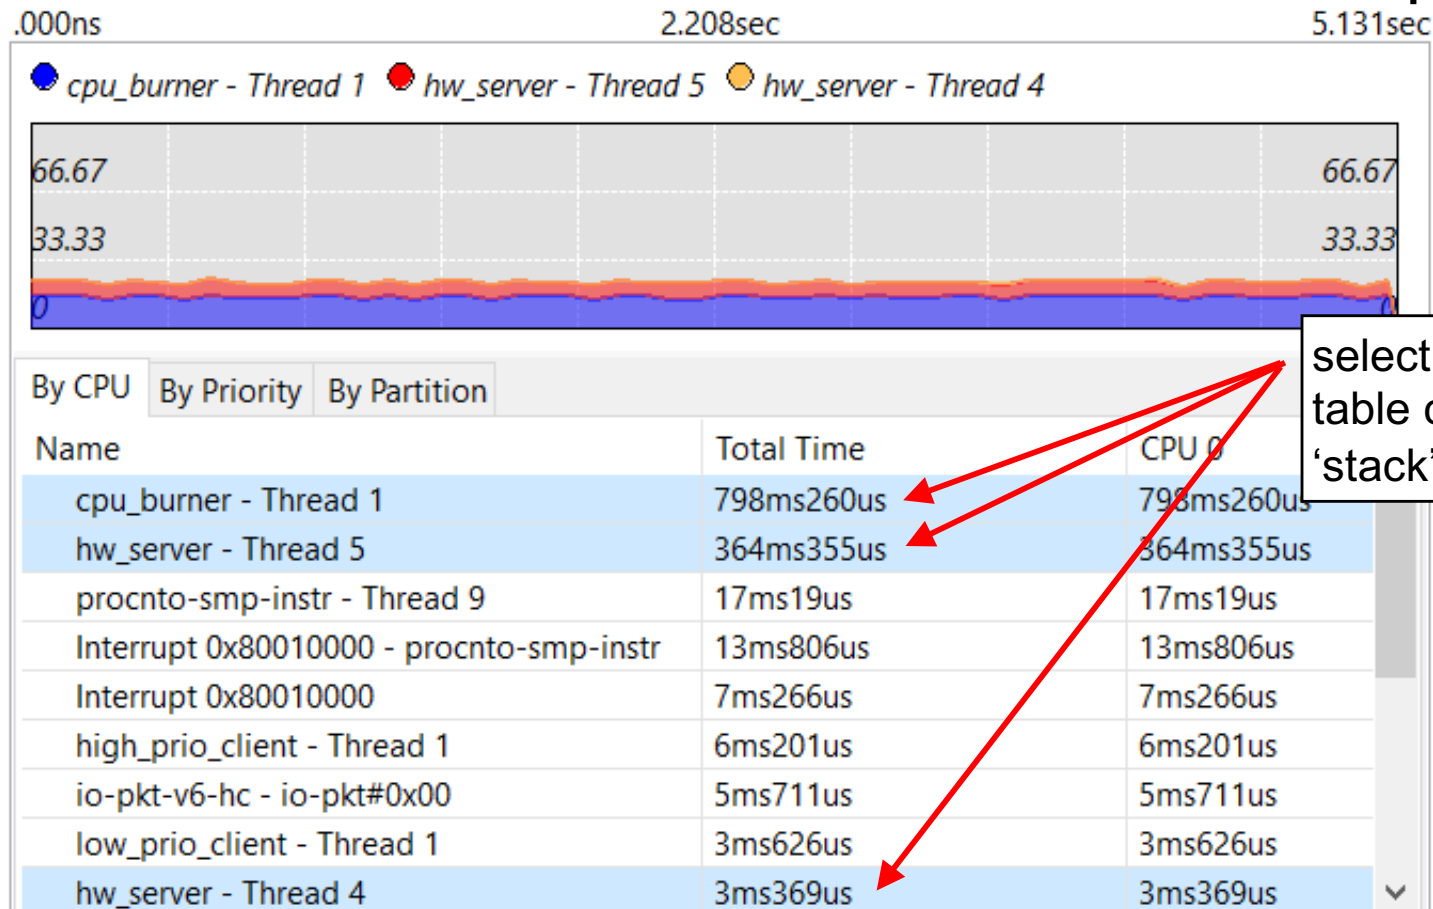


threads are sorted in order of greatest CPU user

CPU Usage

You can select multiple threads in CPU Usage:

- hold down CTRL to select/deselect multiple

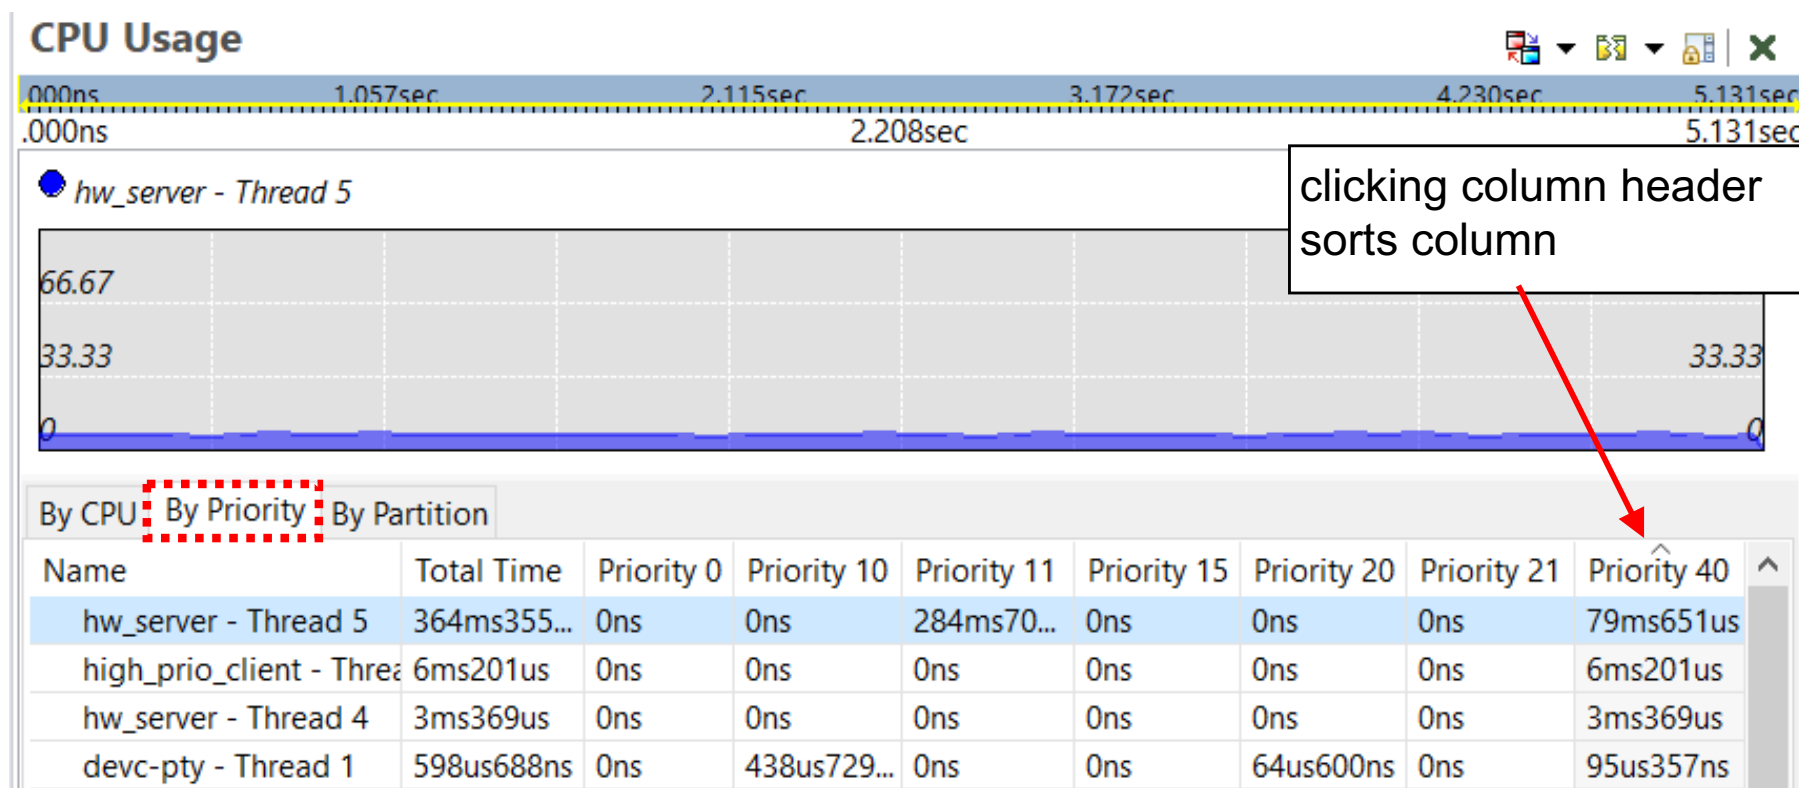


selecting threads in the table causes them to 'stack' visually

CPU Usage

CPU Usage can be sorted 'By Priority':

- good for analyzing circumstances where threads at the same priority are competing for CPU time



The CPU usage of a particular thread may not give the whole story

- Neutrino is a message passing based system
- threads will send messages to server threads to get work done on their behalf
- Client/Server CPU Statistics is a tool that can help you take server execution time into account

Client/Server CPU Statistics

Client/Server CPU Statistics:

list of all threads that ran during the log

how much total time that the thread consumed (itself + server time)

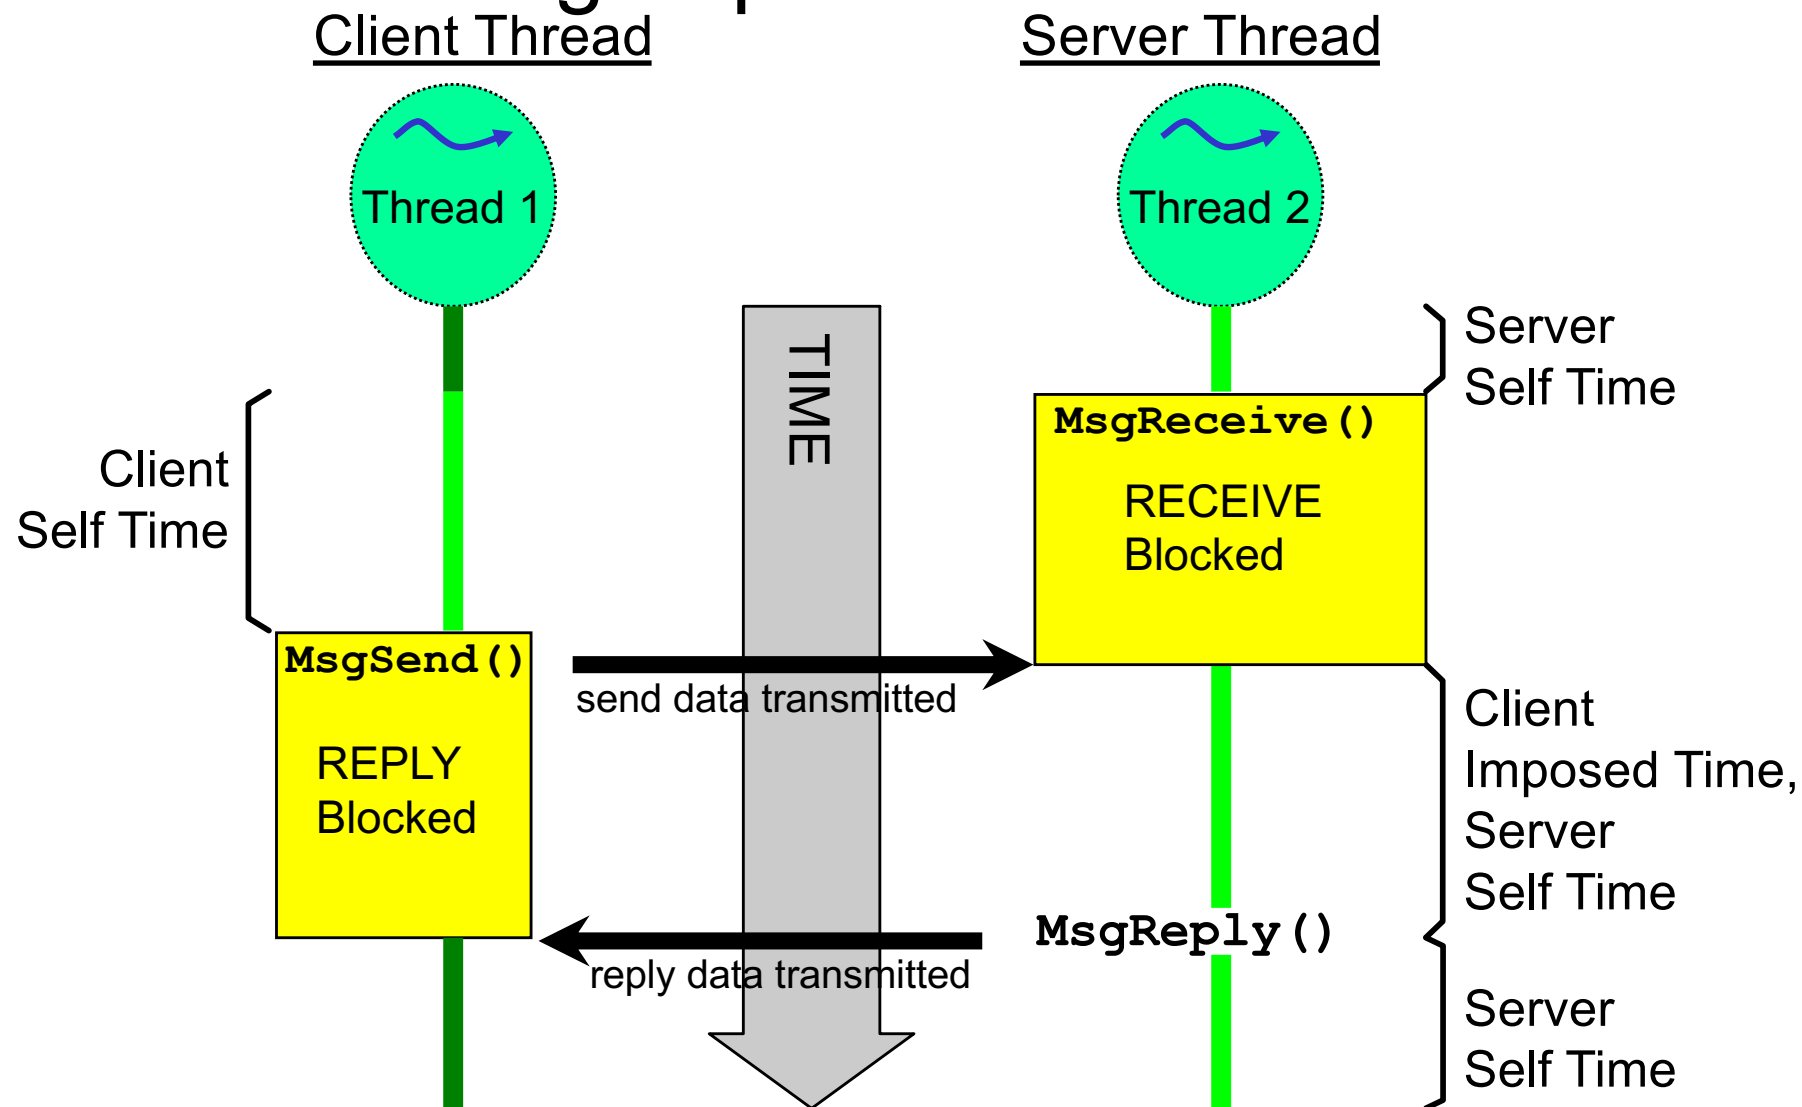
time thread itself ran for

time that any servers consume while doing work on thread's behalf

Owner	Total Time	Self Time	Imposed Time
cpu_burner - Thread 1	799ms168us	799ms168us	0ns
devb-eide - eide_driver_thread	29us228ns	29us228ns	0ns
devc-pty - Thread 1	601us194ns	601us194ns	0ns
high_prio_client - Thread 1	87ms31us	6ms210us	80ms820us
hw_server - Thread 4	3ms370us	3ms370us	0ns
hw_server - Thread 5	364ms577us	364ms577us	0ns
io-pkt-v6-hc - io-pkt#5x00	5ms749us	5ms749us	0ns
low_prio_client - Thread 1	287ms430us	3ms628us	283ms801us
procnto-smp-instr - Thread 7	1ms606us	1ms606us	0ns

Client/Server CPU Statistics

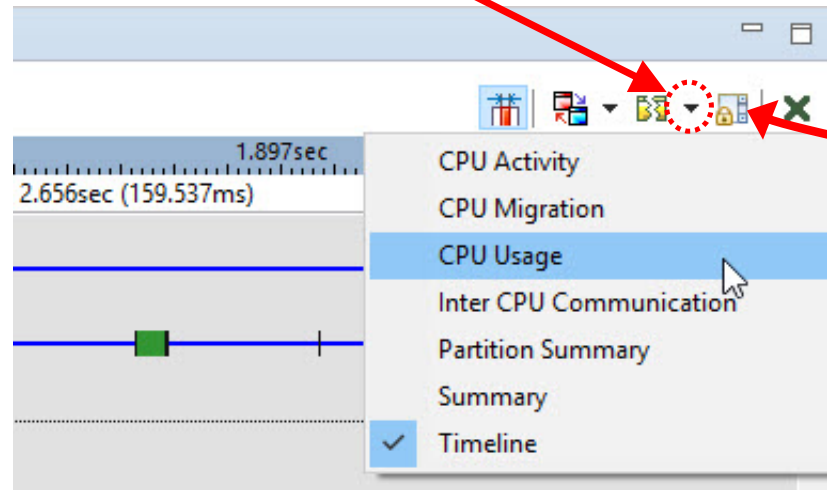
Understanding Imposed Time:



Panes – Splitting and Locking

Splitting and locking panes:

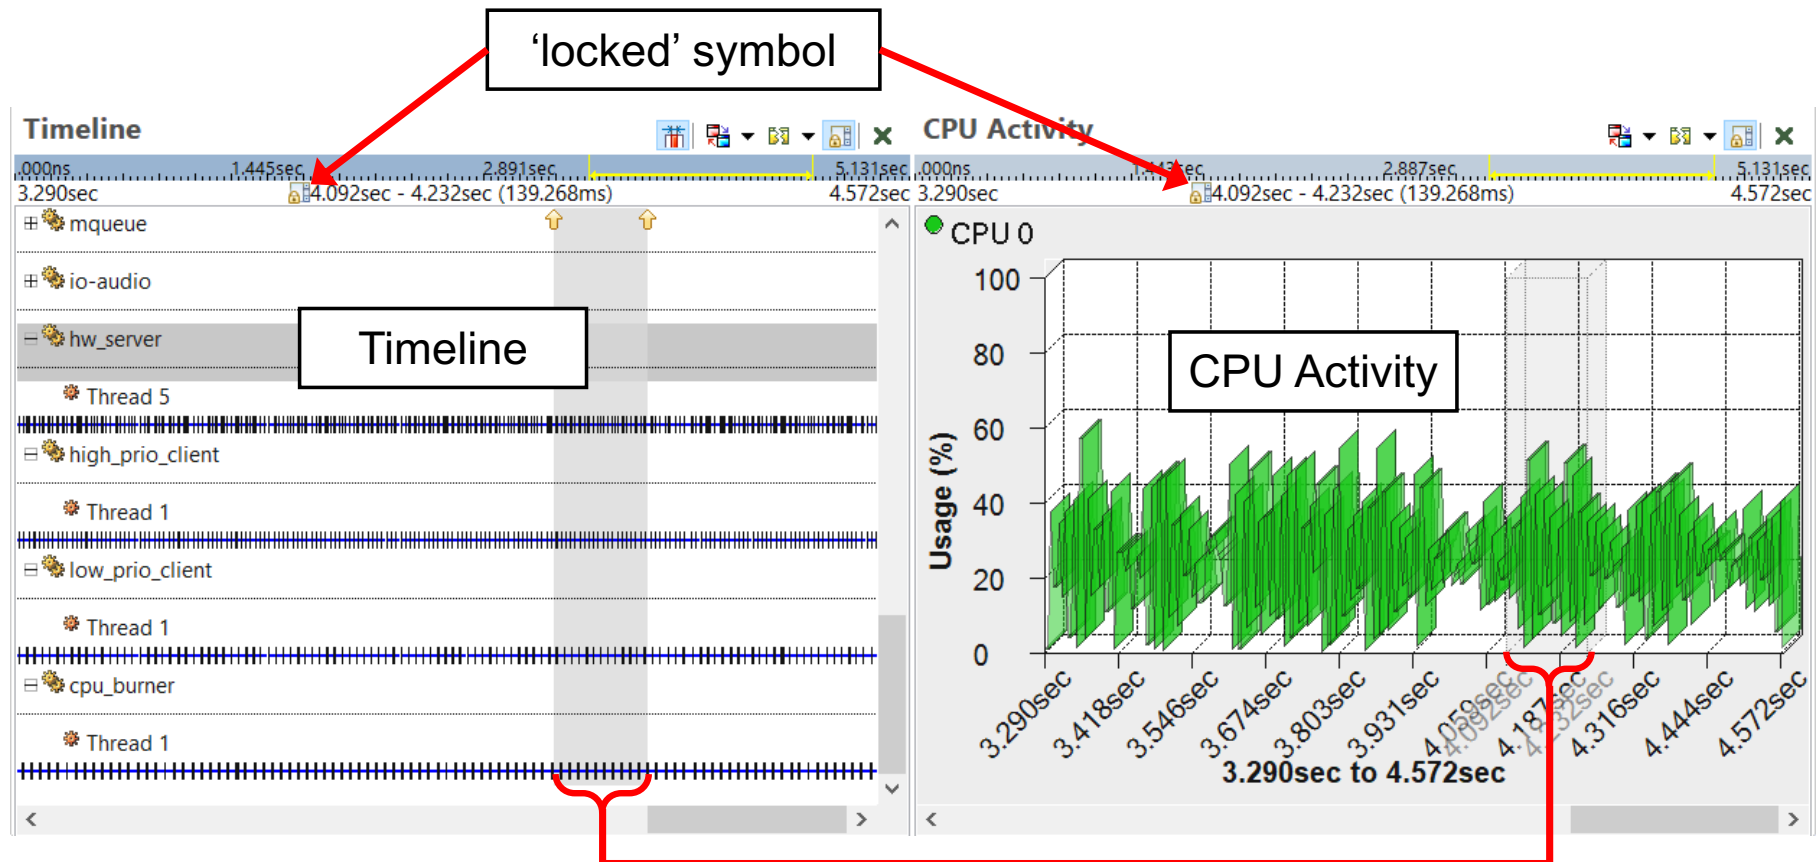
- you can split the panes
 - e.g. pane can be half Timeline and half CPU Usage
 - use this ‘pull-down’ icon



- panes can be ‘locked’
 - any scrolling, selection, or zooming that is done in one panes is reflected in the other(s)
 - all panes need to be locked

Displays – Split and Locked Example

Example of split and locked panes:



- after locking, selection and scrolling for both panes are synchronized

Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane



– Exercise

Tying The Trace to Your Code

Multi-core Related Features

Adaptive Partitioning: Partition Summary

Conclusion

EXERCISE – CPU Usage and CPU Activity

Examine CPU usage:

1. use CPU Activity to gauge how busy the system was overall
2. use CPU Usage to determine which (non-idle) thread consumed the most CPU time
 - hold CTRL down and select the 2nd and 3rd busiest threads to see their CPU usage as well
3. use Client/Server Stats to look at time that clients forced servers to run
 - do you still consider the thread you found in 2. to have been the busiest?
4. split and lock CPU Usage with Timeline
 - zoom in on a peak of usage and look at the events in the Timeline to try and determine the reason for the peak

Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

→ Tying The Trace to Your Code

– Exercise

Multi-core Related Features

Adaptive Partitioning: Partition Summary

Conclusion

Tying Your Code to the Trace

So your client sent a message to the server...

- where in the client did it send this?
- there are two main ways to try and get this information:
 - annotate your code by inserting trace log events yourself
 - build your application(s) with Function Instrumentation Profiling and incorporate that data into your trace log and analysis

Logging your own Events

The *trace_log*()* functions can be used to log your own events:

- benefits:
 - general method of instrumenting your code
 - can be useful for aligning log events to your code
 - can be used from an interrupt handler
- different user event types:
 - simple user events - the data is 2 integers
 - string user events - the data is a string
 - complex user events – the data is a buffer



Logging your own Events

Examples of logging user events:

```
/* log a simple user event */  
trace_logi(_NTO_TRACE_USERFIRST, 3, 5);
```

```
/* log a string user event */  
trace_logf( _NTO_TRACE_USERFIRST+1,"Hello world");
```

- the first argument is a code, which must be between `_NTO_TRACE_USERFIRST` and `_NTO_TRACE_USERLAST`
- for our two events we've used different but consecutive event IDs starting from the value of `_NTO_TRACE_USERFIRST`
- when looking at the log, you will see “User Event 0” for example, the 0 is `_NTO_TRACE_USERFIRST`
- with `trace_logf()`, it can also be a *printf()* style format string:

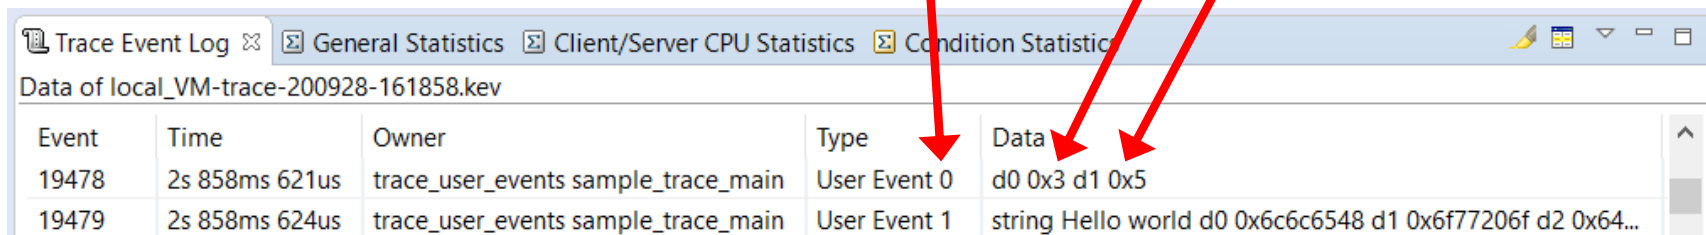
```
trace_logf( _NTO_TRACE_USERFIRST+2,"value: %d", value );
```



Logging your own Events

The result of the previous page:

```
trace_logi(_NTO_TRACE_USERFIRST, 3, 5);
```



Event	Time	Owner	Type	Data
19478	2s 858ms 621us	trace_user_events sample_trace_main	User Event 0	d0 0x3 d1 0x5
19479	2s 858ms 624us	trace_user_events sample_trace_main	User Event 1	string Hello world d0 0x6c6c6548 d1 0x6f77206f d2 0x64...

```
trace_logf(_NTO_TRACE_USERFIRST+1, "Hello world");
```

Note: the IDE doesn't know what type of data you generated, it guesses based on the contents:

- integer elements with zeros in them may be displayed as a string

Or, you can tell the IDE how to display your events:

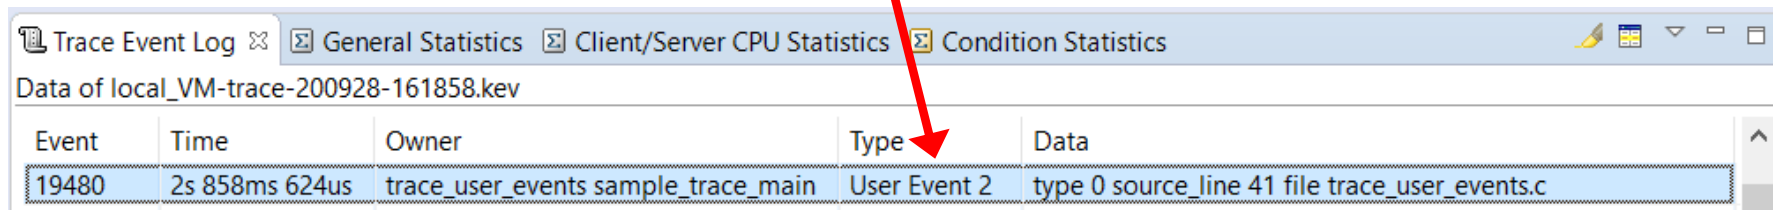
- can be global:
 - Windows-Preferences->QNX->System Profiler->User Event Data
- or per key file:
 - Menu Click -> Properties -> User Event Data
- you will be specifying an xml file that tells the IDE how to parse your events

Parsing User events

For example, to parse:

```
struct mydata
{
    int type; // 0 - trace, 1 - warn, 2 = error
    int line;
    char source_file[20];
} my_data_a;
...
trace_logb( _NTO_TRACE_USERFIRST+2, &my_data_a,
            sizeof(my_data_a));
```

and get:



Trace Event Log				
Data of local_VM-trace-200928-161858.kev				
Event	Time	Owner	Type	Data
19480	2s 858ms 624us	trace_user_events sample_trace_main	User Event 2	type 0 source_line 41 file trace_user_events.c

you would...

Parsing User Events

... use an event definition file like:

```
<?xml version="1.0" encoding="UTF-8" ?>
```

```
<eventdefinitions>
```

```
  <eventclass id="6" name="User Events">
```

```
    <event id="2" sformat="%4uld type %4uld source_line %1s0 file" />
```

```
  </eventclass>
```

```
</eventdefinitions>
```

- the **sformat** specifies a series of data and name specifications, each delimited by a space



Function Tracing is primarily an Application Profiling tool

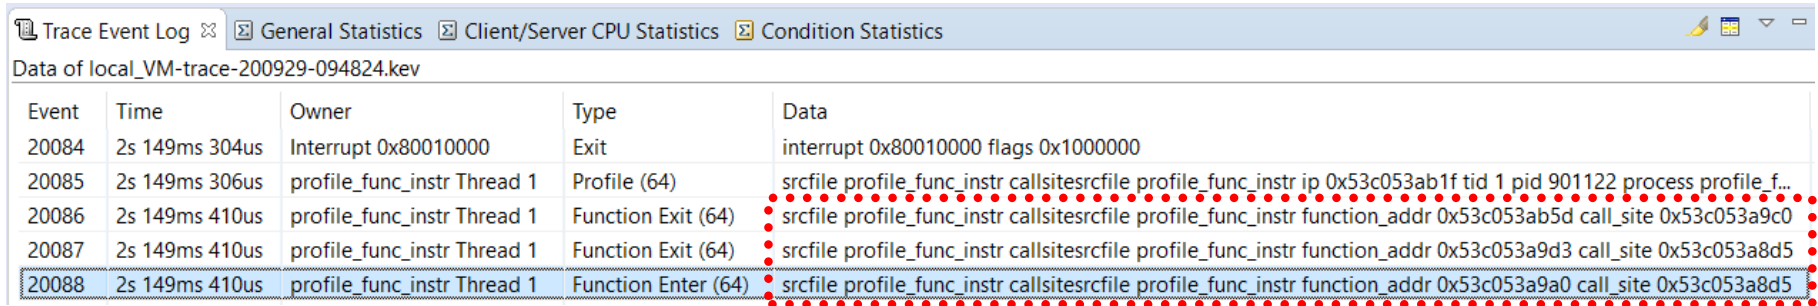
- it can be used with System Profiling
 - the data output must be configured to go to the kernel trace log
- data from multiple applications can be collected in the same log
- all source that you wish traced must be specially compiled
- all applications must be linked to the profiling library

Preparing for Function Tracing:

- compile your source files to generate the tracing code:
 - `CFLAGS += -finstrument-functions`
 - this adds code to the start and end of every function call that generates data
- link against the tracing library:
 - `-lprofilingS`
 - this is a library, it must be after the objects on the link line, e.g.
 - `$(LD) $(LDFLAGS) main.o funcs.o -lprofilingS -o my_prog`
- tell the tracing library to send its data to the kernel event log:
 - set the environment variable `QPROF_KERNEL_TRACE=1`, e.g.:
 - `QPROF_KERNEL_TRACE=1 my_prog &`
 - you may wish to export this early or always set it, since it won't affect uninstrumented applications

Function Tracing

In the log it looks like:




Event	Time	Owner	Type	Data
20084	2s 149ms 304us	Interrupt 0x80010000	Exit	interrupt 0x80010000 flags 0x1000000
20085	2s 149ms 306us	profile_func_instr Thread 1	Profile (64)	srcfile profile_func_instr callsitesrcfile profile_func_instr ip 0x53c053ab1f tid 1 pid 901122 process profile_f...
20086	2s 149ms 410us	profile_func_instr Thread 1	Function Exit (64)	srcfile profile_func_instr callsitesrcfile profile_func_instr function_addr 0x53c053ab5d call_site 0x53c053a9c0
20087	2s 149ms 410us	profile_func_instr Thread 1	Function Exit (64)	srcfile profile_func_instr callsitesrcfile profile_func_instr function_addr 0x53c053a9d3 call_site 0x53c053a8d5
20088	2s 149ms 410us	profile_func_instr Thread 1	Function Enter (64)	srcfile profile_func_instr callsitesrcfile profile_func_instr function_addr 0x53c053a9a0 call_site 0x53c053a8d5

- for finding these events, they are in the System category
- but we only have raw function addresses
 - if the application was built with symbol information we can translate this to function names and lines
 - we can do this...

Function Tracing

To add source code information:

In the Project Explorer view, bring up the Properties of the key file

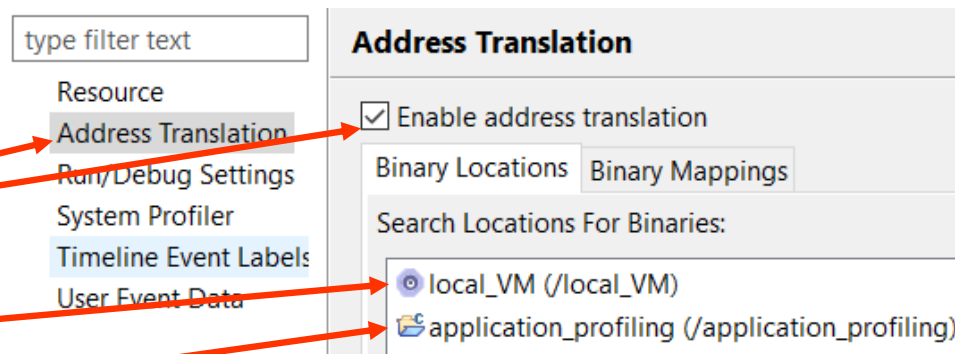


Select Address Translation

Click to Enable

Remove the Target Project


And add source locations



Enter binary load address

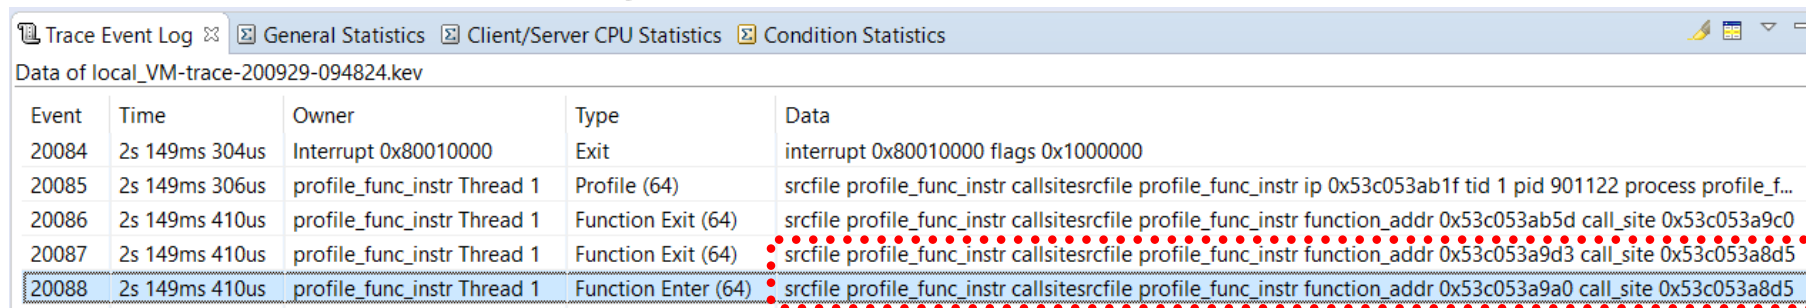
Then add the binaries for which you want translations

You can add libraries to a binary as well



Function Tracing

So, after adding translation:

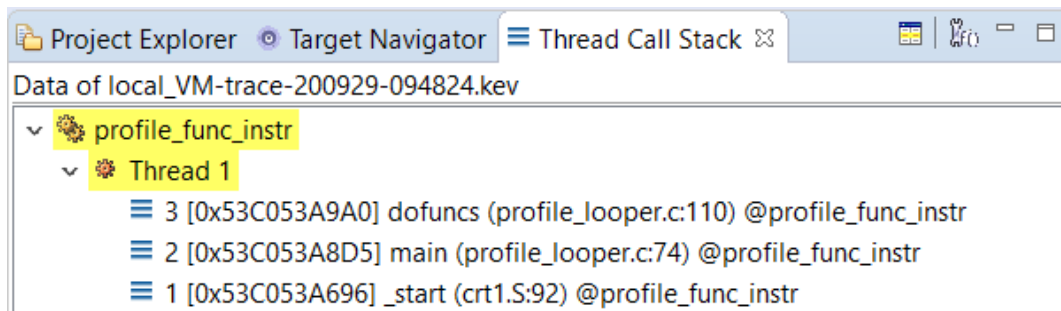


Event	Time	Owner	Type	Data
20084	2s 149ms 304us	Interrupt 0x80010000	Exit	interrupt 0x80010000 flags 0x1000000
20085	2s 149ms 306us	profile_func_instr Thread 1	Profile (64)	srcfile profile_func_instr callsitesrcfile profile_func_instr ip 0x53c053ab1f tid 1 pid 901122 process profile_f...
20086	2s 149ms 410us	profile_func_instr Thread 1	Function Exit (64)	srcfile profile_func_instr callsitesrcfile profile_func_instr function_addr 0x53c053ab5d call_site 0x53c053a9c0
20087	2s 149ms 410us	profile_func_instr Thread 1	Function Exit (64)	srcfile profile_func_instr callsitesrcfile profile_func_instr function_addr 0x53c053a9d3 call_site 0x53c053a8d5
20088	2s 149ms 410us	profile_func_instr Thread 1	Function Enter (64)	srcfile profile_func_instr callsitesrcfile profile_func_instr function_addr 0x53c053a9a0 call_site 0x53c053a8d5

Now looks like:

Function Exit (64) srcfile C:/ide-7.1-workspace/application_profiling/profile_looper.c srcline 114
Function Enter (64) srcfile C:/ide-7.1-workspace/application_profiling/profile_looper.c srcline 110

And the Thread Call Stack view can tell you where you are:



Naming Threads

You can also name your threads:

- “interrupt_thread” is more useful than “Thread 2”
- for example:

```
pthread_setname_np(0, "my_thread_name" );
```

- this will also name the thread for the System Information Perspective

Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code



– Exercise

Multi-core Related Features

Adaptive Partitioning: Partition Summary

Conclusion

EXERCISE – `trace_log*()` and naming threads

Experiment with *trace_log*()* and naming threads:

- modify a program to add some *trace_log*()* calls to insert traces into the log
 - the compiler built-in variables `__LINE__` and `__FILE__` may be useful
- also name any thread or threads in the program
- run the program and get a log
- find those events in the resulting log
- rebuild a program with function tracing as well, or run a program already built that way, and get that data

Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

→ Multi-core Related Features


Adaptive Partitioning: Partition Summary

Conclusion

Multi-core and SMP Features

Migration of a thread between CPUs:

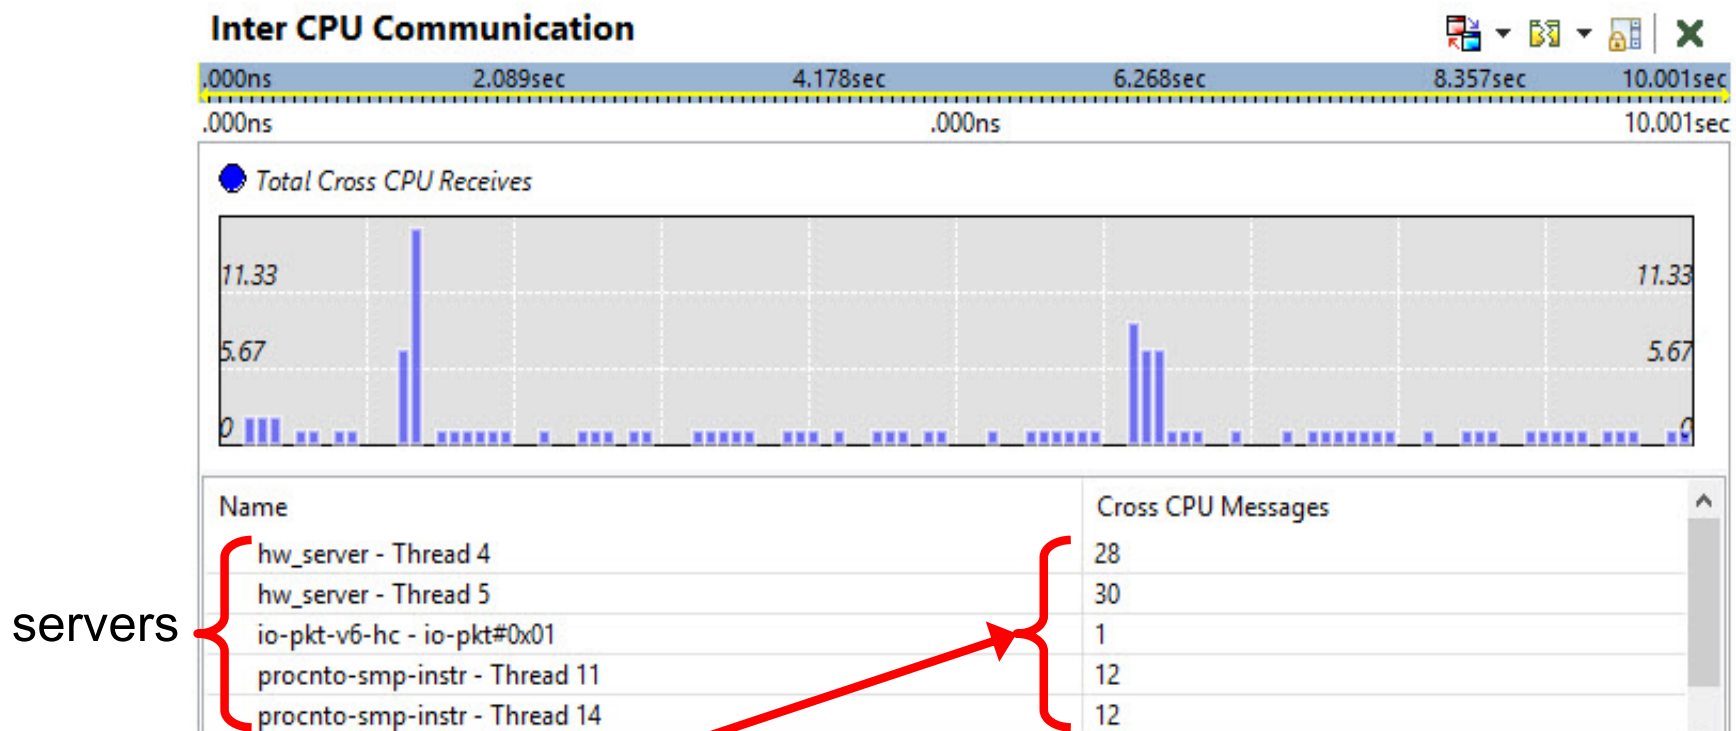
- allows a thread to run when it would not have the opportunity if there was no migration
- may cause a performance penalty due to inefficient use of CPU cache
- CPU Migration Pane will indicate how many migrations occurred



Multi-core and SMP Features

Inter-CPU communication Pane:

- cross-CPU message passing also can cause an inefficient utilization of the CPU cache



number of times clients and servers were scheduled on a different CPU

Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

Multi-core Related Features

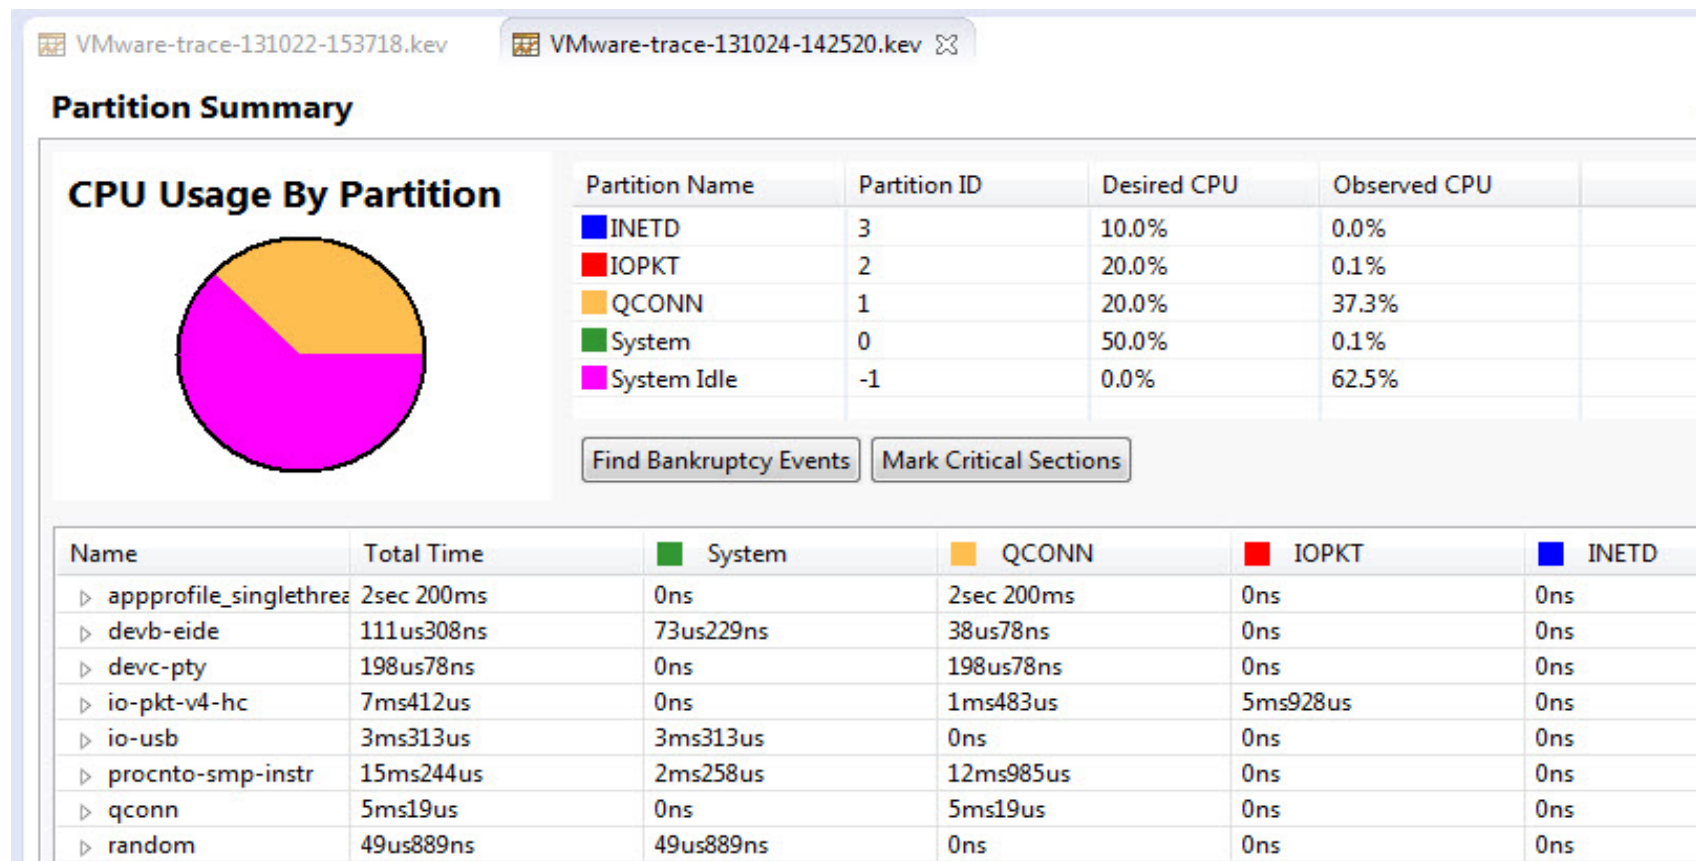
→ Adaptive Partitioning: Partition Summary

Conclusion

Adaptive Partitioning Features

Partition Summary Pane:

- lists partitions and CPU usage within
- helps to find problem areas where critical threads didn't get enough CPU time



Topics:

Overview

Creating a Log

Log Summary

A Quick Tour

Filtering Events and Event Owners

Navigating through a log with the Timeline Pane

Statistics

CPU Activity Pane and CPU Usage Pane

Tying The Trace to Your Code

Multi-core Related Features

Adaptive Partitioning: Partition Summary

→ Conclusion

Conclusion

You learned:

- how the instrumented kernel, tracelogger and the IDE can be used to gather events
- how to use *trace_log*()* to log your own data
- how to control logging of this data from:
 - the IDE
 - the command line using **tracelogger** and
 - from your own code
- how to analyze this data using the IDE